

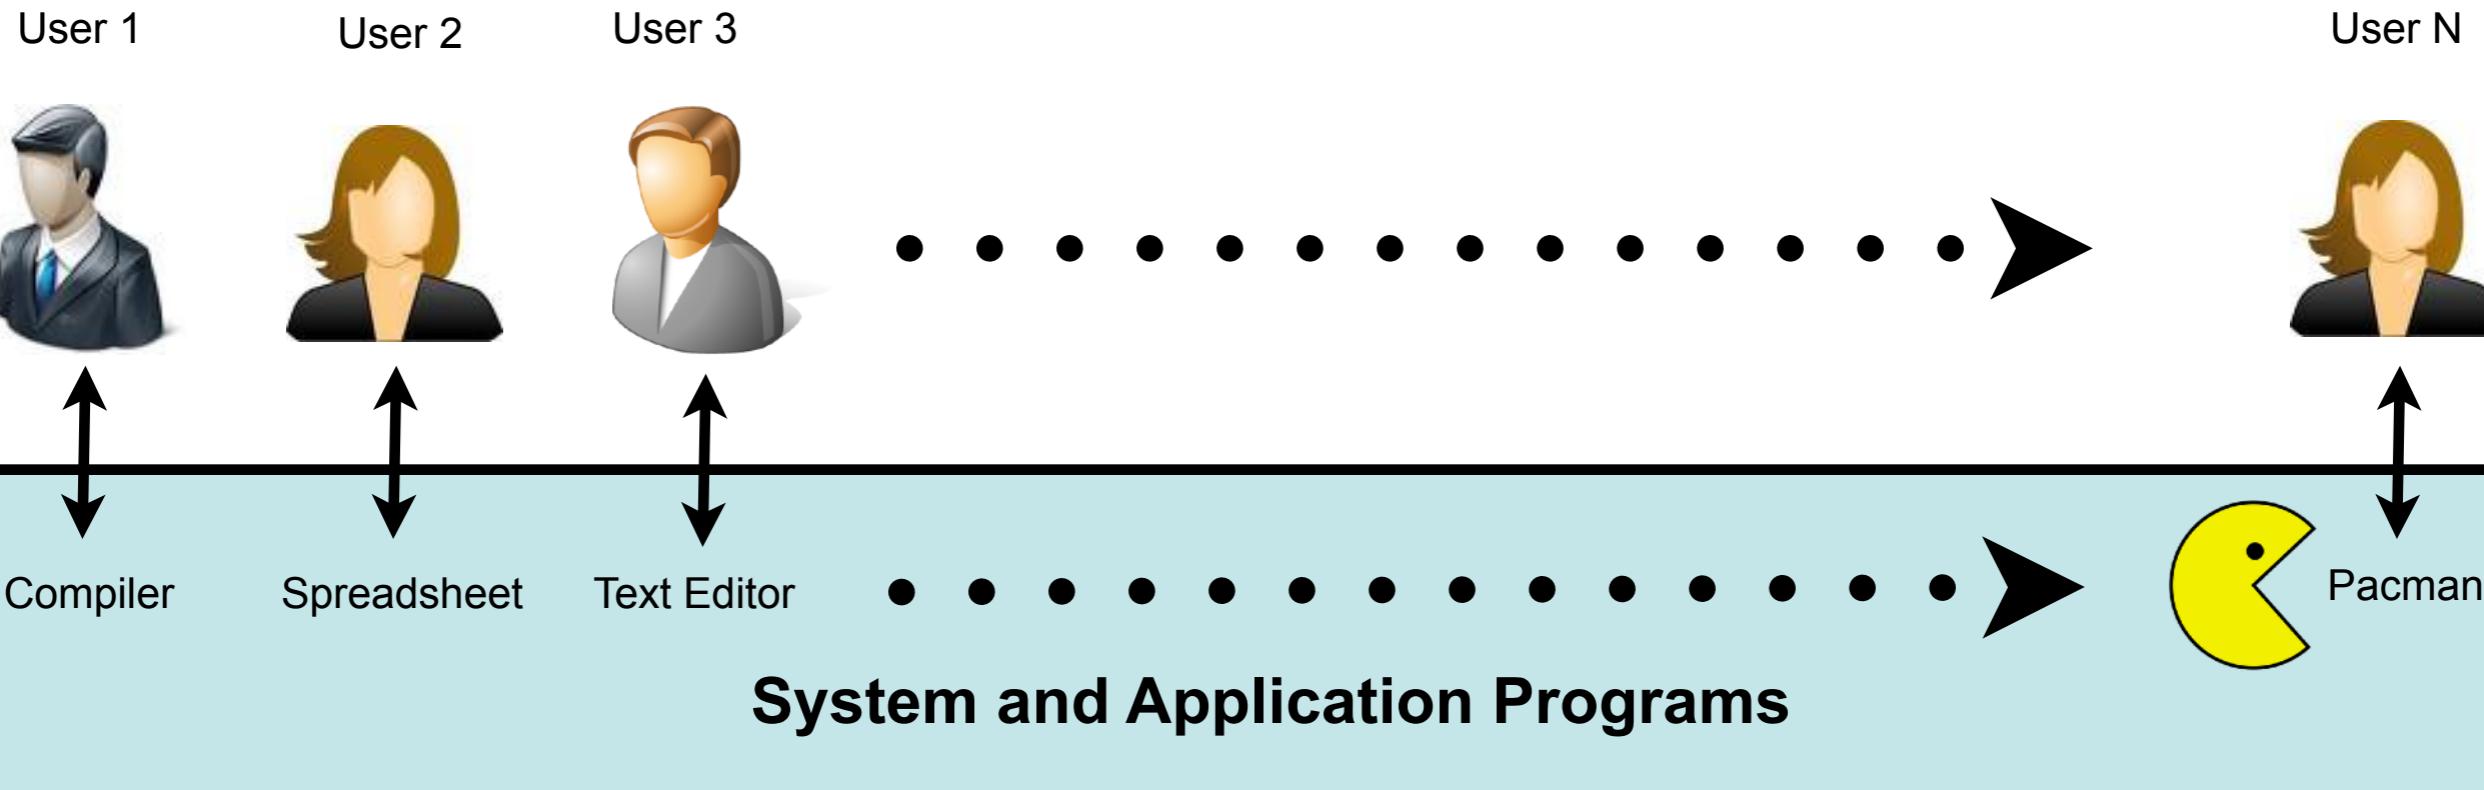
The process concept and inter process communication

Module 2

Lecture

Operating systems 2019

1DT003



Operating System

Controls the hardware and coordinates its use among the various application programs for the various user.

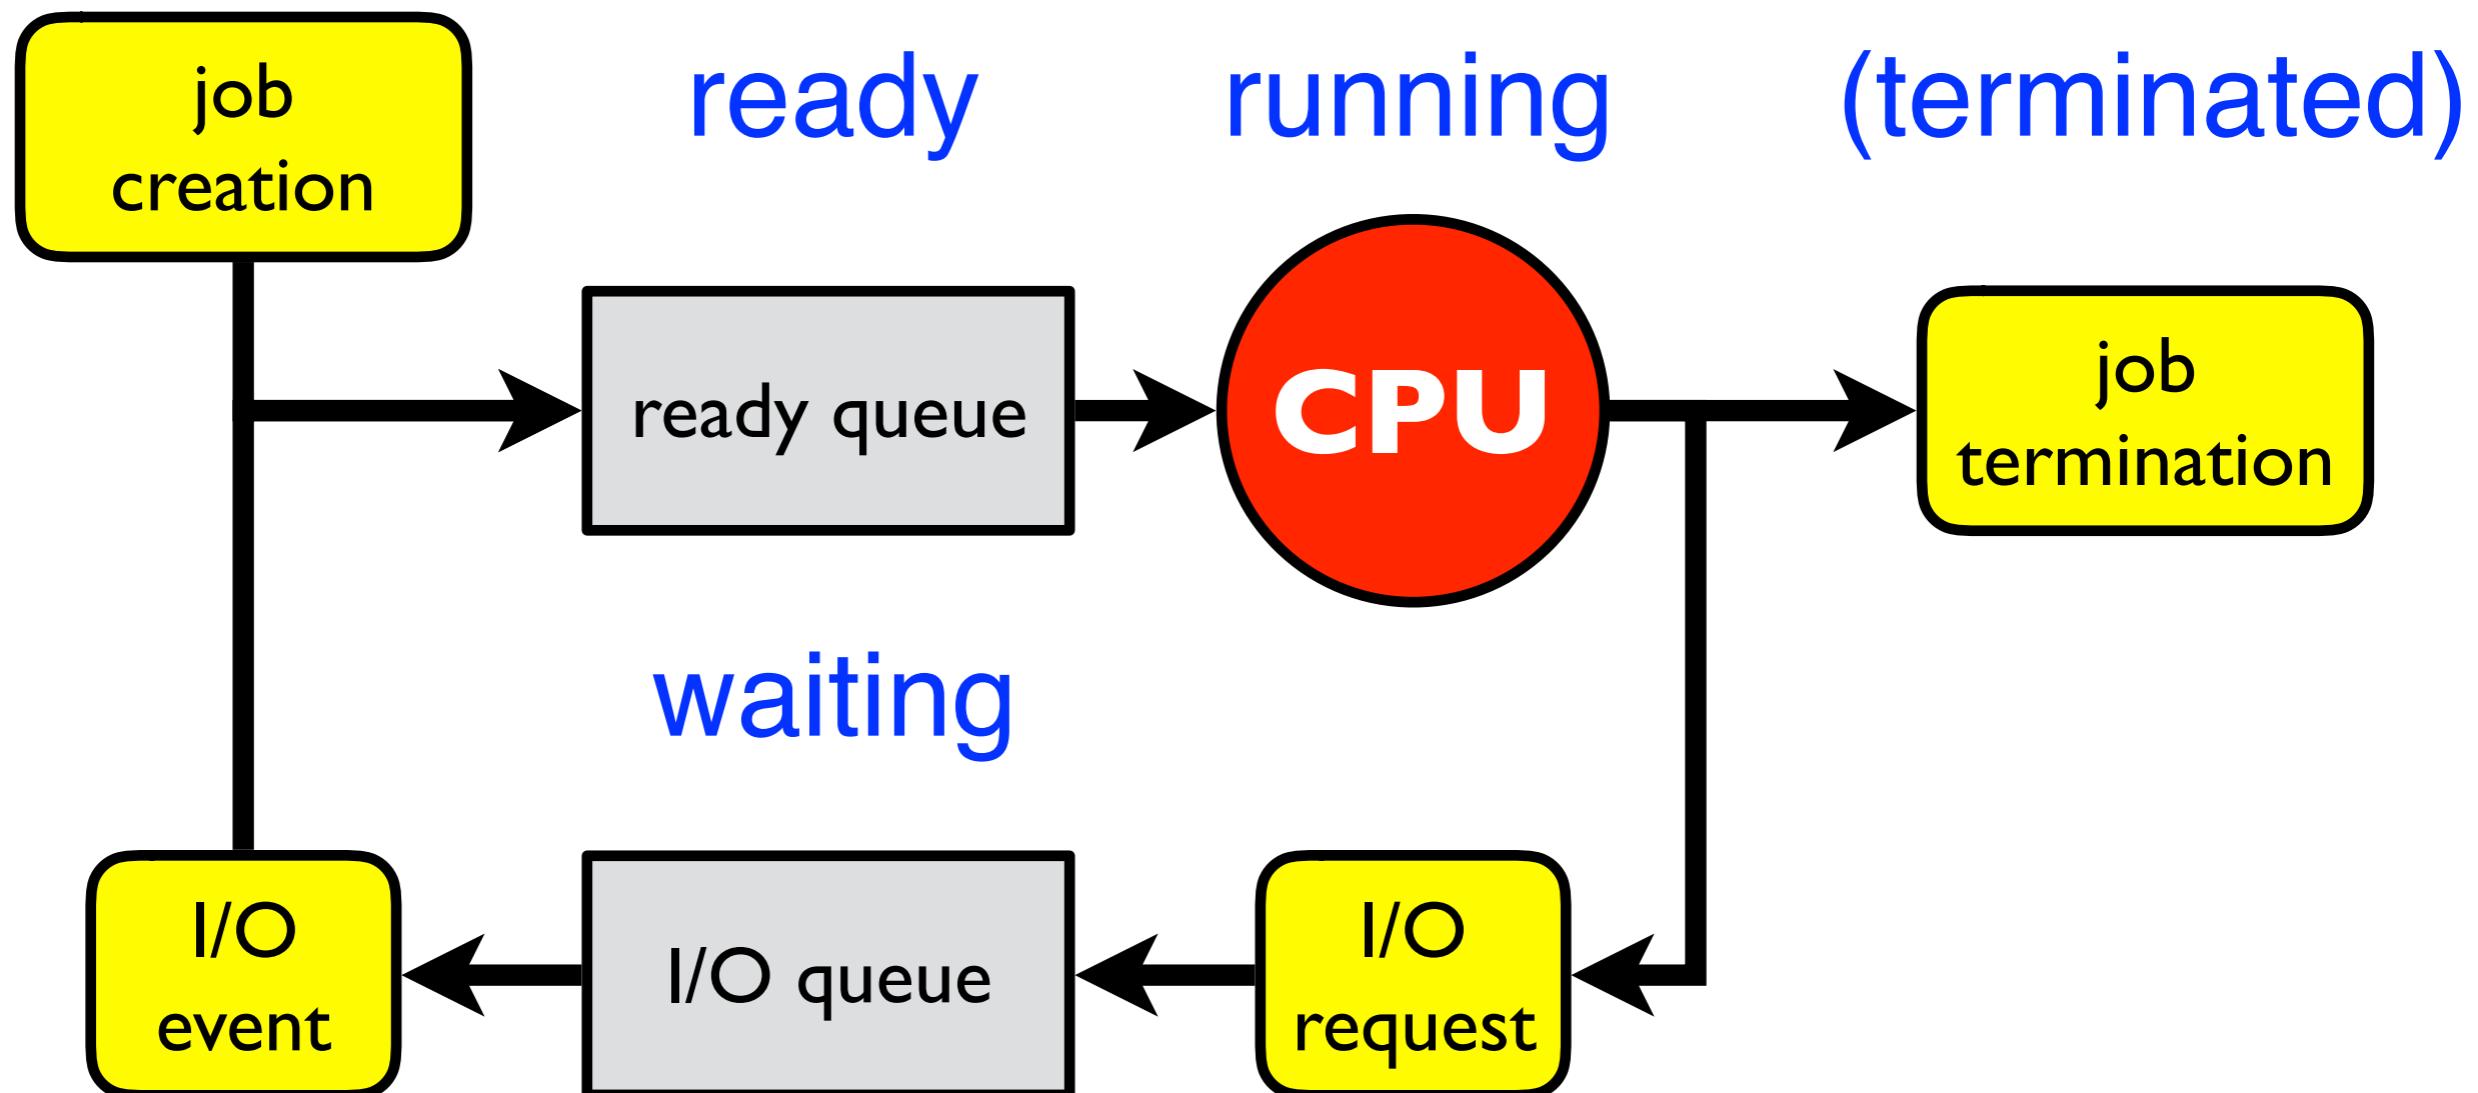
Computer Hardware



Important definitions

Multiprogramming

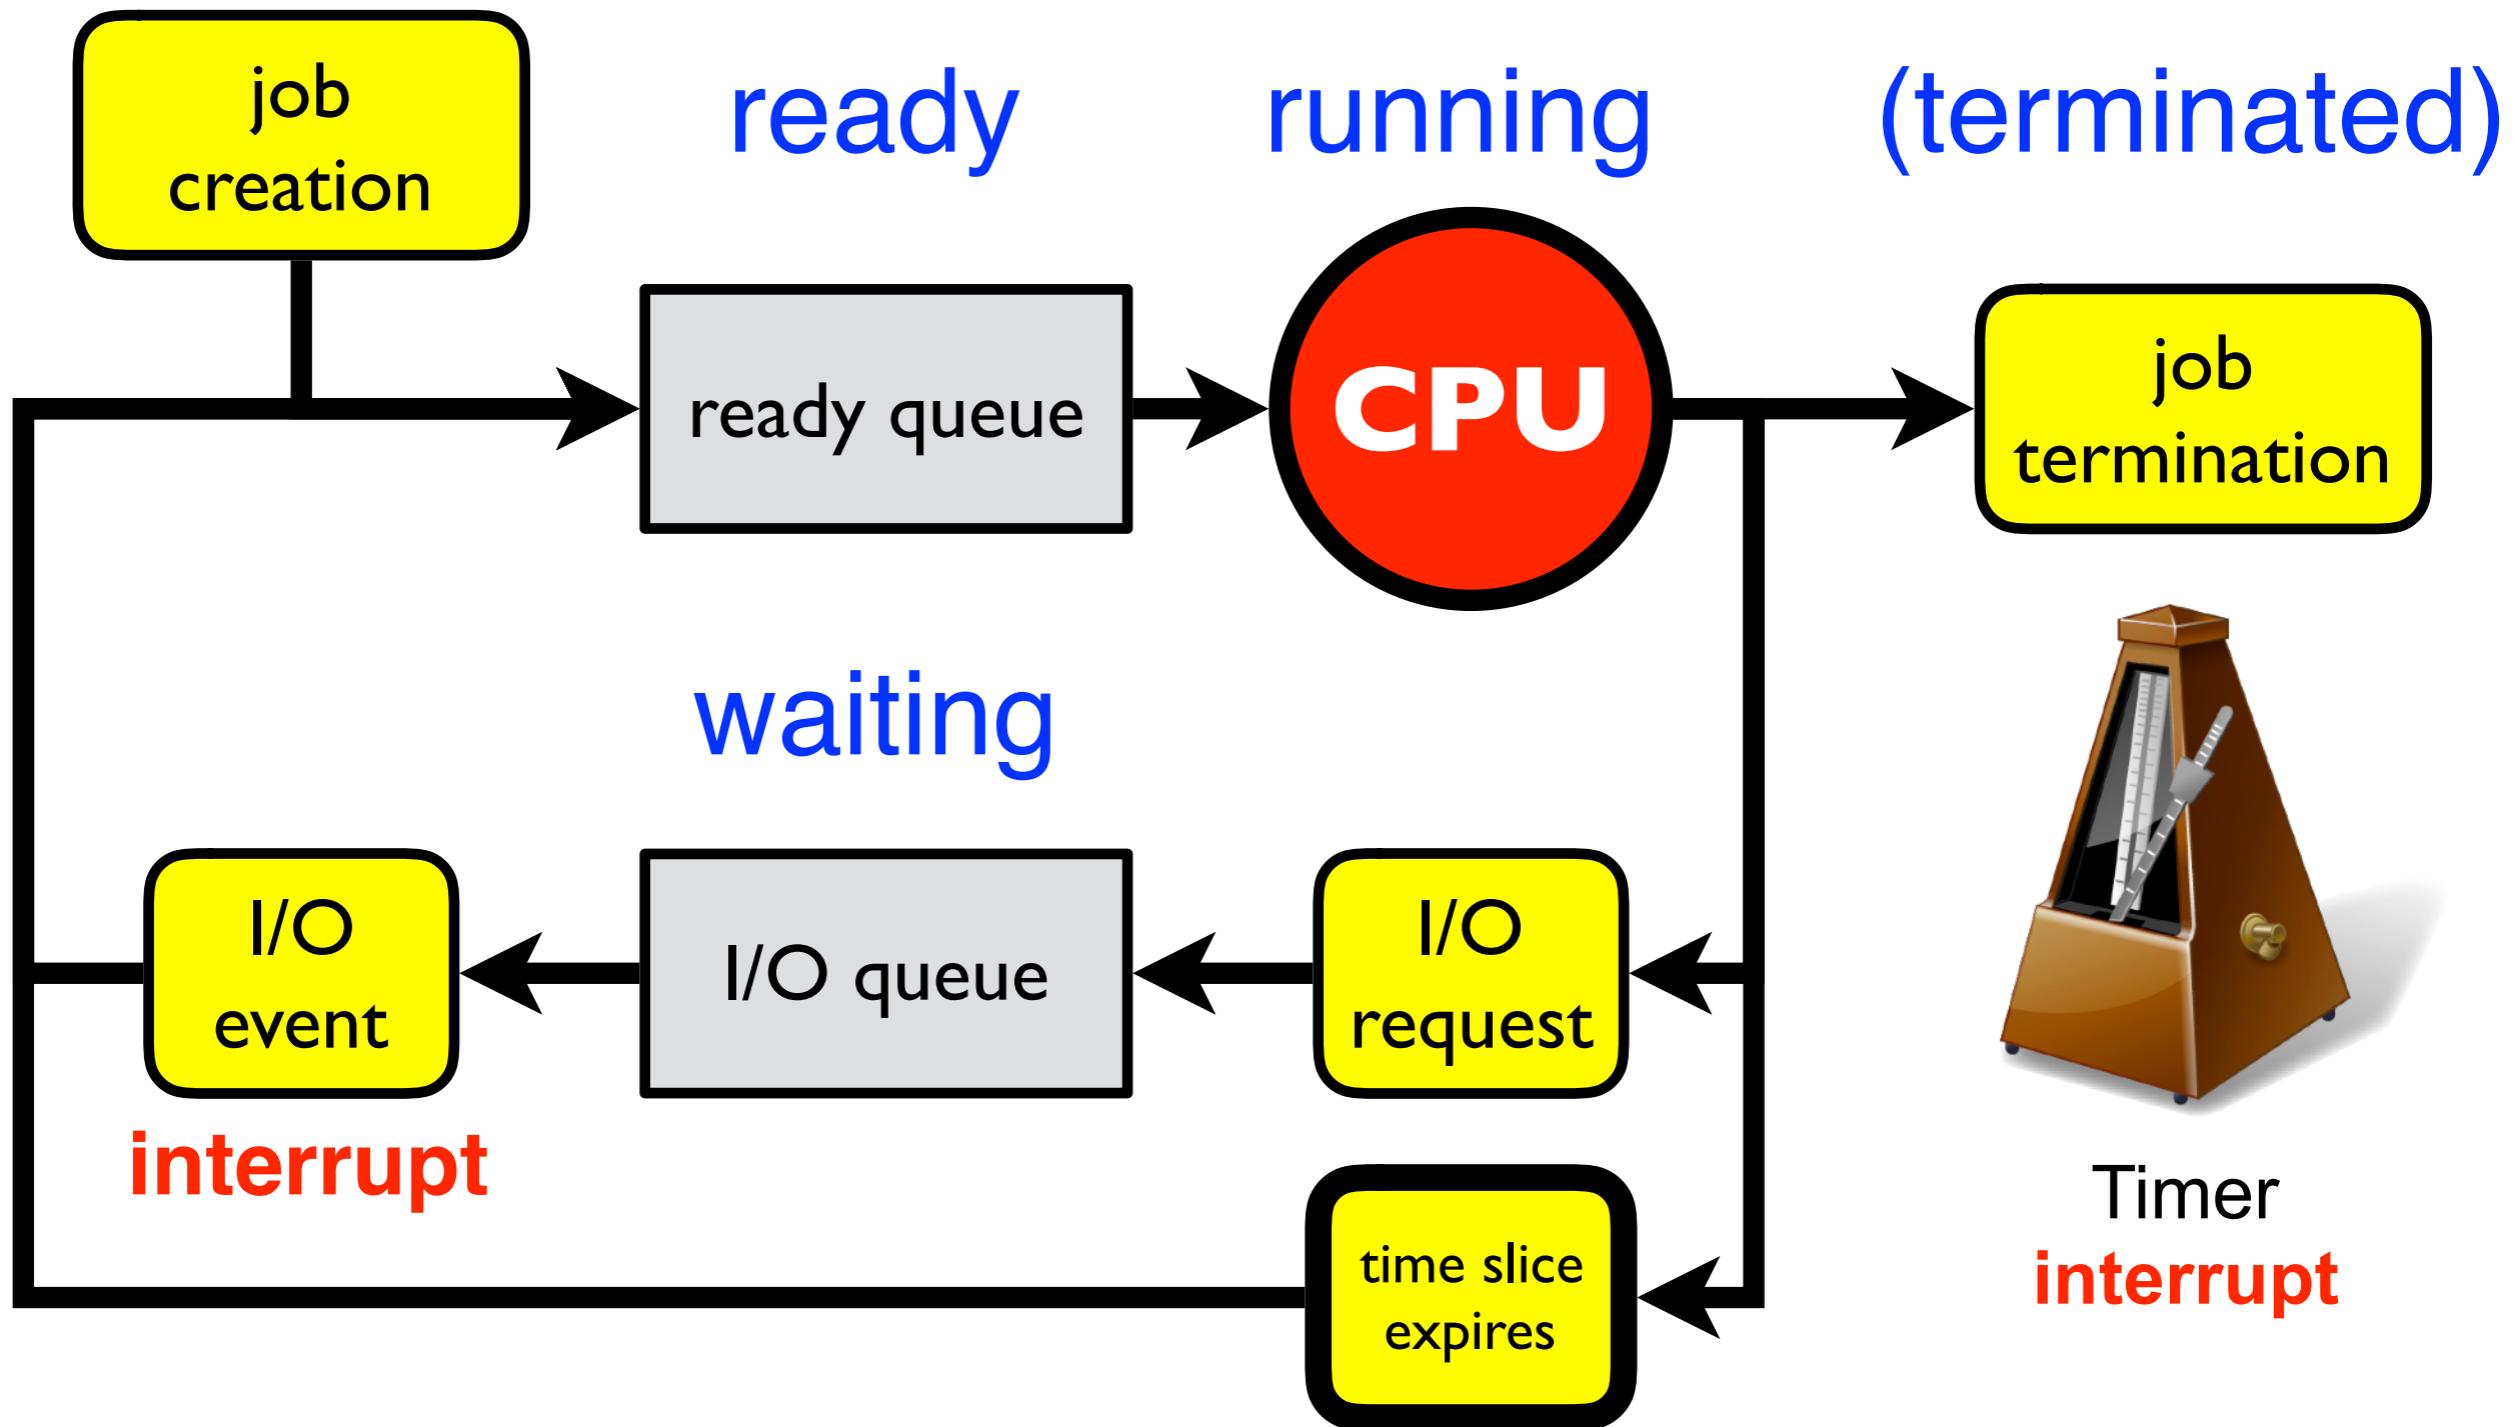
A schematic view of multiprogramming

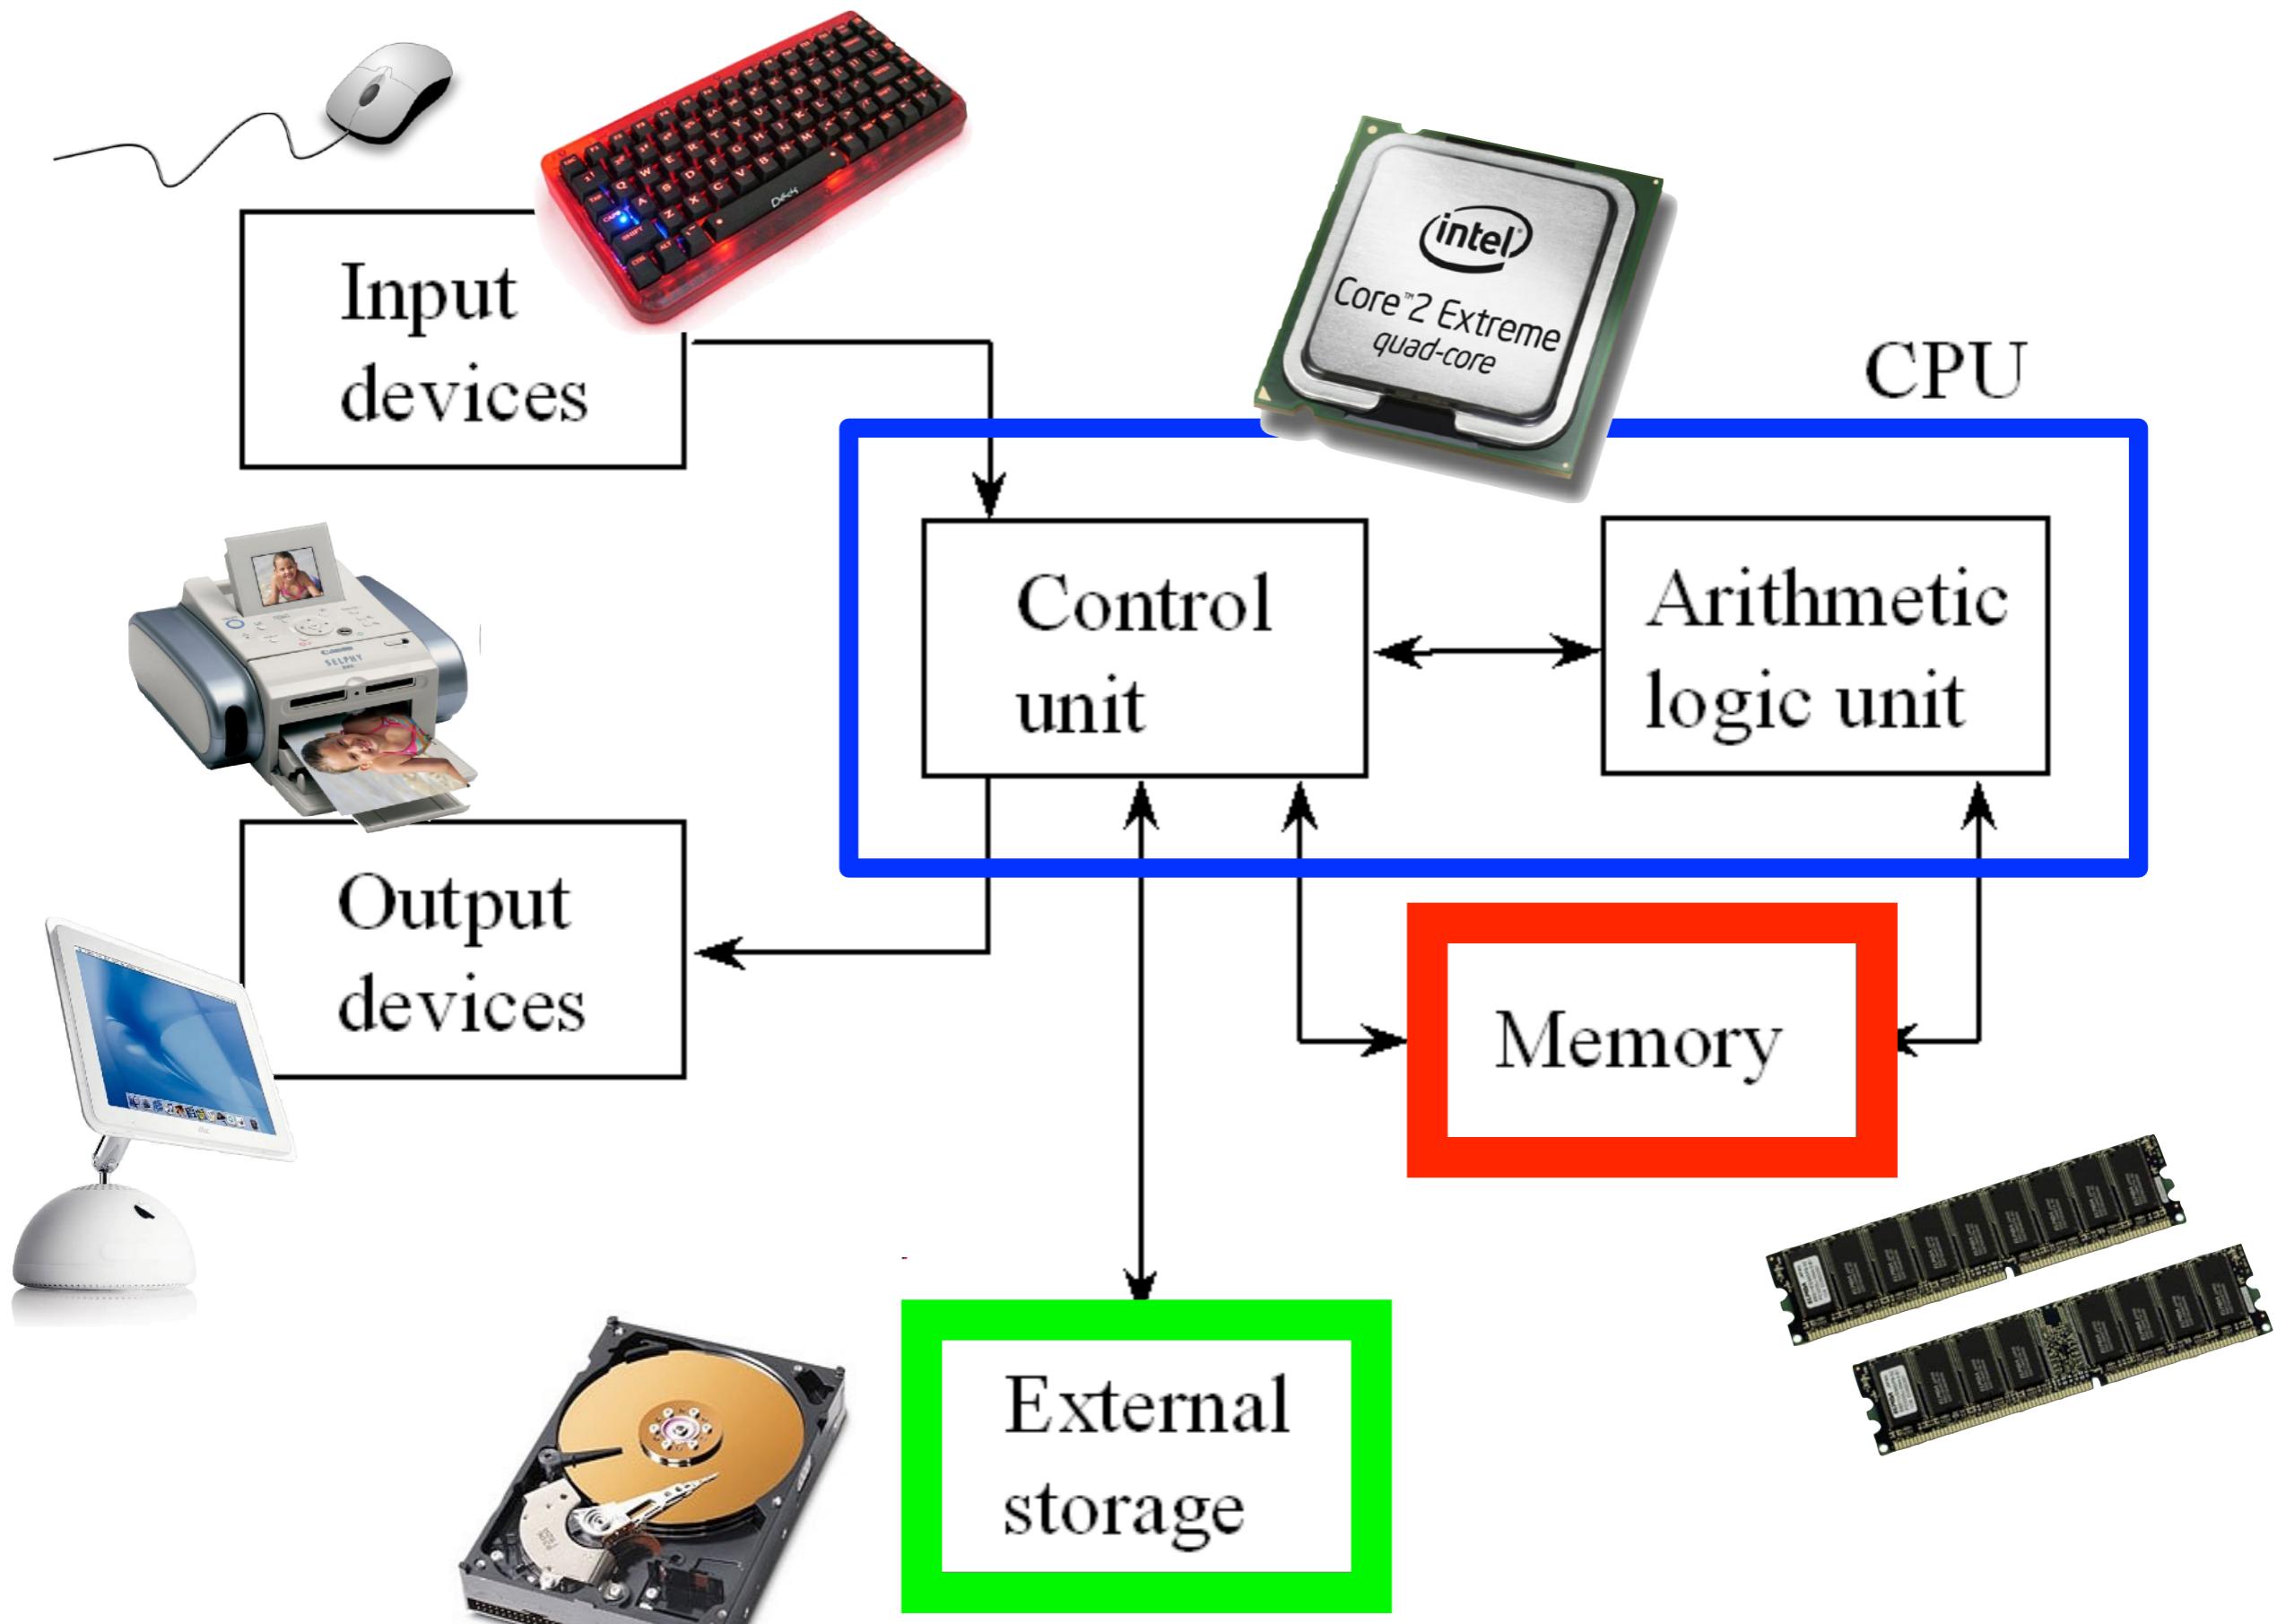


Interrupt used to
signal an external and
asynchronous I/O
event

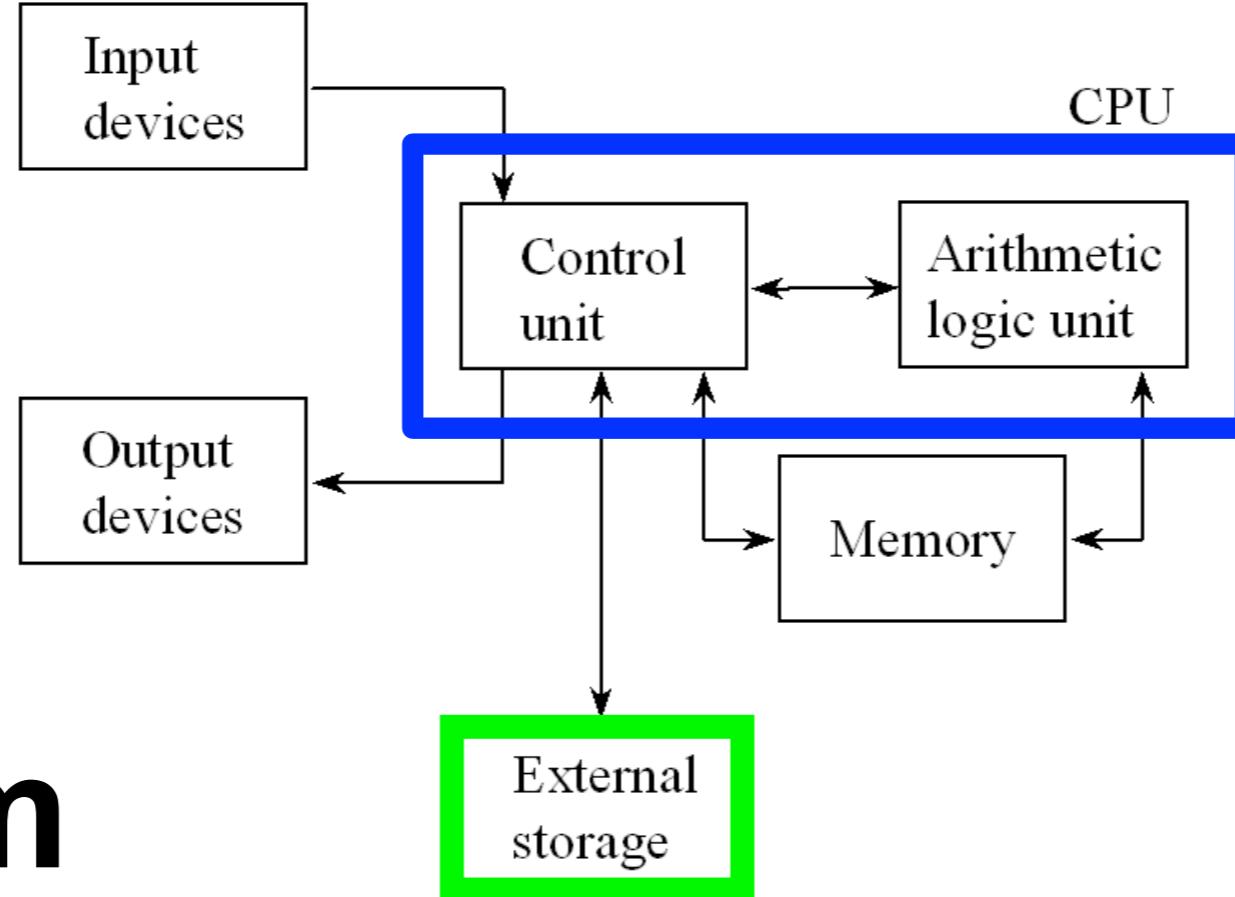
Multitasking

A schematic view of multitasking





The von Neuman modell

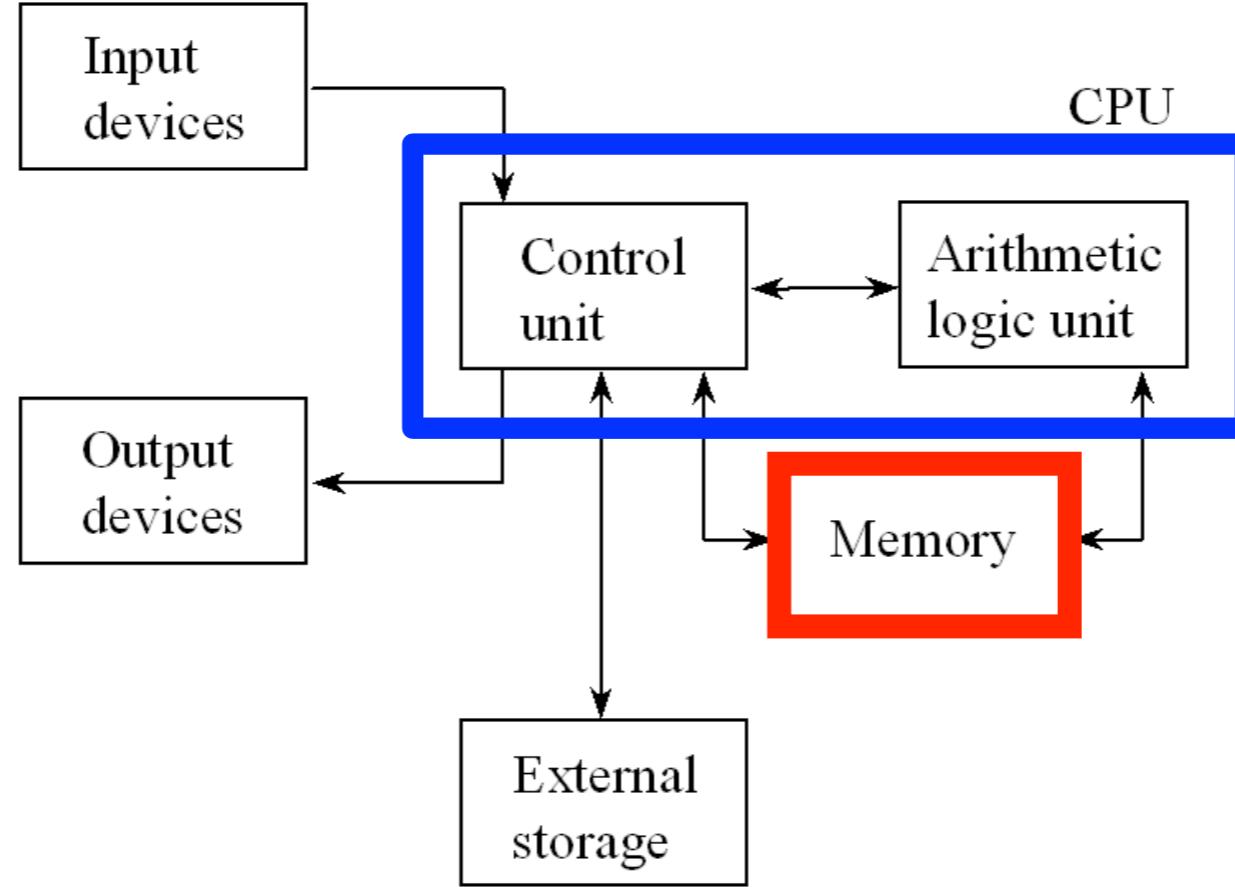


Program

A set of instructions which is in human readable format. A passive entity stored on secondary storage (external storage).

Executable

A compiled form of a program including machine instructions and static data that a computer can load and execute. A passive entity stored on secondary storage (external storage).

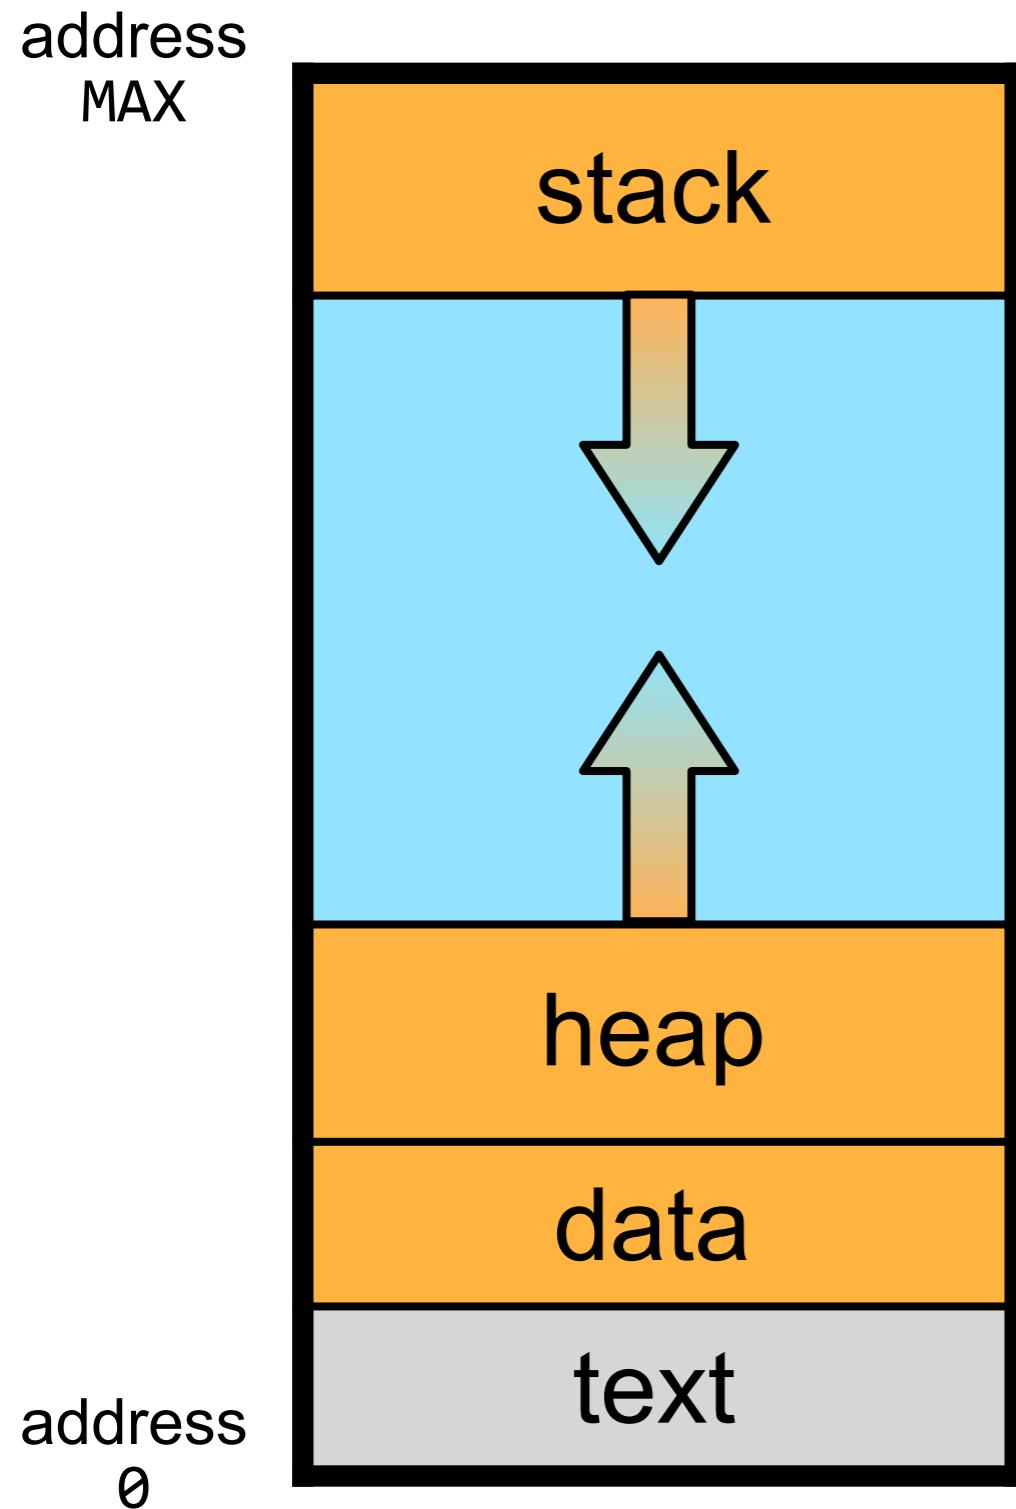


Process

A **program** loaded into **memory** and executing or waiting. A process typically executes for only a short time before it either finishes or needs to perform I/O (waiting).

A process is an active entity and **needs resources** such as **CPU time, memory** etc to execute.

A process in memory



Stack: Temporary data such as function parameters, local variables and return addresses.

The stack grows from high addresses towards lower address.

Heap: Dynamically allocated (malloc) by the program during runtime.

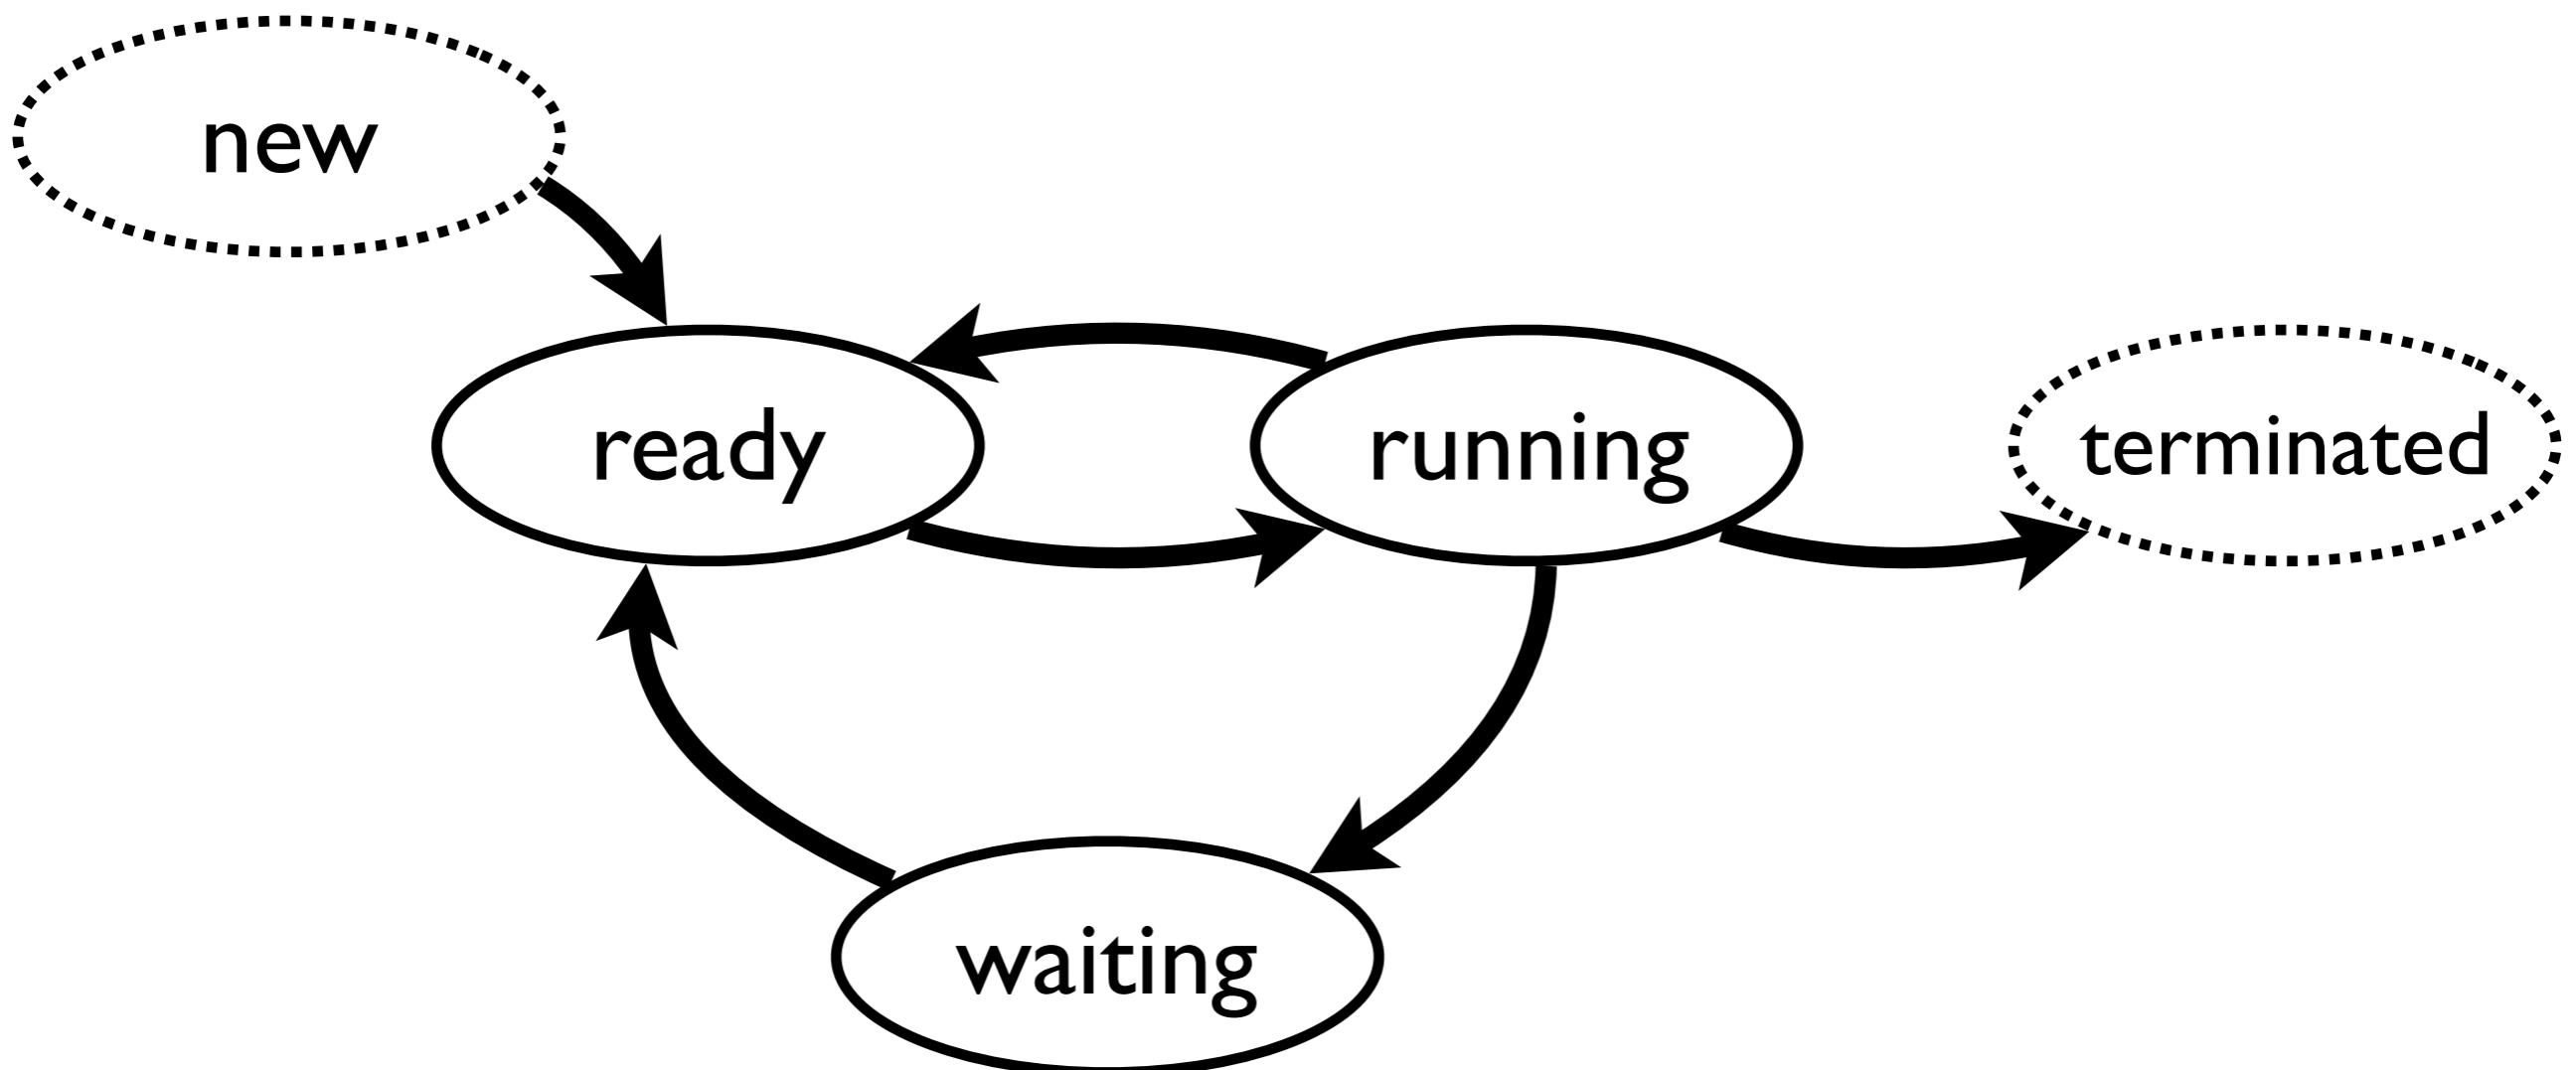
The heap grows from low addresses towards higher addresses.

Data: Statically (known at compile time) global variables and data structures.

Text: The program executable machine instructions.

Process state transitions

In general, a process can be in one of five states: new, ready, running, waiting or terminated. A process transitions between the states according to the following diagram.



Example program

example.c

```
// Example C program

#include <stdlib.h>

int x;

int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```

In this example we start with a single file **example.c** file.

```
$ ls  
example.c
```

file is a standard Unix/Linux program for recognizing the type of data contained in a computer file.

```
$ file example.c  
example.c: C source, ASCII text
```

Here we see that **example.c** is an ordinary C source file.

Compiler

The name compiler is primarily used for programs that translate source code from a high-level programming language to a lower level language (e.g., assembly language, object code, or machine code) to create an executable program.

Executable

A compiler (for example gcc) transforms a program to an executable.

```
$ gcc example.c
```

In this example the compiler **gcc** transformed the **example.c** program into the file **a.out**.

```
$ ls  
a.out  
example.c
```

The compiler **gcc** transformed the **example.c** program into the **executable a.out**.

```
$ file a.out  
  
a.out: ELF 64-bit LSB executable, x86-64,  
version 1 (SYSV), dynamically linked,  
interpreter /lib64/ld-linux-x86-64.so.2,  
for GNU/Linux 2.6.32,  
BuildID[sha1]=4a384d17163f5f8c65cc65b87875  
2ad38827d845, not stripped
```

Process

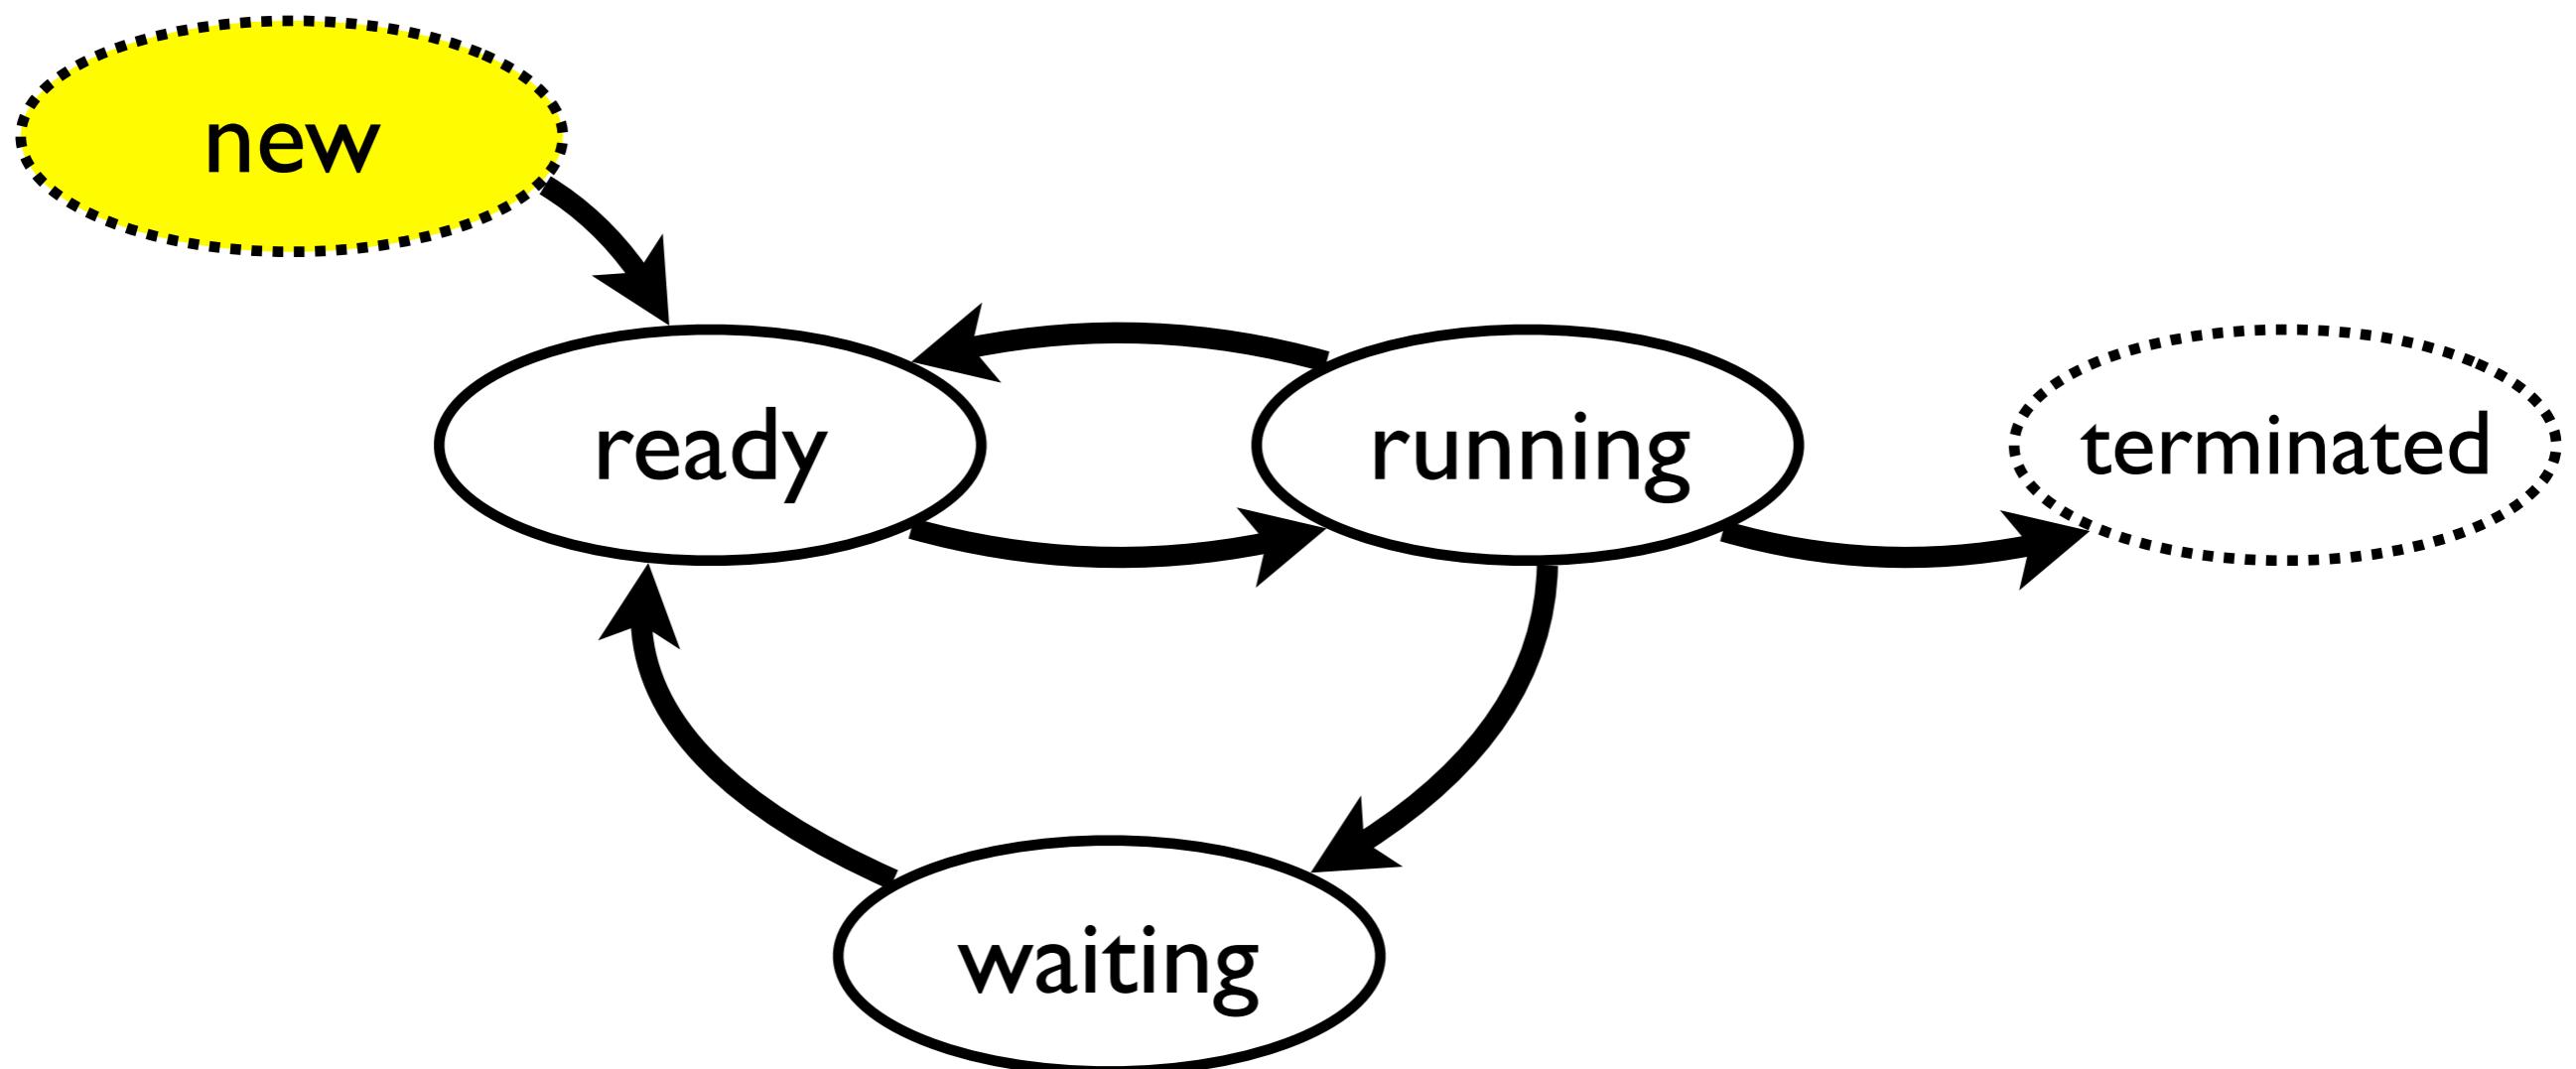
To run the program, the operating system must first create a process.

```
$ ./a.out
```

To create a process, the operating system must allocate memory for the process.

New

To run a program, the operating system must first allocate memory for the process.



Static memory allocation

address MAX

The operating system allocates a blob of memory for the new process.

```
// Example C program

#include <stdlib.h>

int x;

int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```

memory

address 0

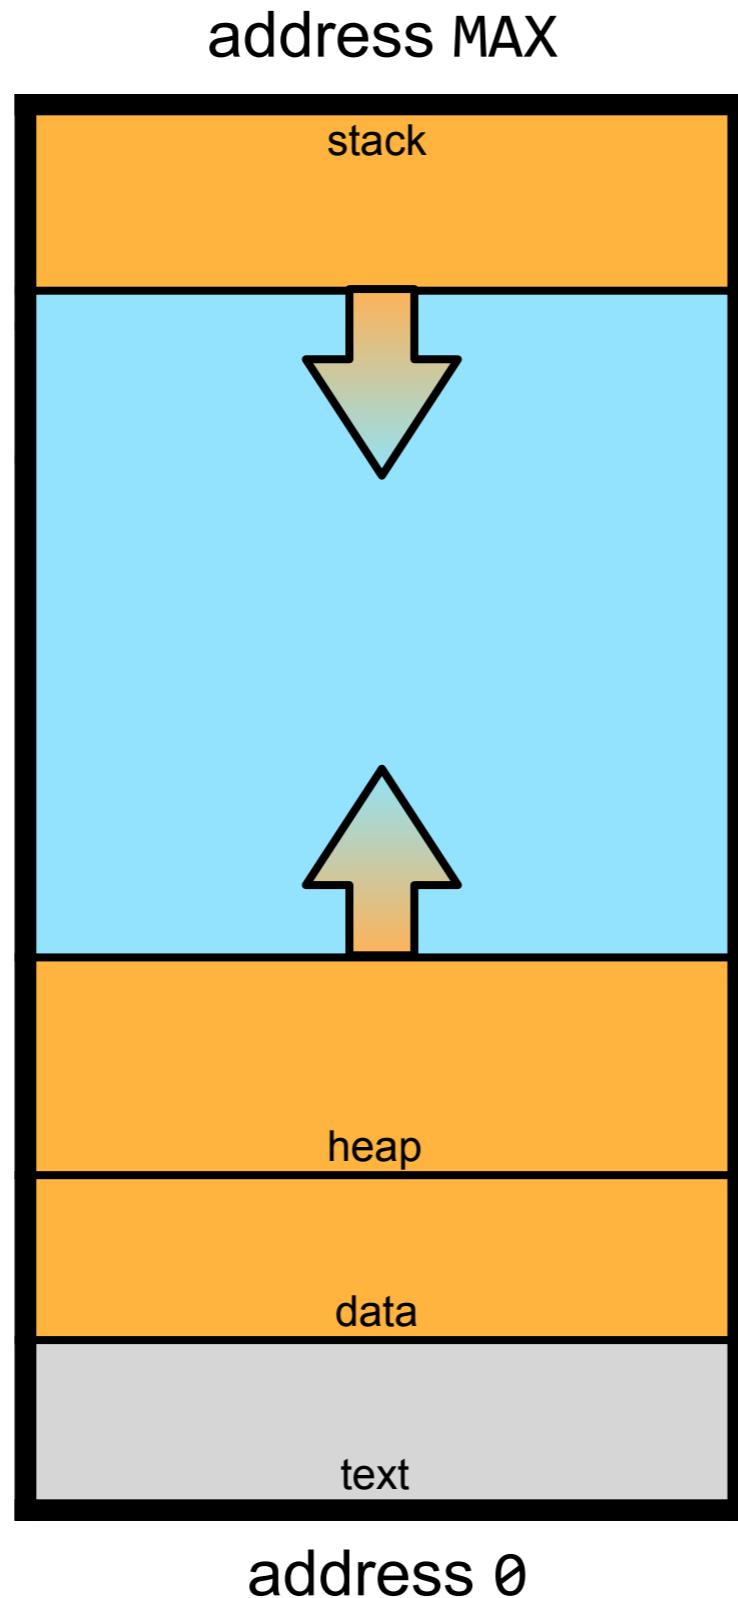
Static memory allocation

```
// Example C program

#include <stdlib.h>

int x;

int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```



The allocated memory is divided into the following segments:

- text
- data
- heap
- stack

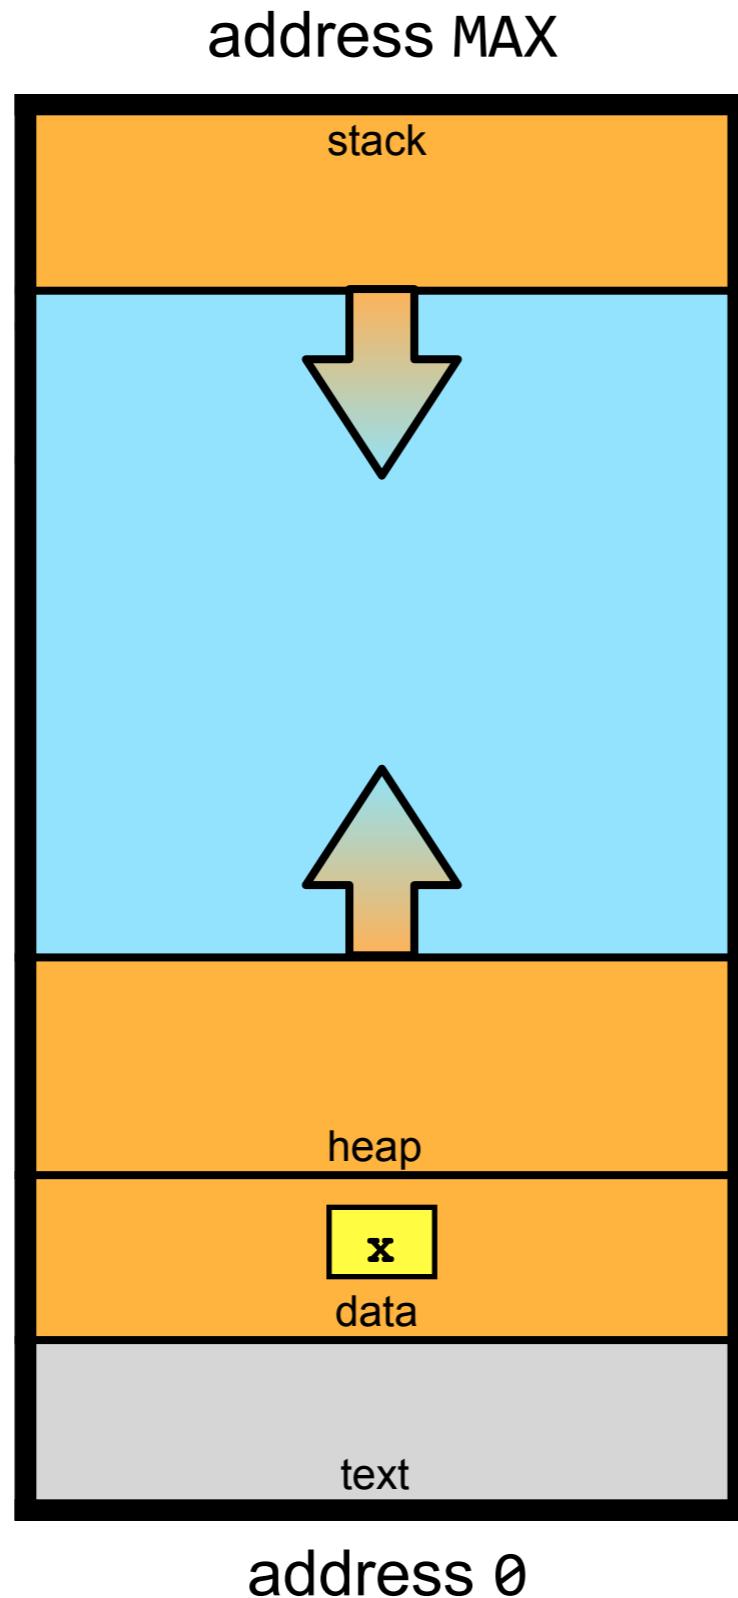
Static memory allocation

```
// Example C program

#include <stdlib.h>

int x;

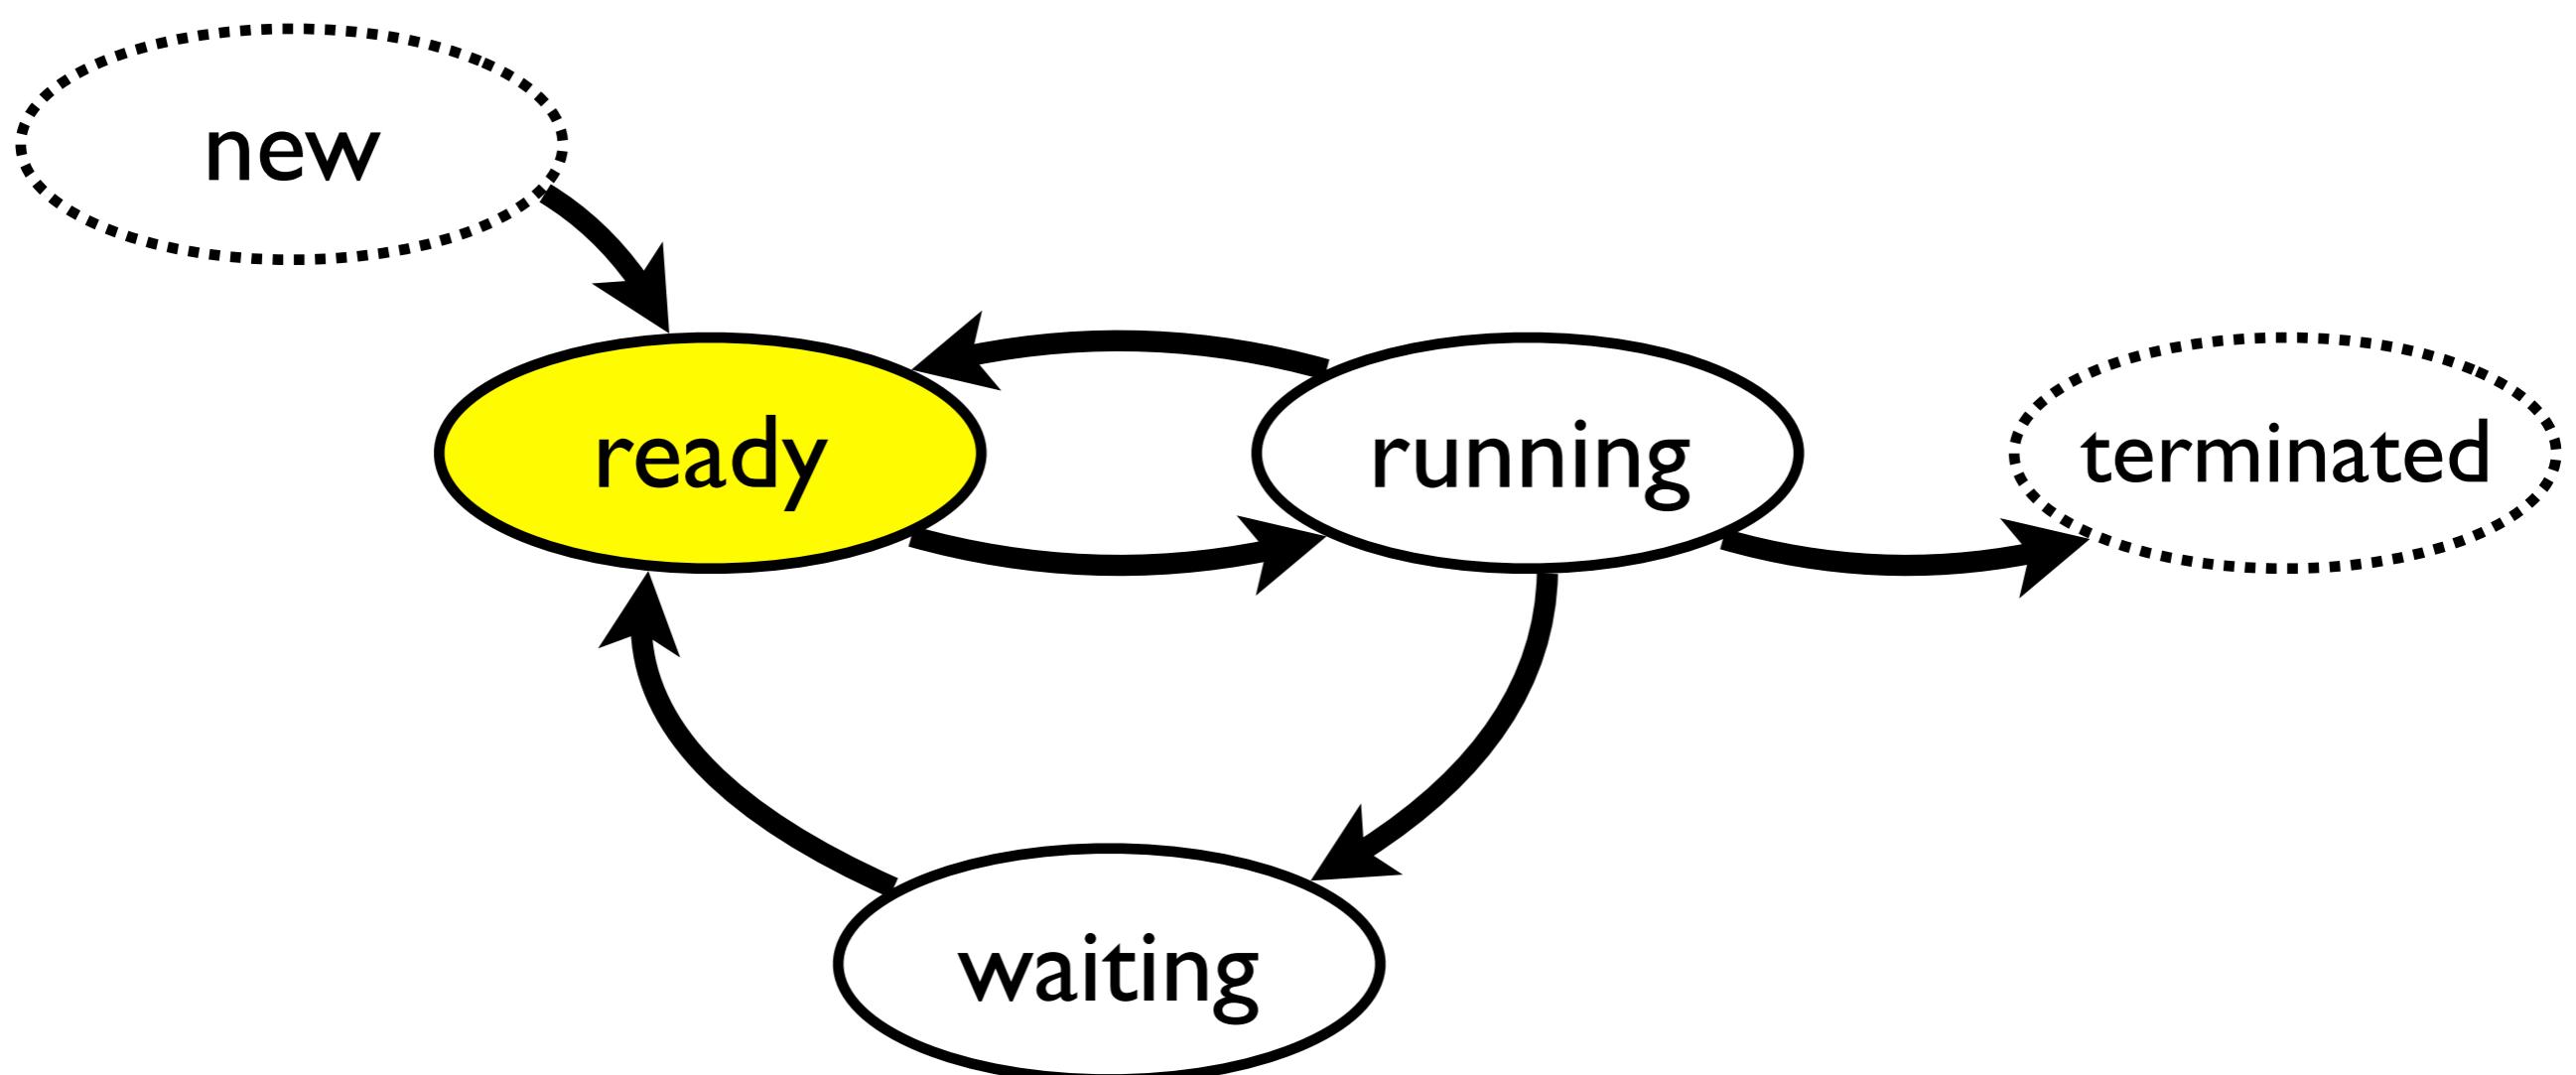
int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```



The size of the needed storage for the **global variable x** is known at compile time and storage for x is allocated in the static **data segment**.


Ready

After the operating system has allocated memory for the process it is ready to execute and changes state from new to ready.



Running

When the process is selected to execute it changes state from ready to running.




Dynamic memory allocation

```
// Example C program

#include <stdlib.h>

int x;

int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```




The size of the needed storage for the **local variables** **y** and **str** are also known at compile time but these are only needed within main and storage is allocated on the **stack**.

In general, the value of the **parameter** to the malloc might not be known at compile time. In this example the argument to malloc, the number 50 is pushed onto the **stack**.

Note: A compiler may also decide to put arguments in certain register.

Dynamic memory allocation

```
// Example C program  
  
#include <stdlib.h>  
  
int x;  
  
int main(void) {  
    int y;  
    char* str;  
    str = malloc(50);  
    exit(EXIT_SUCCESS);  
}
```



Now **malloc** has dynamically allocated memory on the heap.

The variable **str** stores the address to the first of the 50 bytes allocated on the heap, i.e., **str** is a **pointer**.

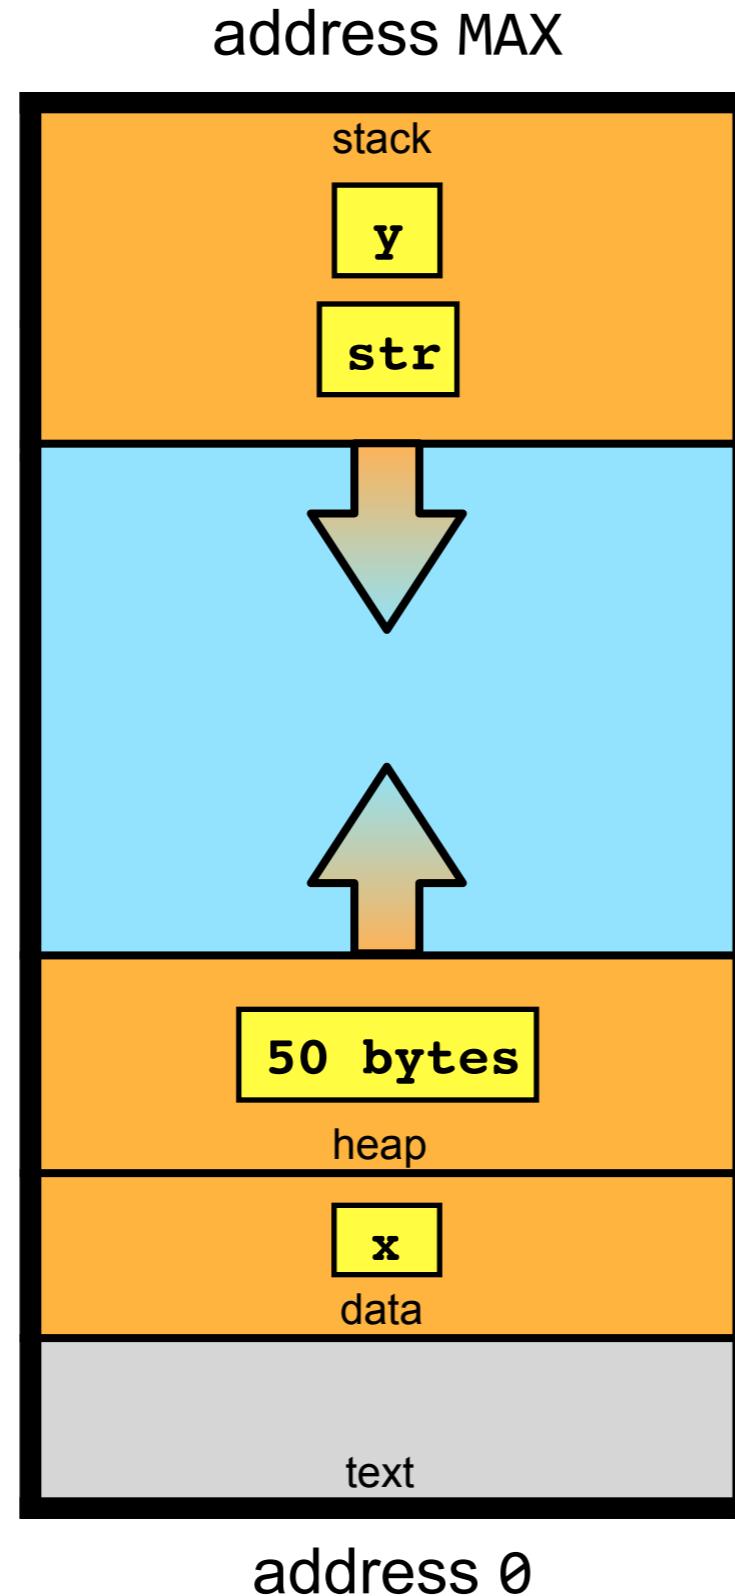
Dynamic memory deallocation

```
// Example C program

#include <stdlib.h>

int x;

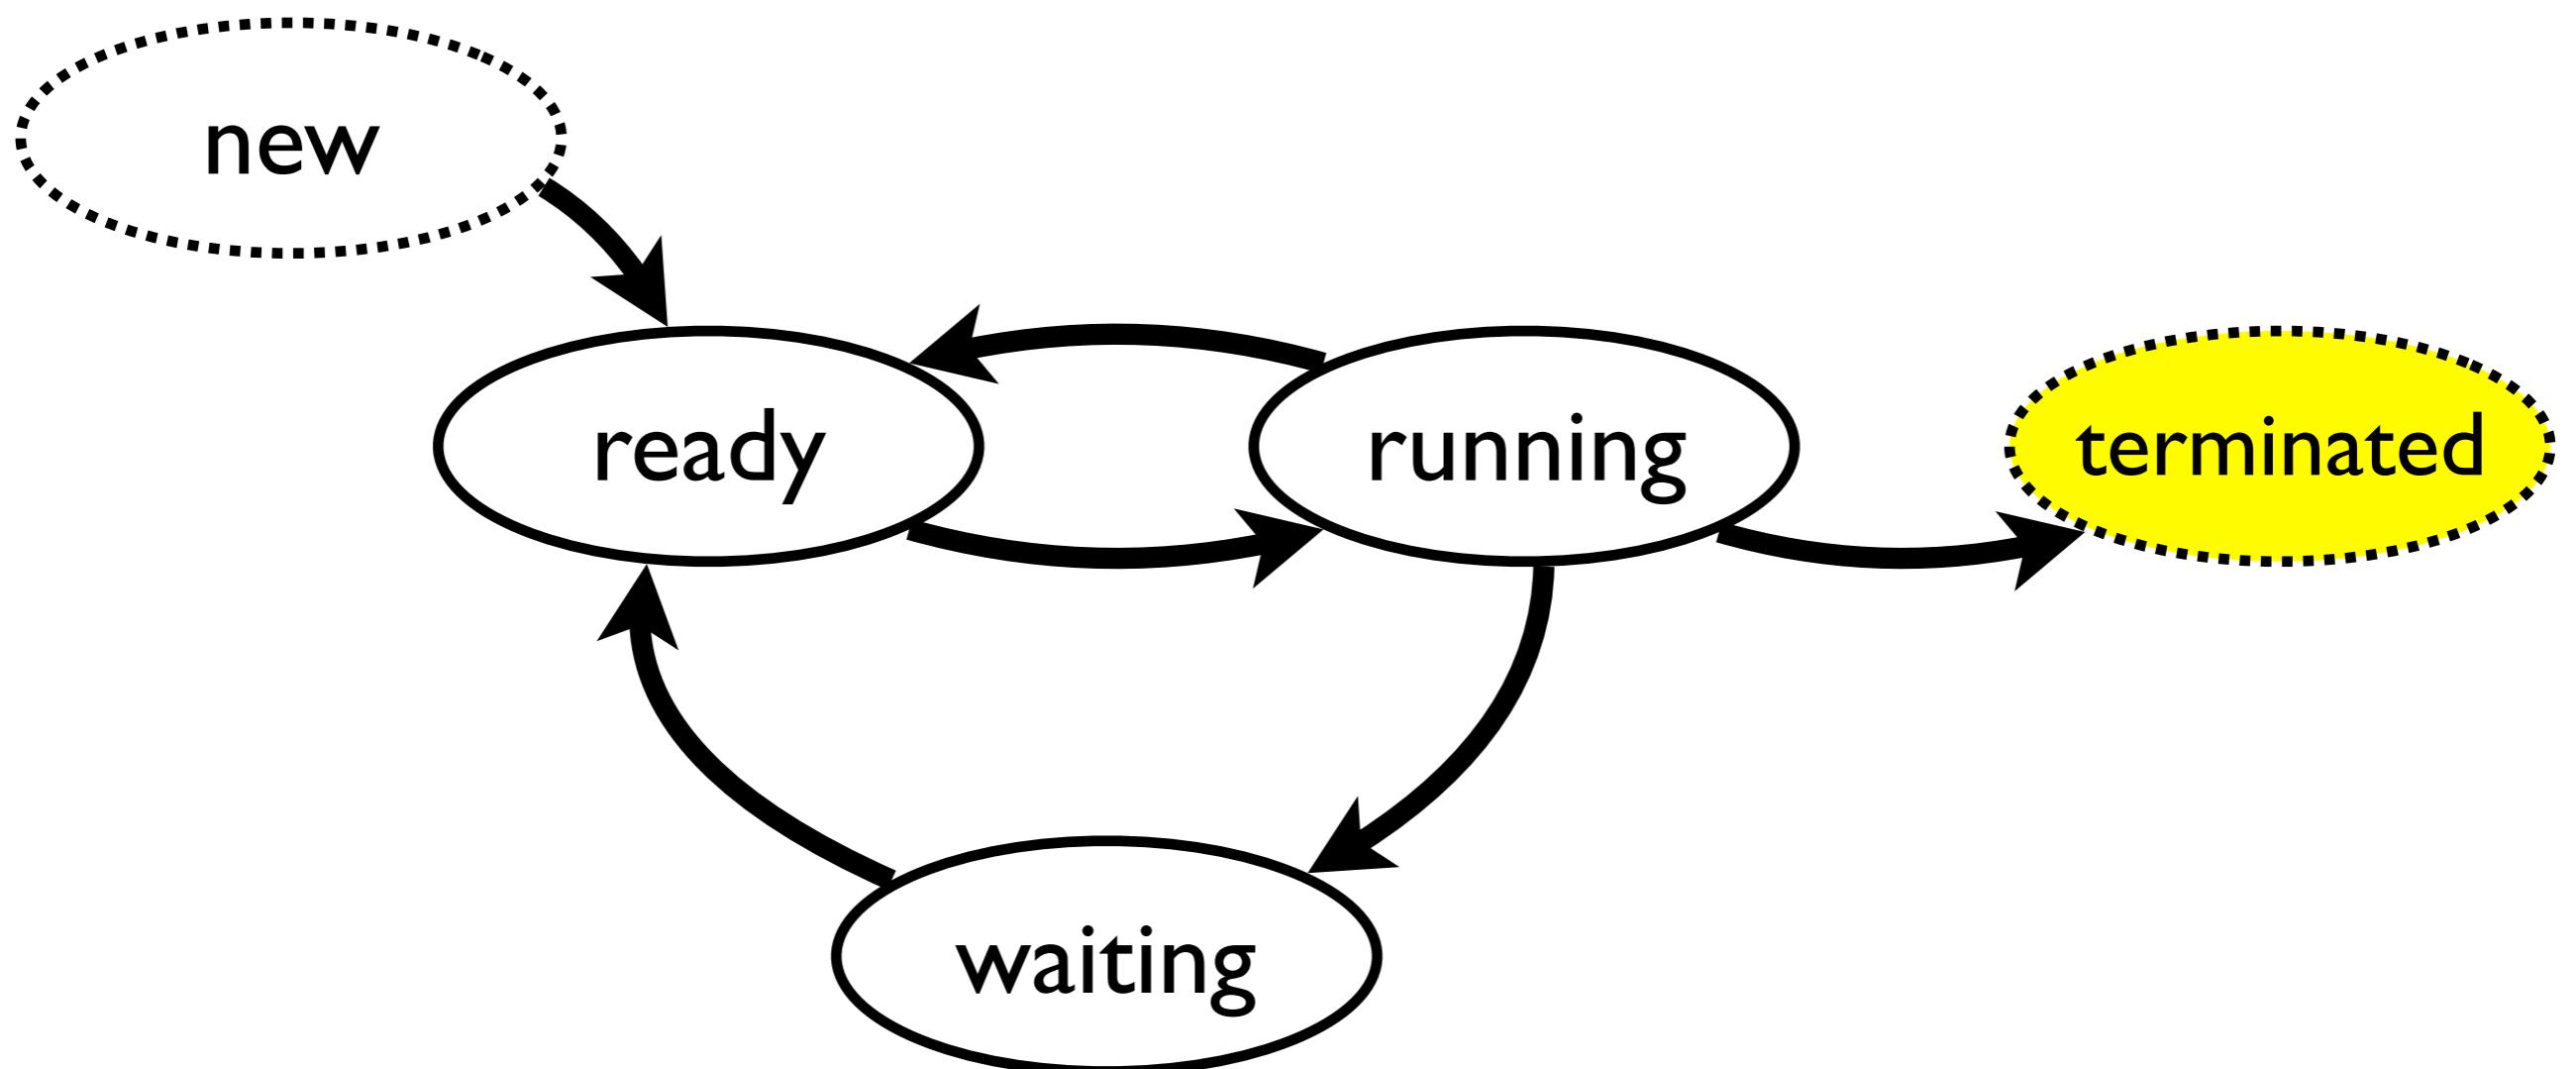
int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```



When malloc returns, the value 50, the argument used when calling malloc is no longer needed and is popped from the stack.

Termination

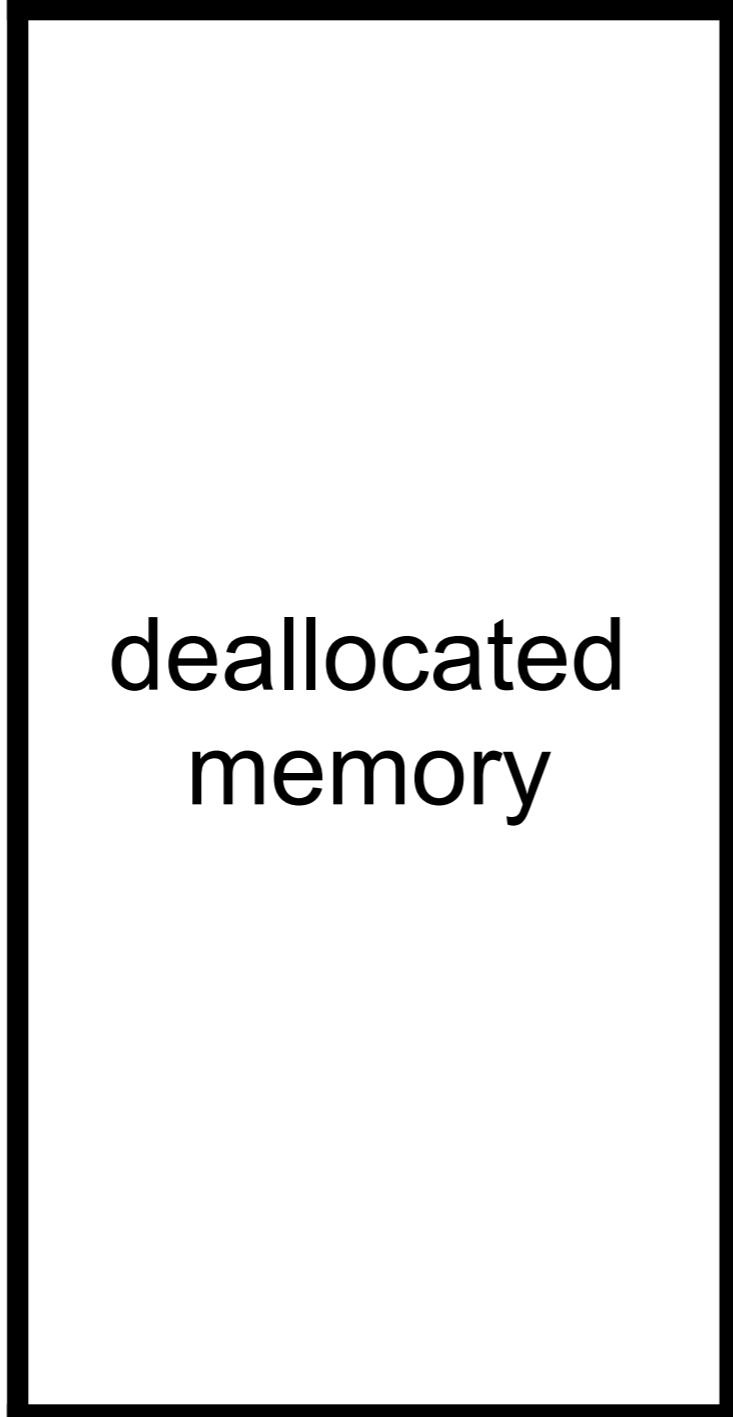
When the process terminates the operating system can deallocate the memory used by the process.



Process termination

address MAX

The program terminates with a call to `exit()` and the memory allocated to the process can be deallocated.



deallocated
memory

A large black rectangular box represents memory. Inside the box, the text "deallocated memory" is centered. Above the box, the text "address MAX" is positioned. Below the box, the text "address 0" is positioned. The box is approximately centered vertically between the two address labels.

address 0

```
// Example C program

#include <stdlib.h>

int x;

int main(void) {
    int y;
    char* str;
    str = malloc(50);
    exit(EXIT_SUCCESS);
}
```

Process Control Block (PCB)

Process Control Block (PCB)

The process control block (PCB) is a data structure in the operating system kernel containing the information needed to manage a particular process.

Source https://en.wikipedia.org/wiki/Process_control_block

2018-01-21

In brief, the PCB serves as the repository for any information that may vary from process to process.

Example of information stored in the PCB

Process Control Block (PCB)
Process id (PID)
Process state (new, ready, running, waiting or terminated)
CPU Context
I/O status information
Memory management information
CPU scheduling information

Context

switch

Context switch (1)

In computing, a context switch is the process of storing the state of a process or of a thread, so that it can be restored and execution resumed from the same point later. This allows multiple processes to share a single CPU, and is an essential feature of a multitasking operating system.

Context switch (2)

It is not enough to save and restore the CPU context. Other information, for example the process state must also be saved and restored.

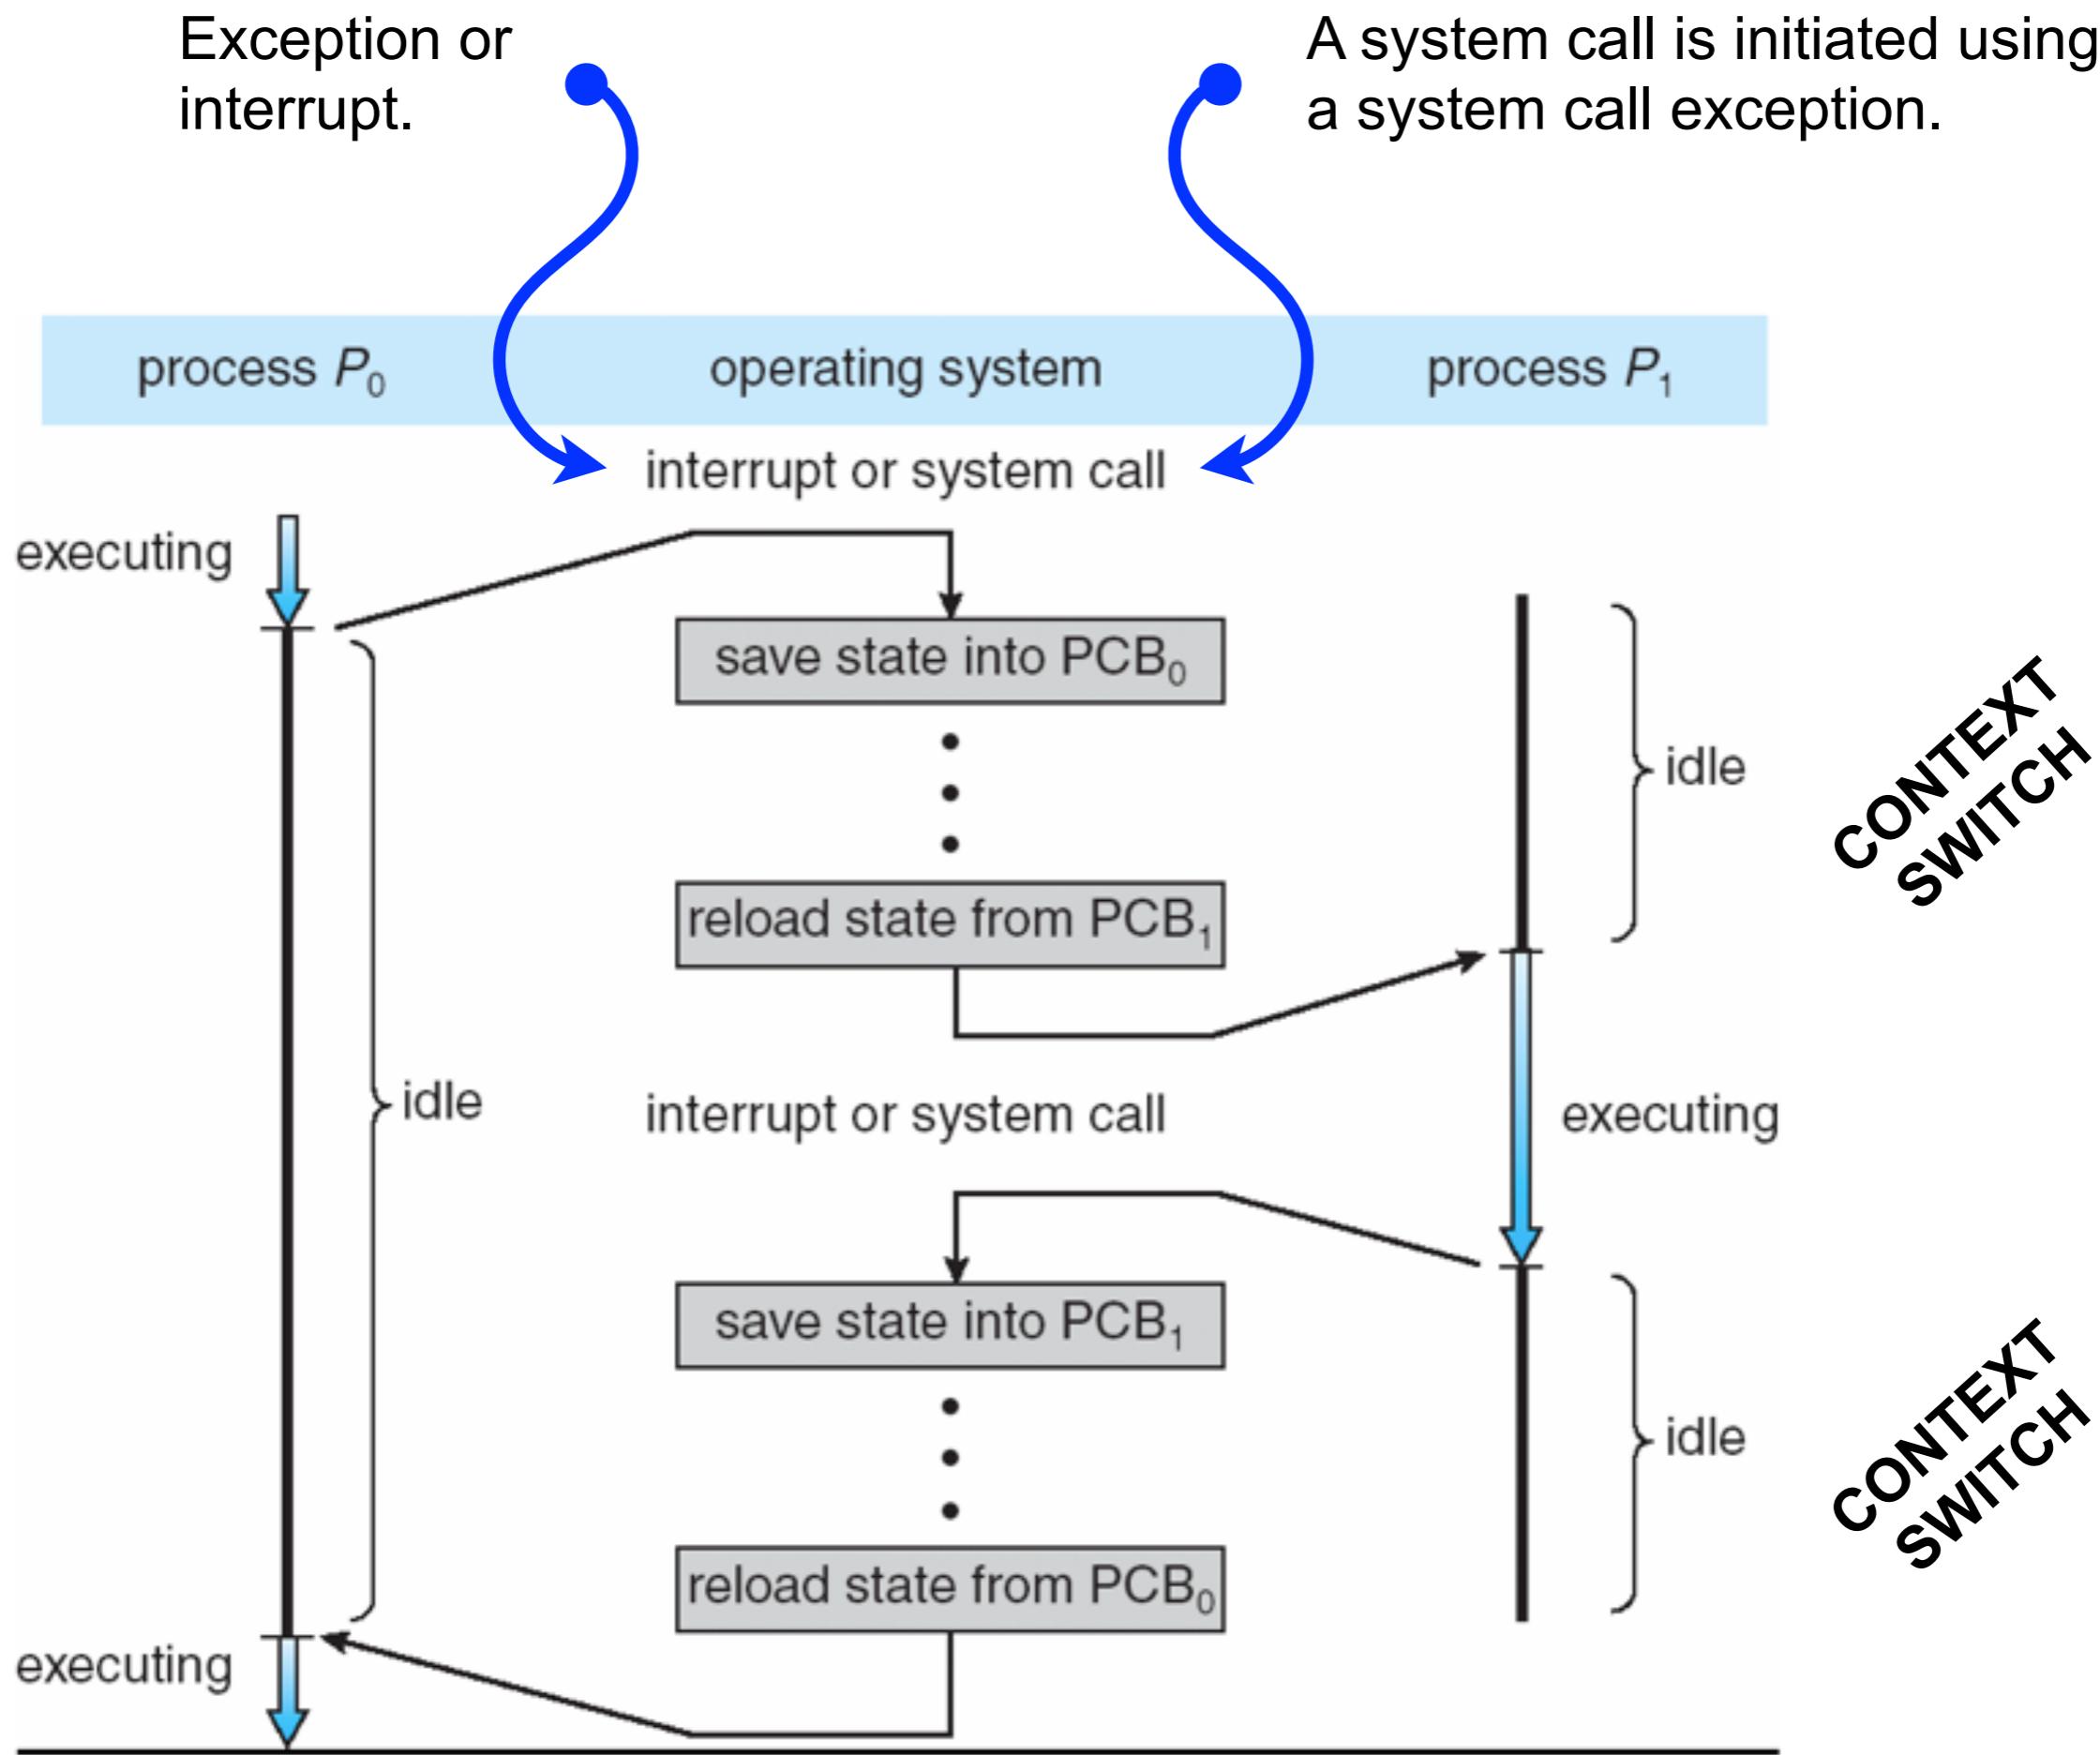
Context switch (3)

When CPU switches to another process, the system must save the state of the old process and restore the state for the new process via a **context switch**.

Context of a process is represented in the PCB.

Context-switch time is **overhead**; the system does no useful work while switching.

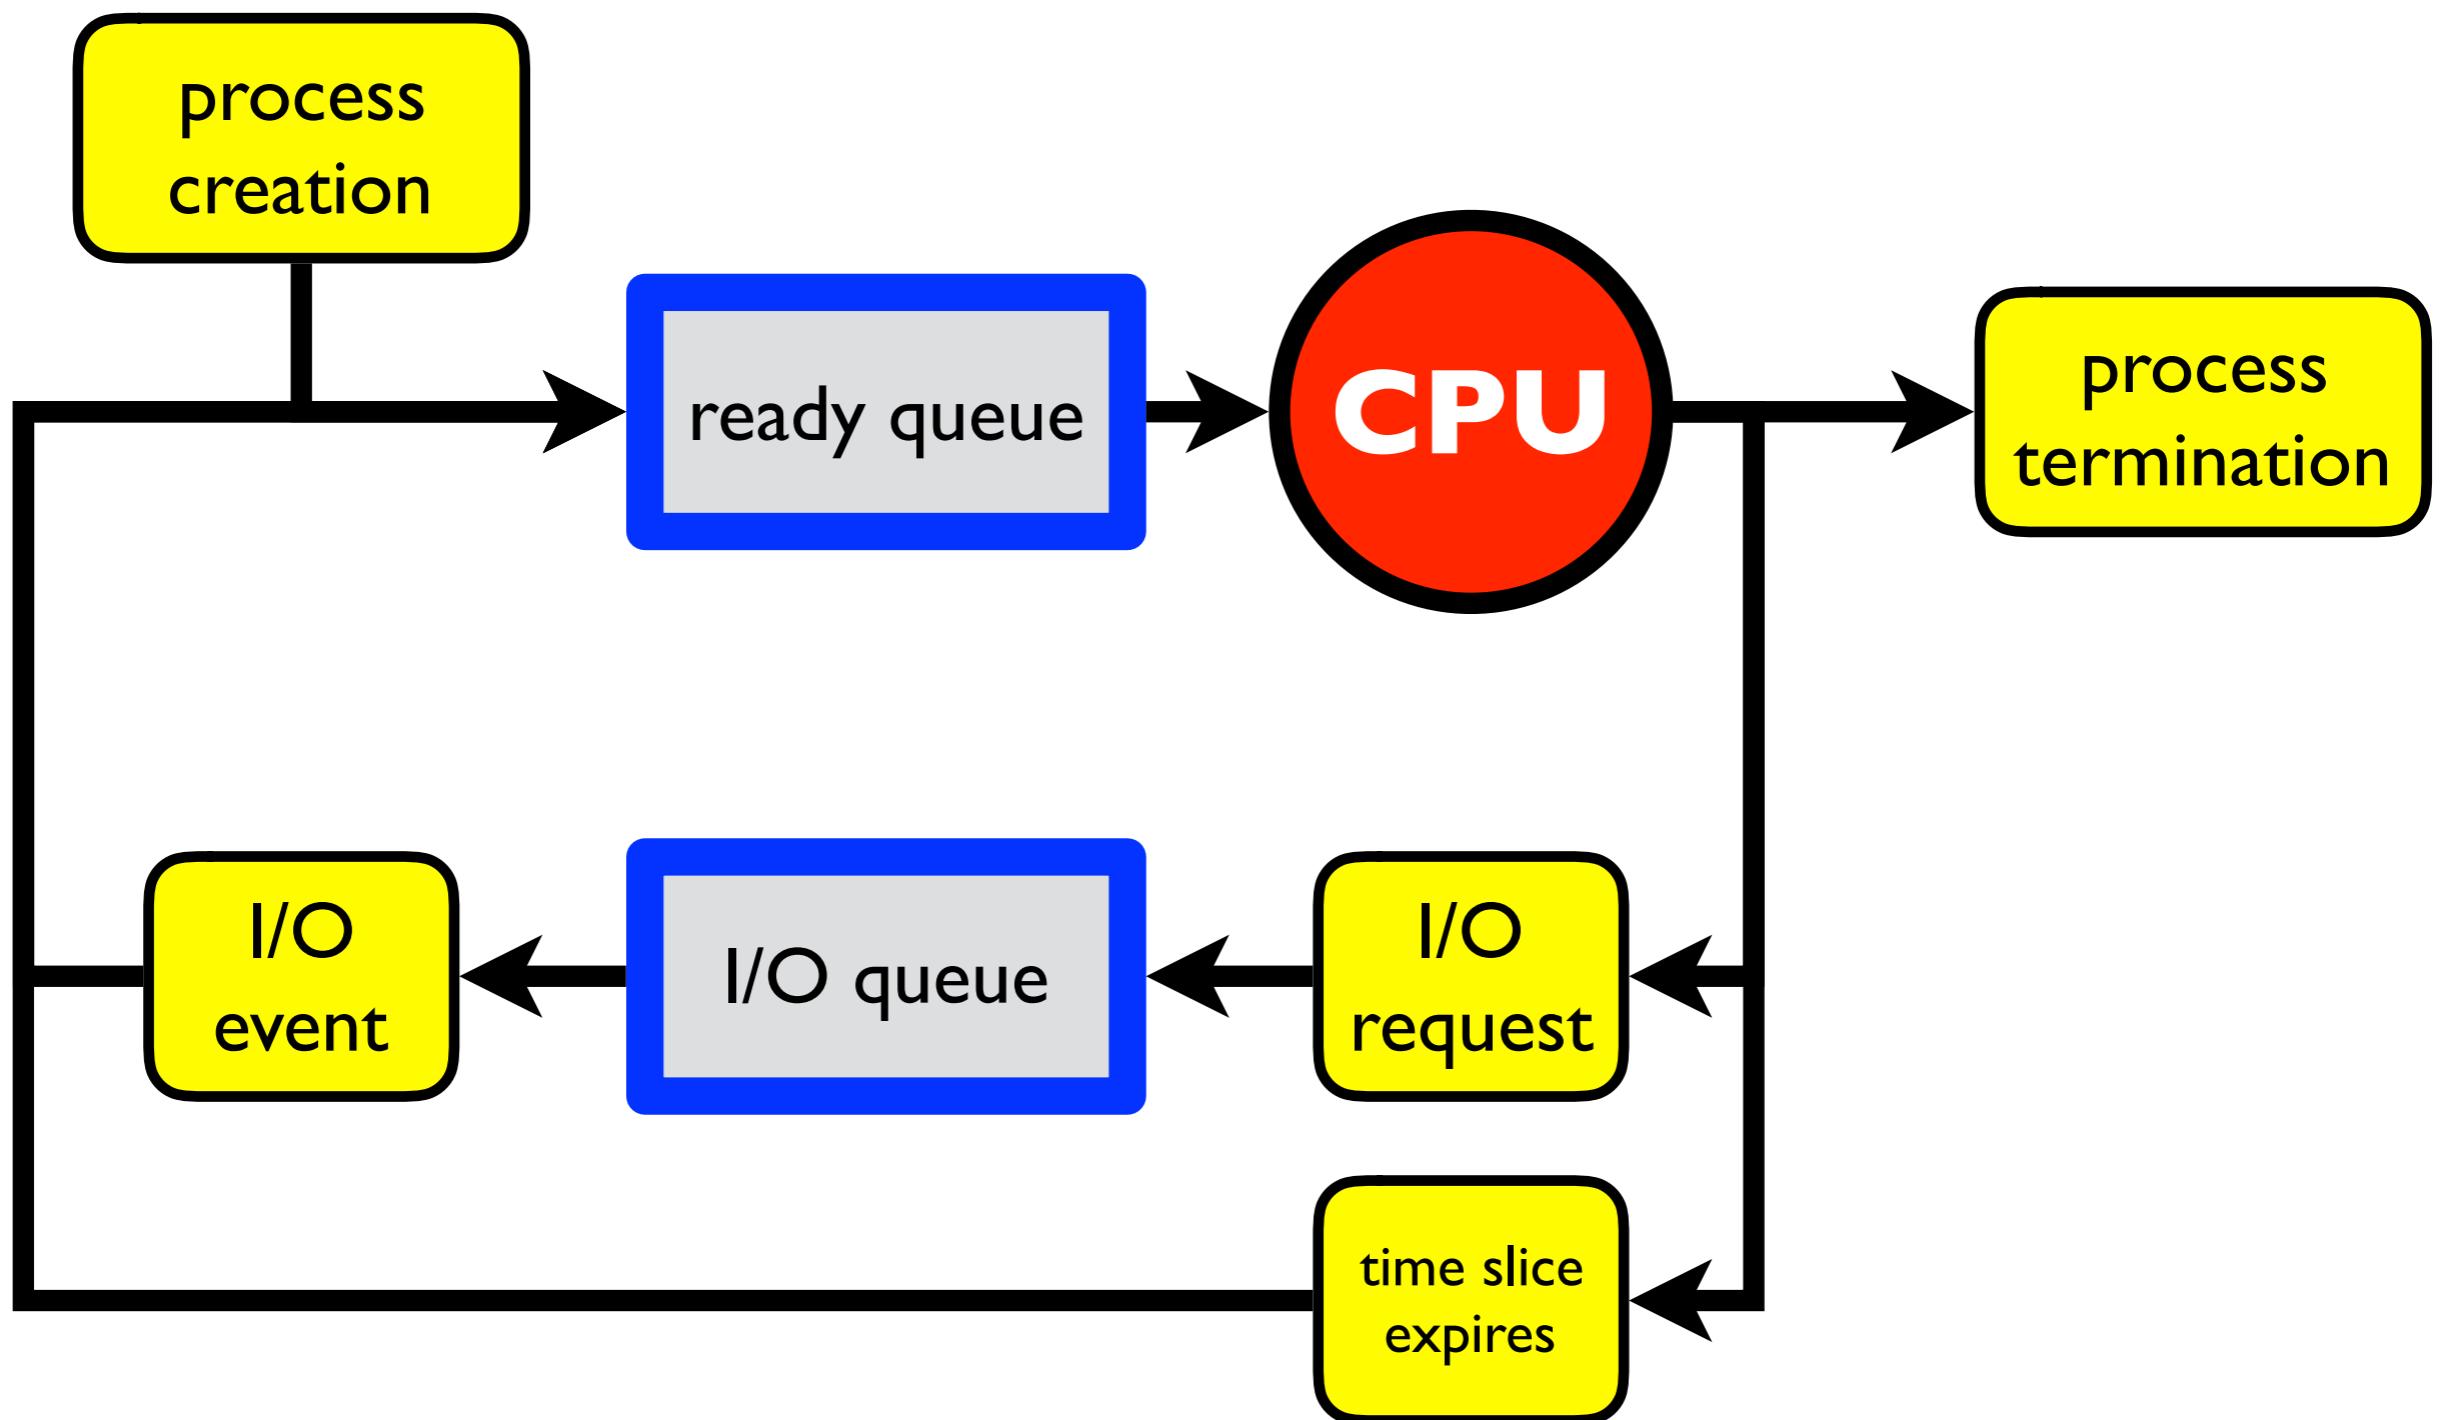
Time for context switch dependent on **hardware support**.



Process queues

Process queues

We have already seen two examples of process queues: the ready queue and the I/O queue.



Process queues


In general we have one ready queue and one wait queue for each I/O device.

Ready queue	Set of all processes residing in main memory, ready and waiting to execute.
Device wait queue	Set of processes, in main memory, waiting for an I/O device.

Process queues can be formed by linking PCBs together


Ready Queue

Device queues



**Process
creation**

How are process created?



System and Application Programs

Operating System

Controls the hardware and coordinates its use among the various application programs for the various users.

Bootstrap program



Kernel



Computer Hardware

System and Application Programs

Operating System

Controls the hardware and coordinates its use among the various application programs for the various users.

Bootstrap program

Kept on chip (ROM or EEPROM), aka **firmware**.

Small program executed on power up or reboot.

Initializes all aspects of the system, from CPU register to device controllers to memory content.

Locates and **loads the kernel into memory** for execution.

Kernel

The part of the operating system that is running at all times.

On boot, starts executing the first process such as **init**.

Waits for some **event** to occur ...


Computer Hardware

Process creation in Unix/Linux



A single **parent** process executes.

The **parent** process uses the **fork()** system call to **create** a new **child** process.



A new process, the **child** process is created. The new **child** is a **copy** of its **parent**.

The **child** **may** use the **exec()** system call after a fork to replace the process's memory space with a new program executable.

The **child** uses the **exit()** system call to terminate.

The single **parent** continue to execute.

fork()

In Unix-like operating system, fork is an operation whereby a process **creates a copy** of itself.



- ★ The process calling fork() is called the **parent**.
- ★ The new process is called the **child** of the parent.
- ★ Fork is usually a **system call**, implemented in the kernel.
- ★ Fork is the primary (and historically, only) method of process creation on Unix-like operating systems.

exec()

In Unix-like operating systems, the **exec family of system calls** runs an executable file in the context of an already existing process, **replacing the previous executable**.

- ★ As a new process is not created, the process identifier (PID) does not change, but the machine code, data, heap, and stack of the process are replaced by those of the new program.
- ★ A file descriptor open when an exec call is made remain open in the new process image, unless was fcntled with FD_CLOEXEC. This aspect is used to specify the standard streams (stdin, stdout and stderr) of the new program.
- ★ As a result of exec, all data in the old program that were not passed to the new program, or otherwise saved, become lost.

wait()

A **parent** process can use the `wait()` system call to suspend its execution until one of its children terminates and get the exit status of the terminated child.

```
pid_t wait(int *status);
```

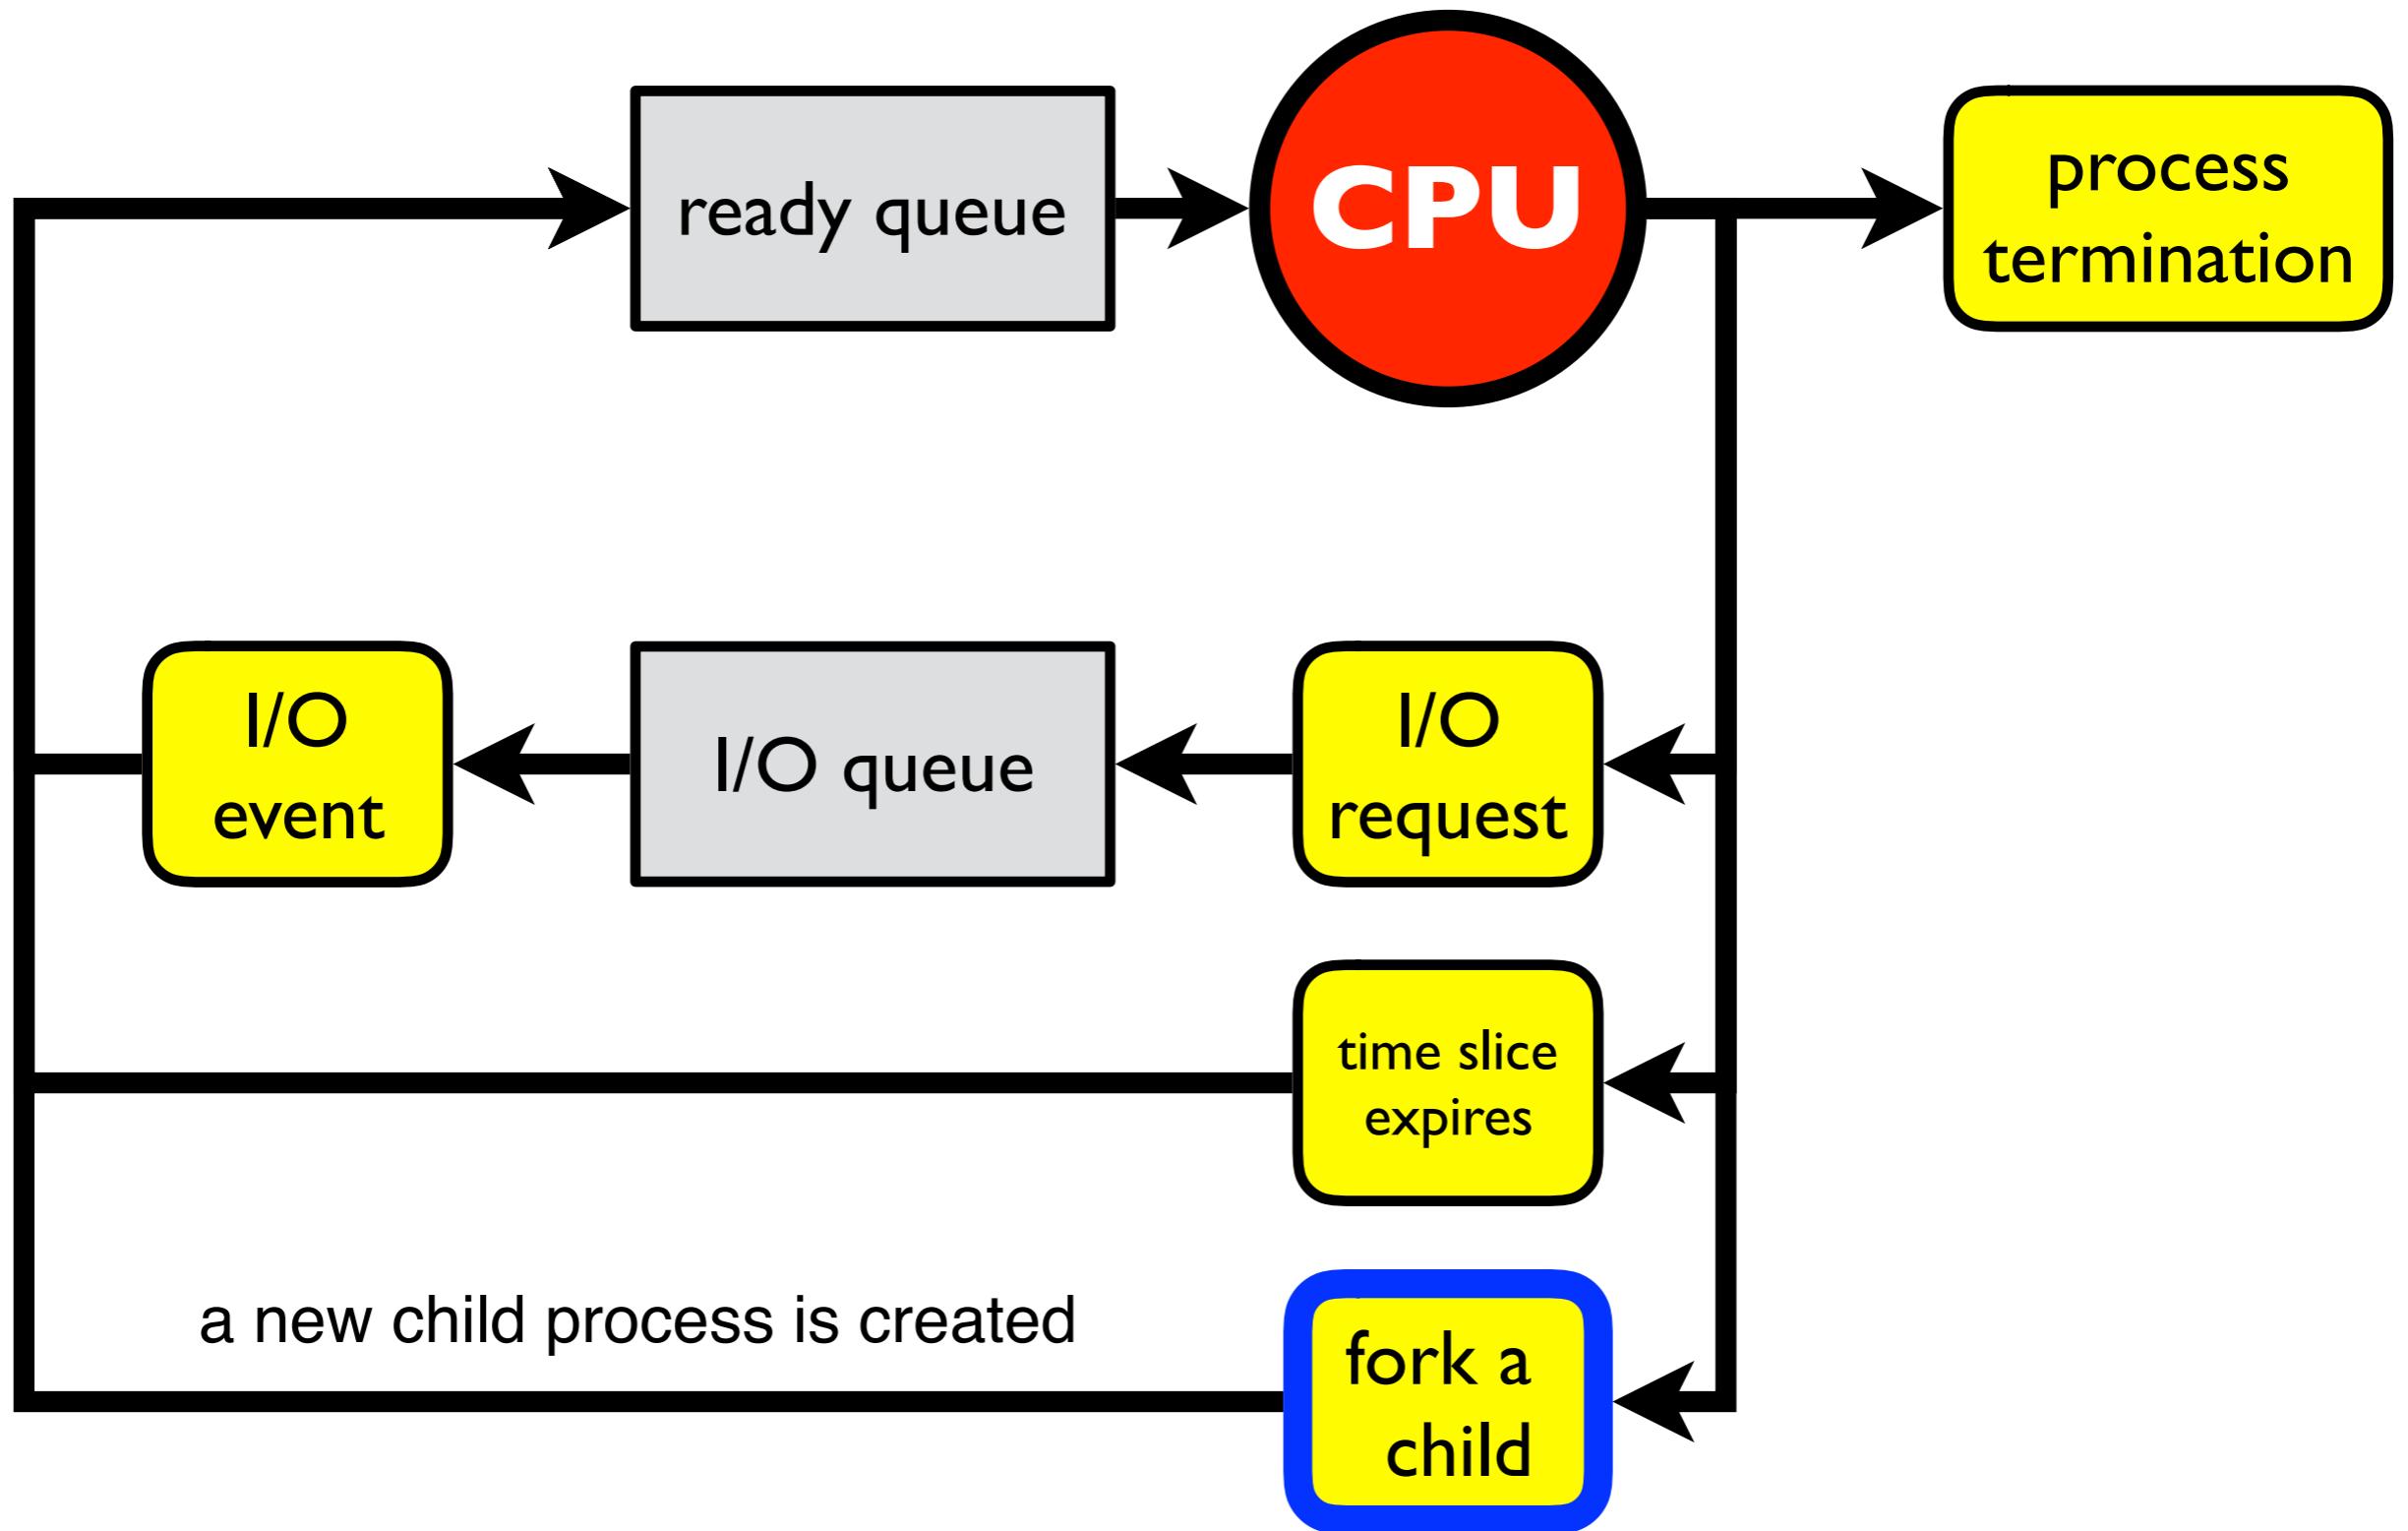
- ★ If there is **no child** process running when the call to `wait()` is made, then this `wait()` has **no effect** at all.
- ★ Otherwise, executing `wait()` **suspends the caller until one of its children terminates**.
- ★ **Returns the PID** of the terminated **child** process.
- ★ If `status` is not a null pointer, the **exit status of the terminated child** will be stored at the location pointed to by `status`.

exit()


When a process terminates it executes an exit() system call, either directly in its own code, or indirectly via library code.

```
void exit(int status);
```

- ★ The exit() call has **no return value** as the process that calls it terminates and so couldn't receive a value anyway.
- ★ The exit() call **resumes** the execution of a **waiting parent** process.
- ★ The exit() call also **communicates** the **status** parameter value **to the parent** process who can get the status using the wait() system call.



Create a new process using fork()



User Space

Kernel

Parent Process

Child Process

TEXT (instructions)
DATA

Parent calls fork()

Resources (open files, etc)

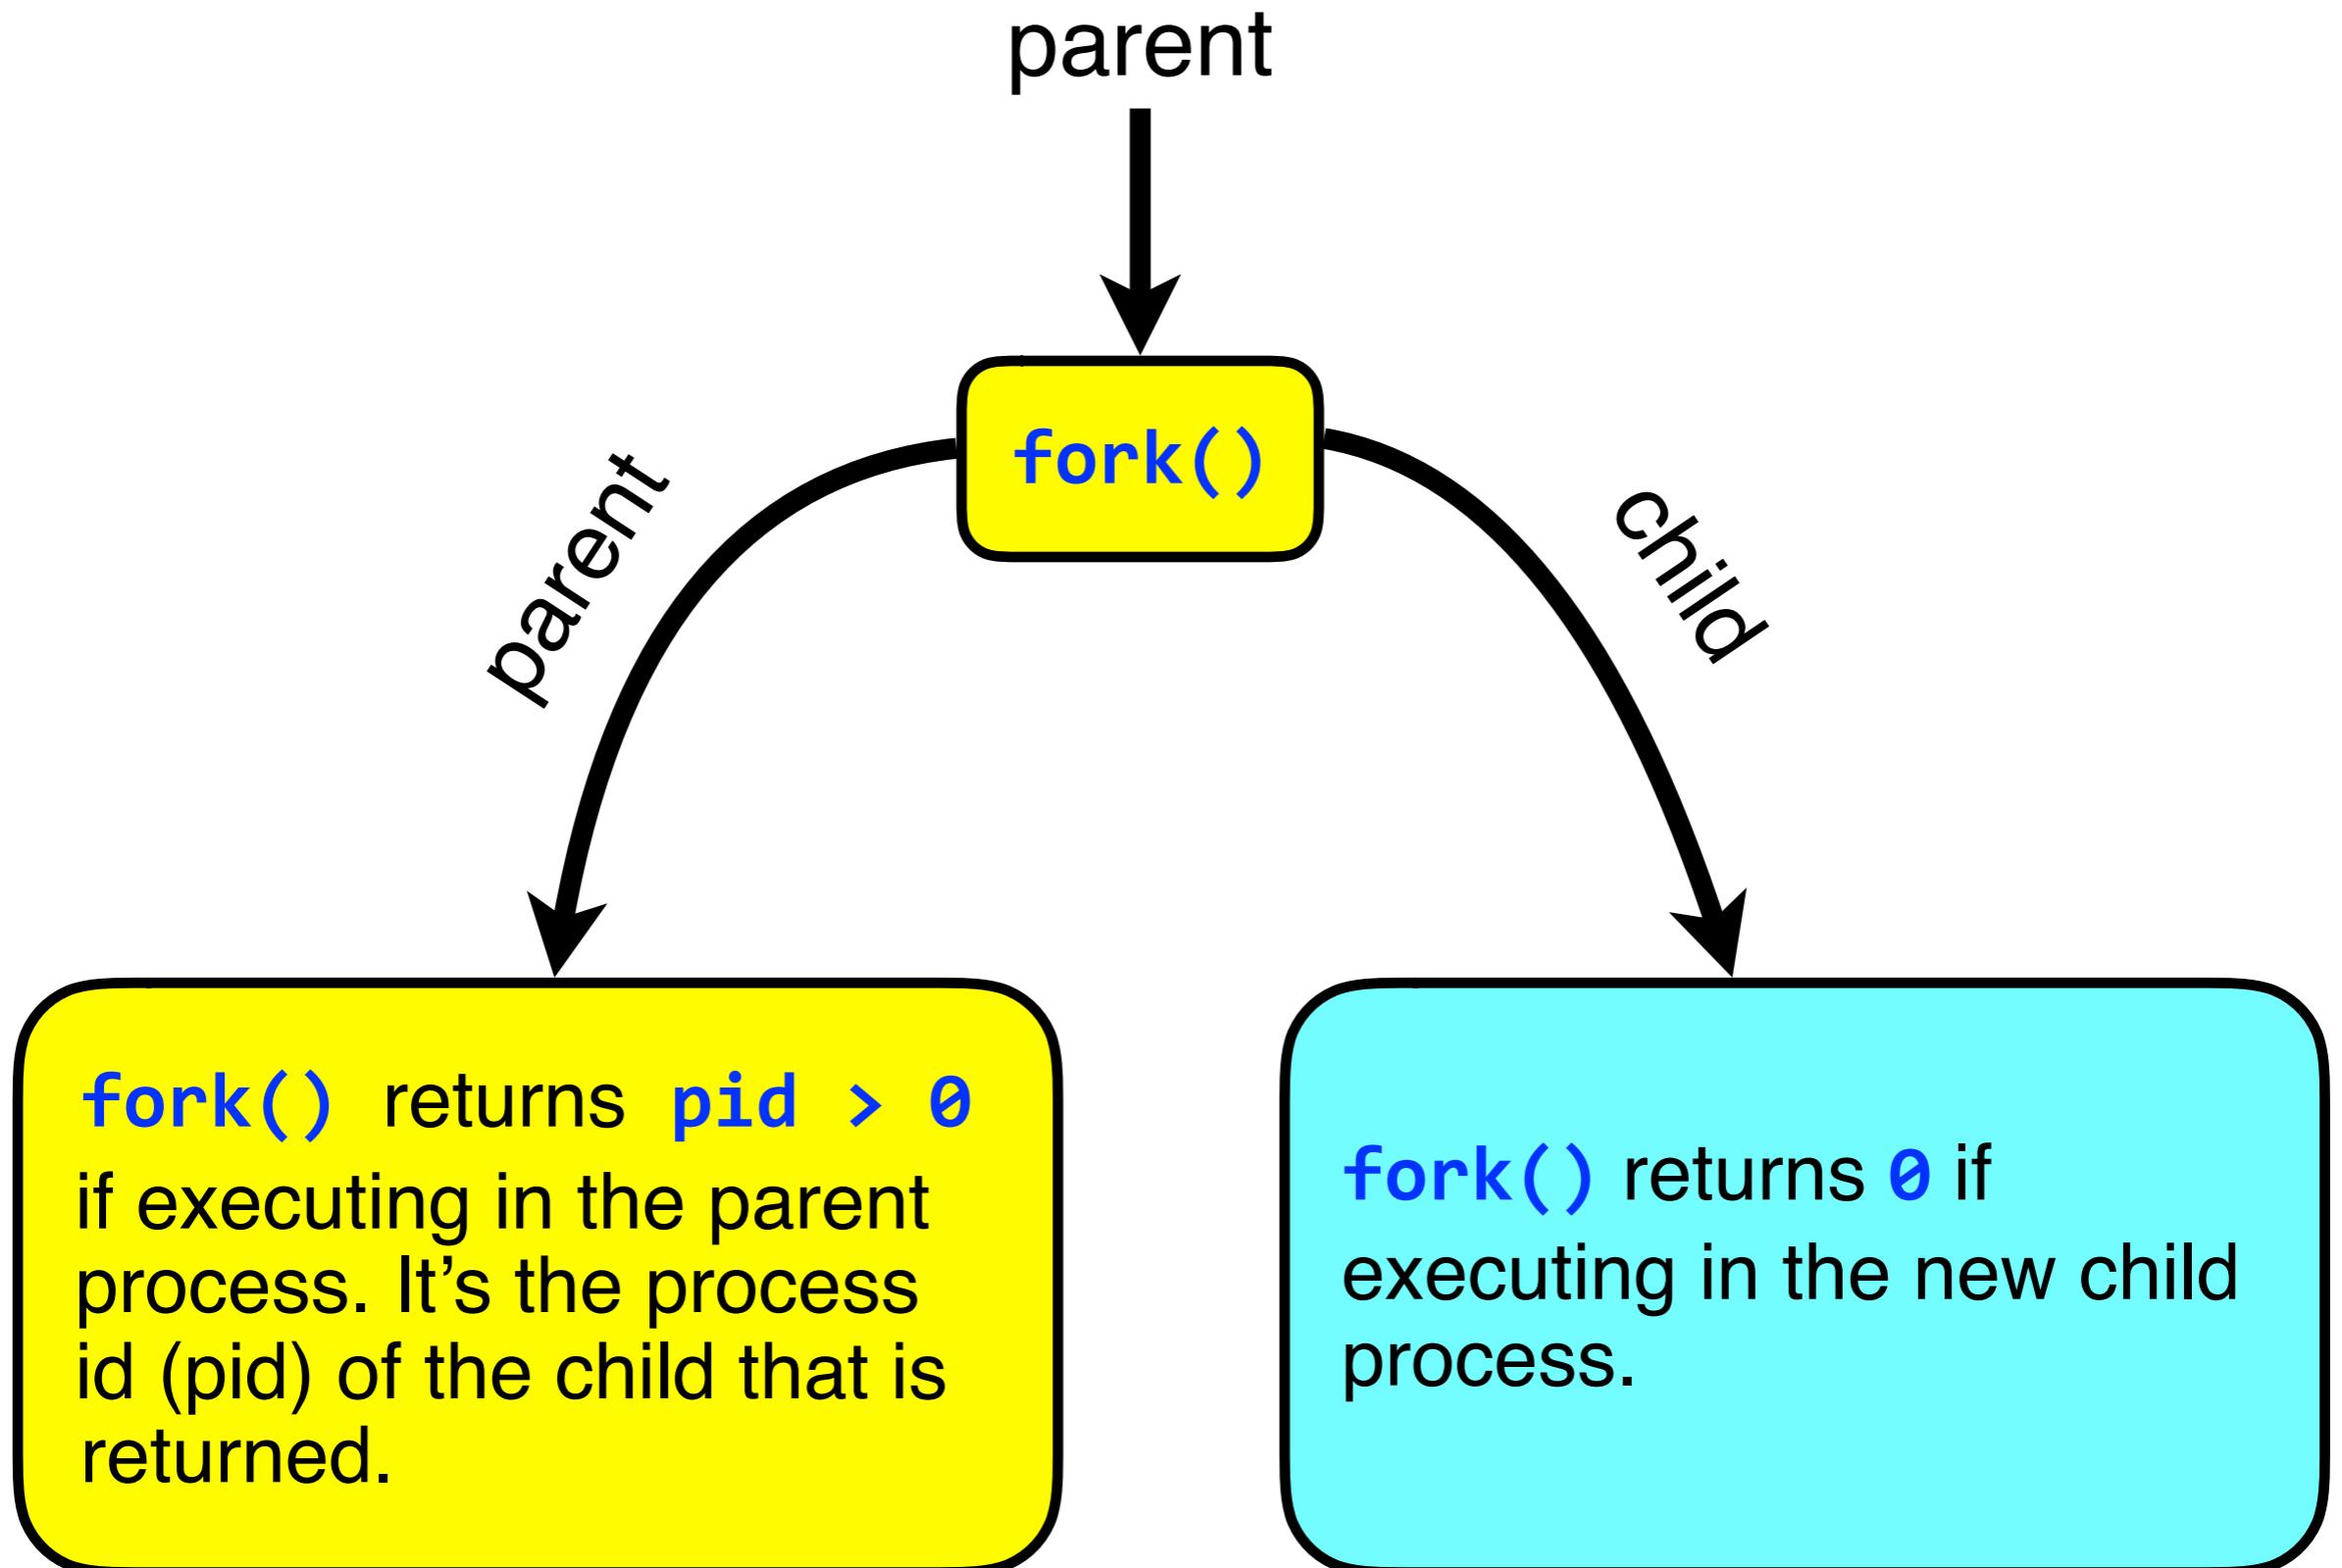
TEXT (instructions)
DATA

The child process is a
copy of the parent

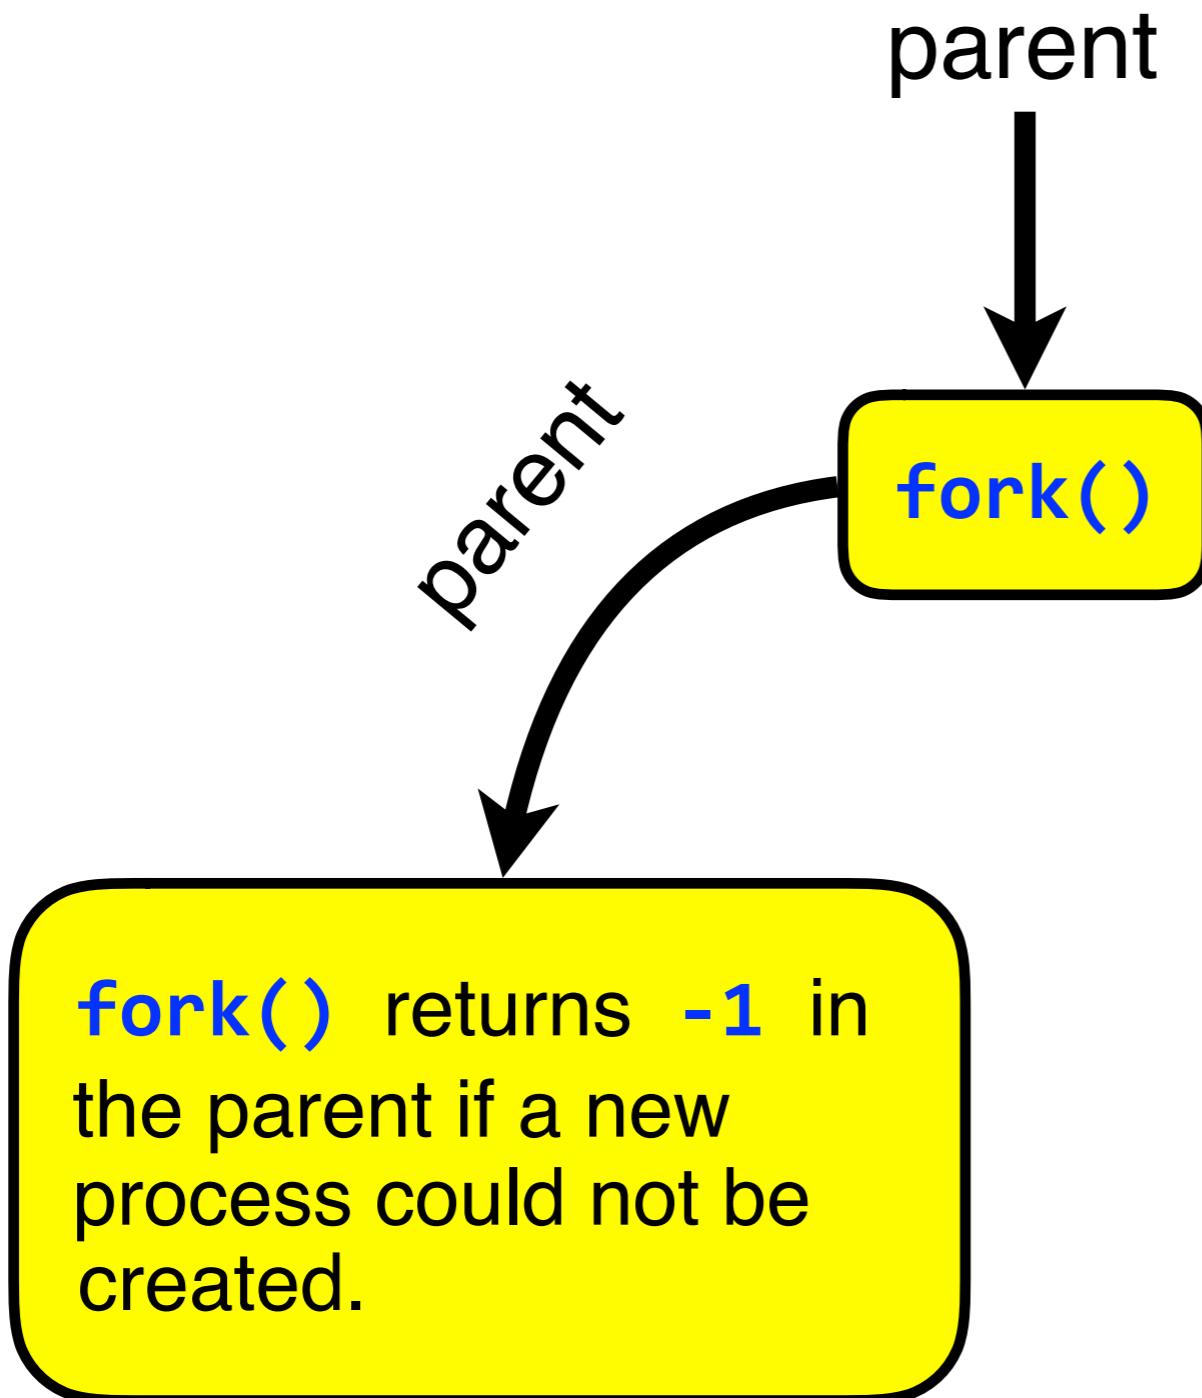
Resources (open files, etc)

After `fork()`, parent and child will execute the same program (TEXT). How can they be distinguished from each other?

On success, `fork()` return twice!



On failure, `fork()` returns -1 in the parent



Parent Process

STACK

pid = 766611

TEXT

pid = fork();

```
switch (pid) {  
case -1:  
    // ERROR  
    exit(EXIT_FAILURE);  
case 0:  
    // CHILD  
    exit(EXIT_SUCCESS);  
default:  
    // PARENT  
    exit(EXIT_SUCCESS);  
}
```

Child Process

STACK

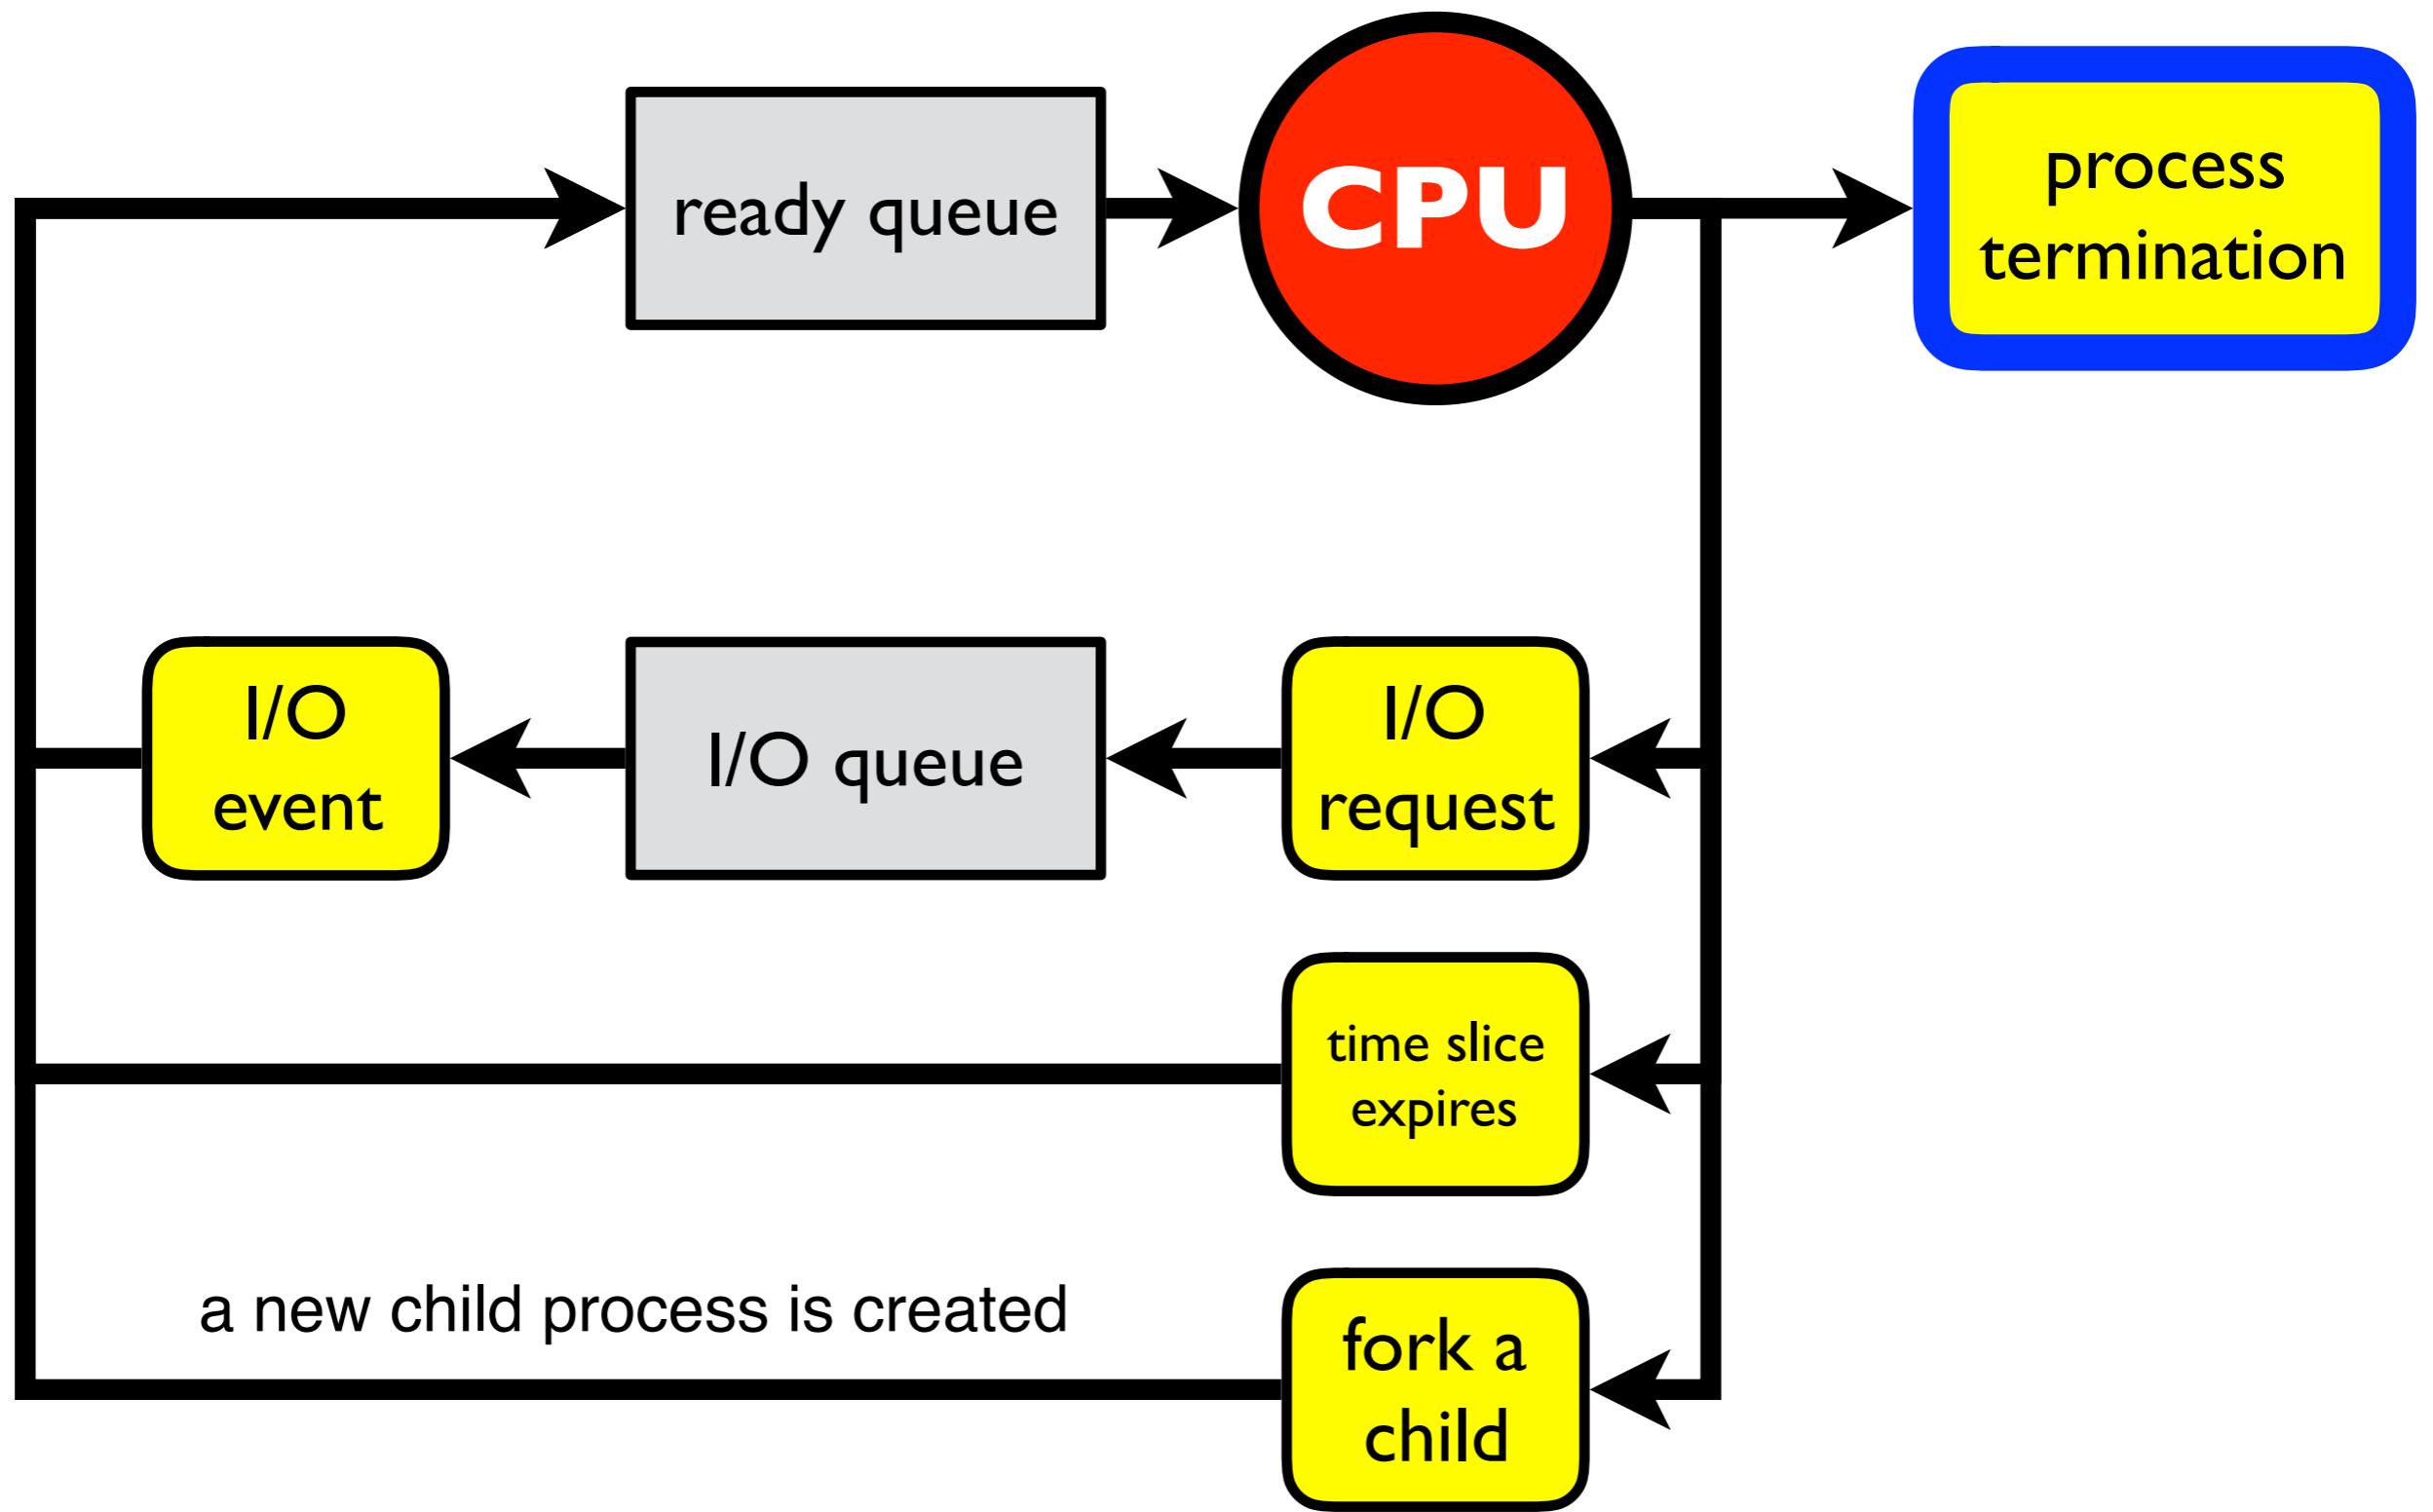
pid = 0

TEXT

pid = fork();


```
switch (pid) {  
case -1:  
    // ERROR  
    exit(EXIT_FAILURE);  
case 0:  
    // CHILD  
    exit(EXIT_SUCCESS);  
default:  
    // PARENT  
    exit(EXIT_SUCCESS);  
}
```

Process termination



exit() and wait()

When a process terminates it executes an **exit()** system call, either directly, or indirectly via library code. This leaves an **exit status** value (typically an integer) **in the PCB** of the terminated process for the parent process to read later.



The parent may use the **wait()** system call to wait for the child to terminate and read the exit status value of the terminated child process.



Zombie processes



Zombie

A terminated process is said to be a zombie or defunct until the parent does **wait()** on the child.

- ★ When a process terminates all of the memory and resources associated with it are deallocated so they can be used by other processes.
- ★ However, the **exit status is maintained in the PCB until the parent picks up the exit status using wait() and deletes the PCB.**
- ★ A child process always first becomes a zombie before being removed from the resource table.
- ★ In most cases, under normal system operation zombies are immediately waited on by their parent.
- ★ Processes that stay zombies for a long time are generally an error and cause a resource leak.

Orphan processes

Orphan

An orphan process is a process whose parent process has terminated, though it remains running itself.


- ★ Any orphaned process will be immediately **adopted** by the special **init** system process.
- ★ Even though technically the process has the **init** process as its parent, it is still called an orphan process since the process that originally created it no longer exists.
- ★ A zombie process is not the same as an orphan process. An orphan process is a process that is still executing, but whose parent has died. An orphan process do not become zombie processes; instead, they are adopted by init (process ID 1), which automatically waits on all its child process.

Inter process communication (IPC)


How can we make processes communicate (sharing information)?

Two methods: **message passing** and **shared memory**.

Message Passing



Shared Memory



Communication take place by means of messages exchanged between the cooperating processes.

A region of memory that is shared by cooperating process is established. Process can then exchange information by reading and writing data to the shared region.

Message Passing

Messages can be exchanged between processes either directly or indirectly using a common mailbox.

The recipient process usually must give its permission for communication to take place with an accept connection system call.

Message-passing is **useful** for exchanging **smaller amounts of data**, because no conflicts need be avoided.

Easier to implement compared to the shared memory approach.

Shared Memory

Processes can communicate by reading and writing to shared memory.

Normally, the OS tries to prevent one process from accessing another process's memory. Shared-memory requires that two or more processes agree to remove this restriction.

Shared memory **allows maximum speed** and convenience of communication, since it can be done at memory transfer speeds.

Problems: protection and synchronization between the processes sharing memory.

signals



Signals

(1)

Signals are a limited form of inter-process communication used in Unix, Unix-like, and other POSIX-compliant operating systems

- ★ A signal is an **notification** sent to a process in order to notify it of an event that occurred.
- ★ When a signal is sent, the operating system **interrupts** the target process's **normal flow of execution** to deliver the signal.
- ★ If the process has previously **registered** a **signal handler**, that routine is executed. Otherwise, the **default signal handler** is executed.

Signals

(2)

Signals can be synchronous or asynchronous.

Synchronous signals

- ★ Are delivered to the same process that performed the operation that caused the signal.
- ★ Examples of synchronous signals: **illegal memory** access and **division by zero**.

Asynchronous signals

- ★ Are generated by an event external to a running process.
- ★ Typically, an asynchronous signal is sent to another process.
- ★ Examples of asynchronous signals: **terminating** a process (ctrl-c), **timer expire**.

Signals

(3)

Typing certain key combinations at the controlling terminal of a running process causes the system to send it certain signals

- ★ **Ctrl-C** sends an INT signal (**SIGINT**); by default, this causes the process to **terminate**.
- ★ **Ctrl-Z** sends a TSTP signal (**SIGTSTP**); by default, this causes the process to **suspend** execution.

Signals


(4)

The kernel can generate a signal to notify the process of an event.

- ★ **Exceptions such as division by zero or a segmentation violation will generate signals.**
- ★ Division by zero will generate a **SIGFPE** signal.
- ★ Segmentation violation will generate a **SIGSEGV** signal.
- ★ If not explicitly caught, SIGFPE and SIGSEGV will cause a core dump and a program exit.

kill()

In Unix-like operating systems, the kill system call is used to send signals to processes.



There are many different signals that can be sent, although the signals in which users are generally most interested are **SIGTERM** and **SIGKILL**.

- ★ The **SIGKILL** signal is sent to a process to cause it to **terminate immediately** (**kill**). In contrast to SIGTERM and SIGINT, this signal cannot be caught or ignored, and the receiving process cannot perform any clean-up upon receiving this signal.

- ★ The **SIGTERM** signal is sent to a process to request its **termination**. Unlike the SIGKILL signal, it can be caught and interpreted or ignored by the process. This allows the process to perform nice termination releasing resources and saving state if appropriate. SIGINT is nearly identical to SIGTERM.

Pipes



Pipes

An (anonymous) pipe is a **simplex FIFO communication channel** that may be used for one-way interprocess communication (IPC).



- ★ An implementation is often integrated into the operating system's file IO subsystem.
- ★ Pipes were one of the first IPC mechanisms in early UNIX systems.
- ★ Typically a parent program opens anonymous pipes, and creates a new process that inherits the other ends of the pipes, or creates several new processes and arranges them in a pipeline.

Producer and consumer


Anonymous pipes allow two processes to communicate in standard producer consumer fashion.

- ★ The **producer** writes to one end of the pipe (the write end).
- ★ The **consumer** reads from one end of the pipe (the read end)
- ★ As a result, ordinary pipes are **unidirectional**, allowing only **one-way communication**.
- ★ If two-way communication is required, two pipes must be used, with each pipe sending data in a different direction.

pipe()

In Unix-like operating systems pipes are created using the **pipe()** system call.

- ★ As a result of the pipe() system call, a pipe object is created.
- ★ A pipe is a FIFO buffer that can be used for inter process communication (IPC).
- ★ The pipe() system call **returns a pair of file descriptors** referring to the **read** and **write** ends of the pipe.



Using pipe()

User Space

Parent Process

```
int pfd[2];
pipe(pfd);
// pfd[0] = 3
// pfd[1] = 4
```

Descriptors	
0	stdin
1	stdout
2	stderr
3	pipe read
4	pipe write



(1)

As a result of the `pipe()` system call, a pipe object is created. A pipe is a FIFO buffer that can be used for inter process communication (IPC).

(2)

As a result of the `pipe()` system call, the elements of the array `pfd` are assigned descriptor numbers to the created pipe.

Now a single process can write data to, and later read data from the pipe ...

Pipe and Fork



Using pipe() together with fork()

User Space
Kernel Space

Parent Process

```
int pfd[2];
pipe(pfd);
// pfd[0] = 3
// pfd[1] = 4
fork();
```

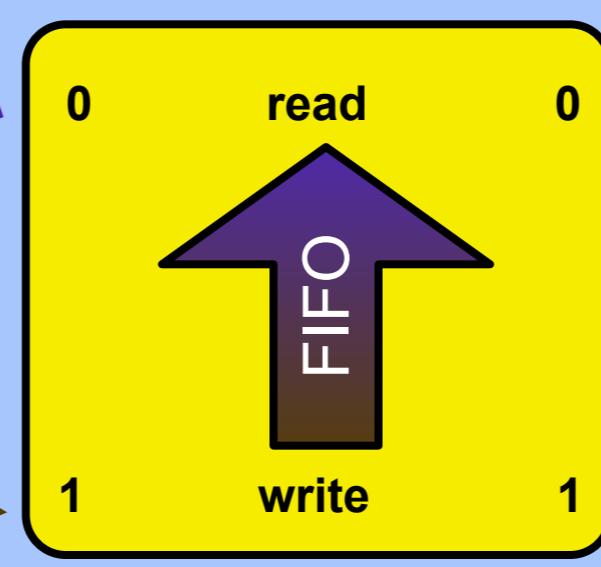
Descriptors	
0	stdin
1	stdout
2	stderr
3	pipe read
4	pipe write

fork()

Child Process

After `fork()`, the child process have the same set of descriptors including read and write descriptors to the pipe.

Descriptors	
0	stdin
1	stdout
2	stderr
3	pipe read
4	pipe write



Using pipe() together with fork()

User Space

Parent Process

```
int pfd[2];
pipe(pfd);
// pfd[0] = 3
// pfd[1] = 4
fork();
close(pfd[0]);
```

The parent can close the read descriptor to the pipe.


Child Process

The child can close the write descriptor to the pipe.

```
close(pfd[1]);
```

Descriptors

	Descriptors
0	stdin
1	stdout
2	stderr
3	pipe read
4	pipe write



Descriptors

	Descriptors
0	stdin
1	stdout
2	stderr
3	pipe read
4	pipe write

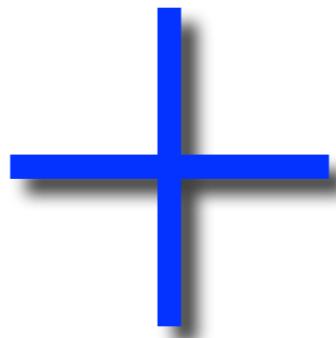
Kernel Space

Blocking

Reading and writing to a pipe may block a process.

Attempt	Conditions	Result
Read	Empty pipe, writer attached	Reader blocked
Write	Full pipe, reader attached	Writer blocked
Read	Empty pipe, no writer attached	EOF returned
Write	No reader	?

Pipes and signals



SIGPIPE

The SIGPIPE signal is used to notify a producer (writer) when there is no longer a consumer (reader) attached to a pipe.

Attempt	Conditions	Result
Read	Empty pipe, writer attached	Reader blocked
Write	Full pipe, reader attached	Writer blocked
Read	Empty pipe, no writer attached	EOF returned
Write	No readers attached	SIGPIPE

By default, SIGPIPE causes the process to terminate.
Don't forget to close unused pipe file descriptors using
the `close()` system call.

