

# Implementing the ls command

**Module 5 self study material**

---

**Operating systems 2019**

**1DT044, 1DT096 and 1DT003**



The Unix command interpreter (aka shell) does not understand the commands in any way.

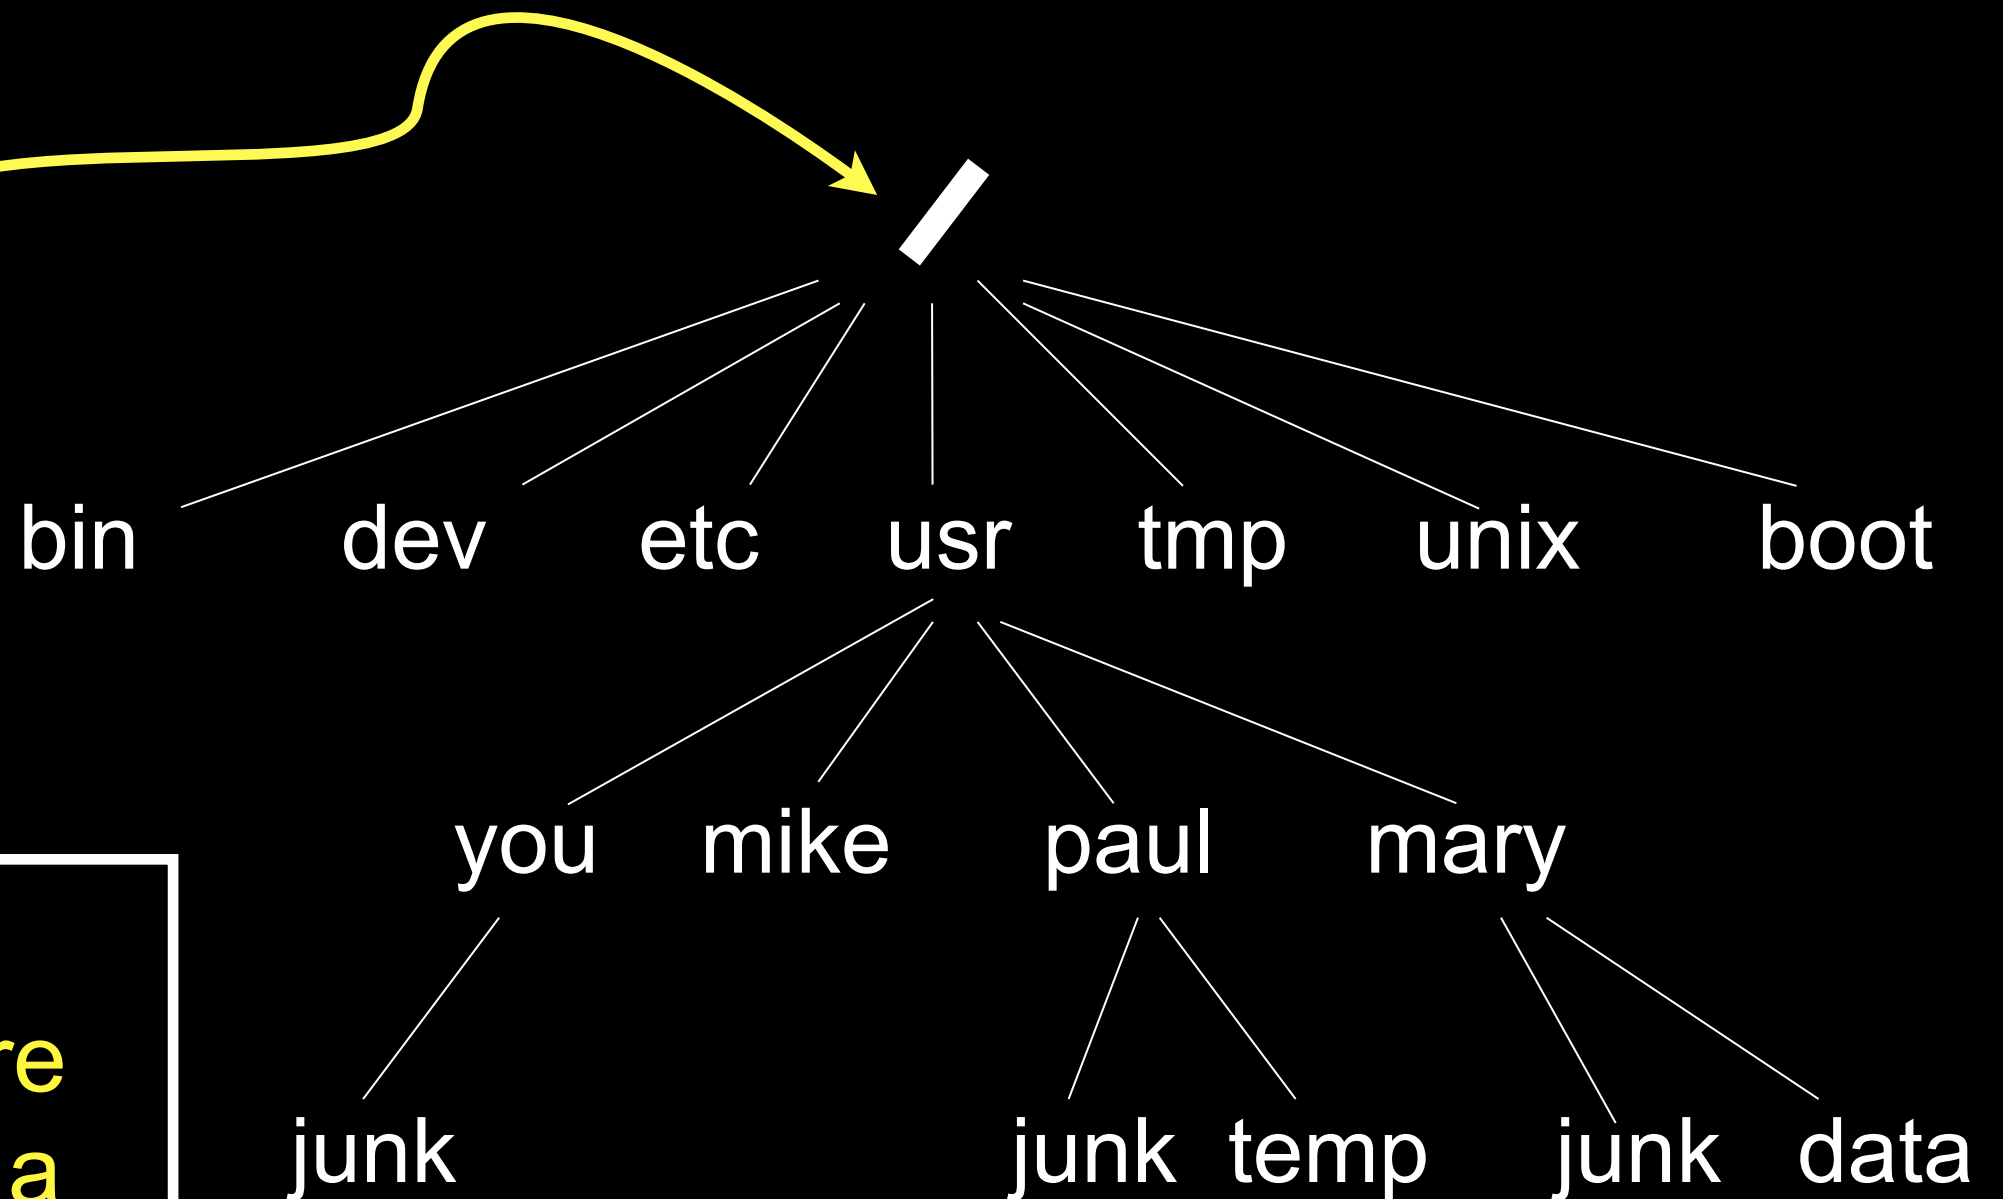
Commands are implemented through system programs.

The shell creates a new process using `fork()` and have the new process use `exec()` to run the identified system program executable.

We should be able to implement our own (simplified) version of the `ls` command.

# The Unix file system

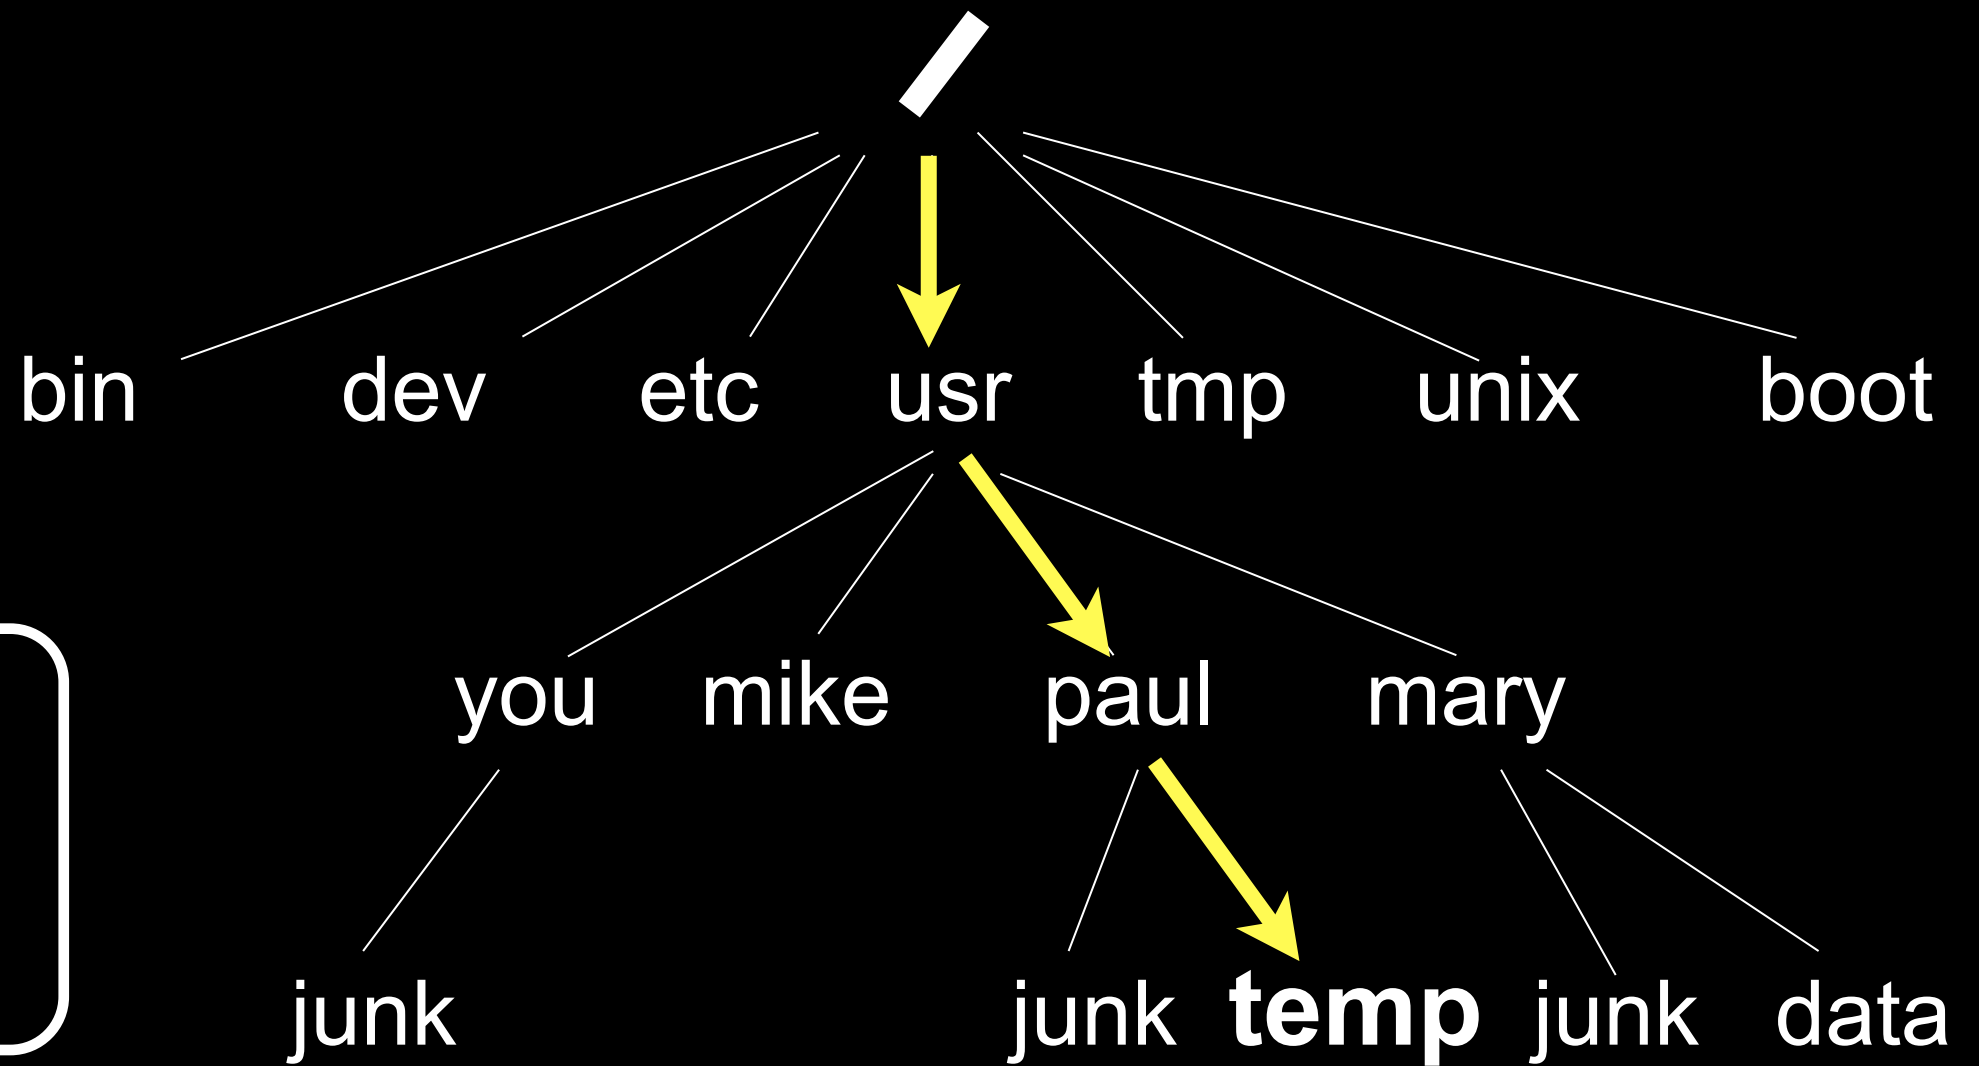
---



Files and directories are organized in a **tree structure** – starting at the **root**.


# The Unix file system

---



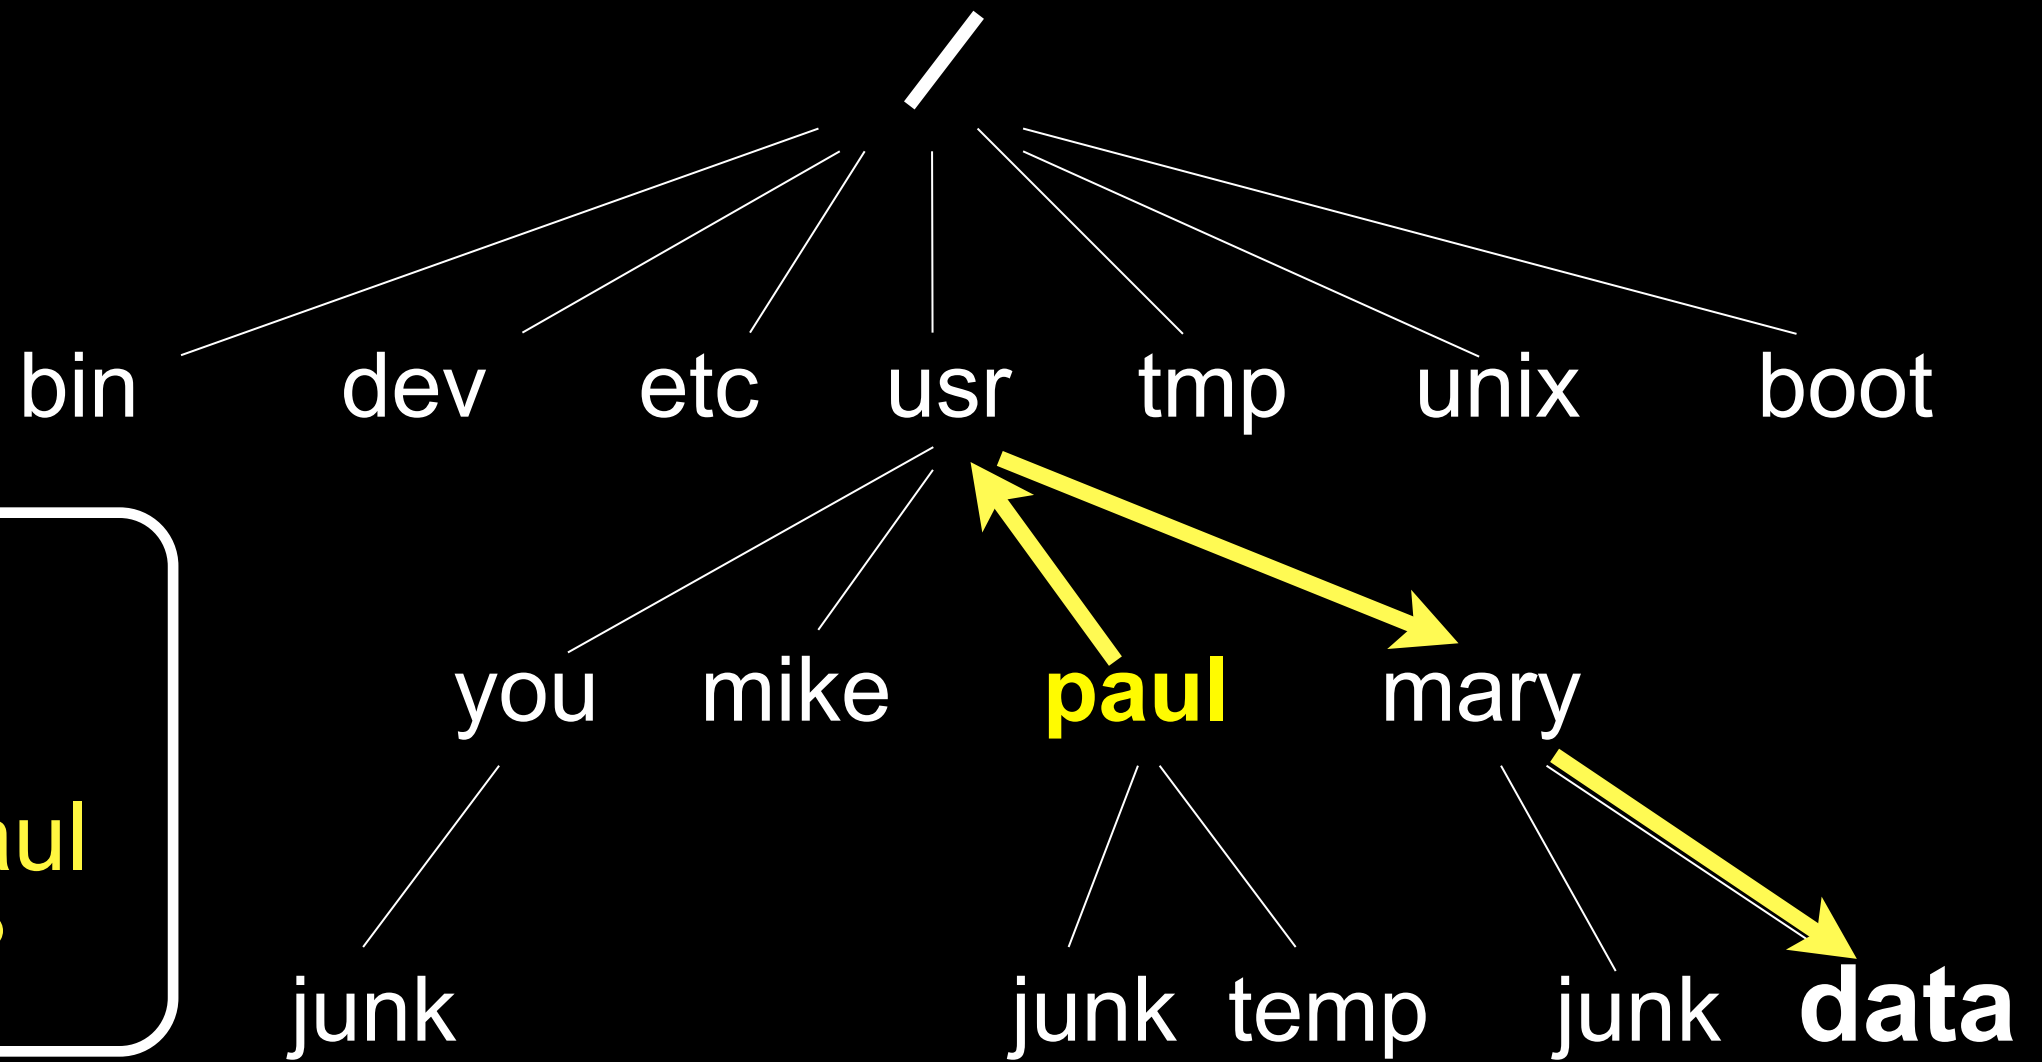
**Absolute  
path to file  
temp?**

**/usr/paul/temp**




# The Unix file system

---



Relative  
path from  
directory paul  
to file data?

../mary/data



# The **ls** command

In Unix-like systems you can use the ls command to get information about files in a directory.

The **owner** of the file

Highest **group** that the owner belongs to

```
$> ls -l
```

```
-rw----- 1 hans it readme.txt
```

```
-rwx-r--r-- 1 karl it script.sh
```

File mode - **permissions**

**Name** of File or Directory

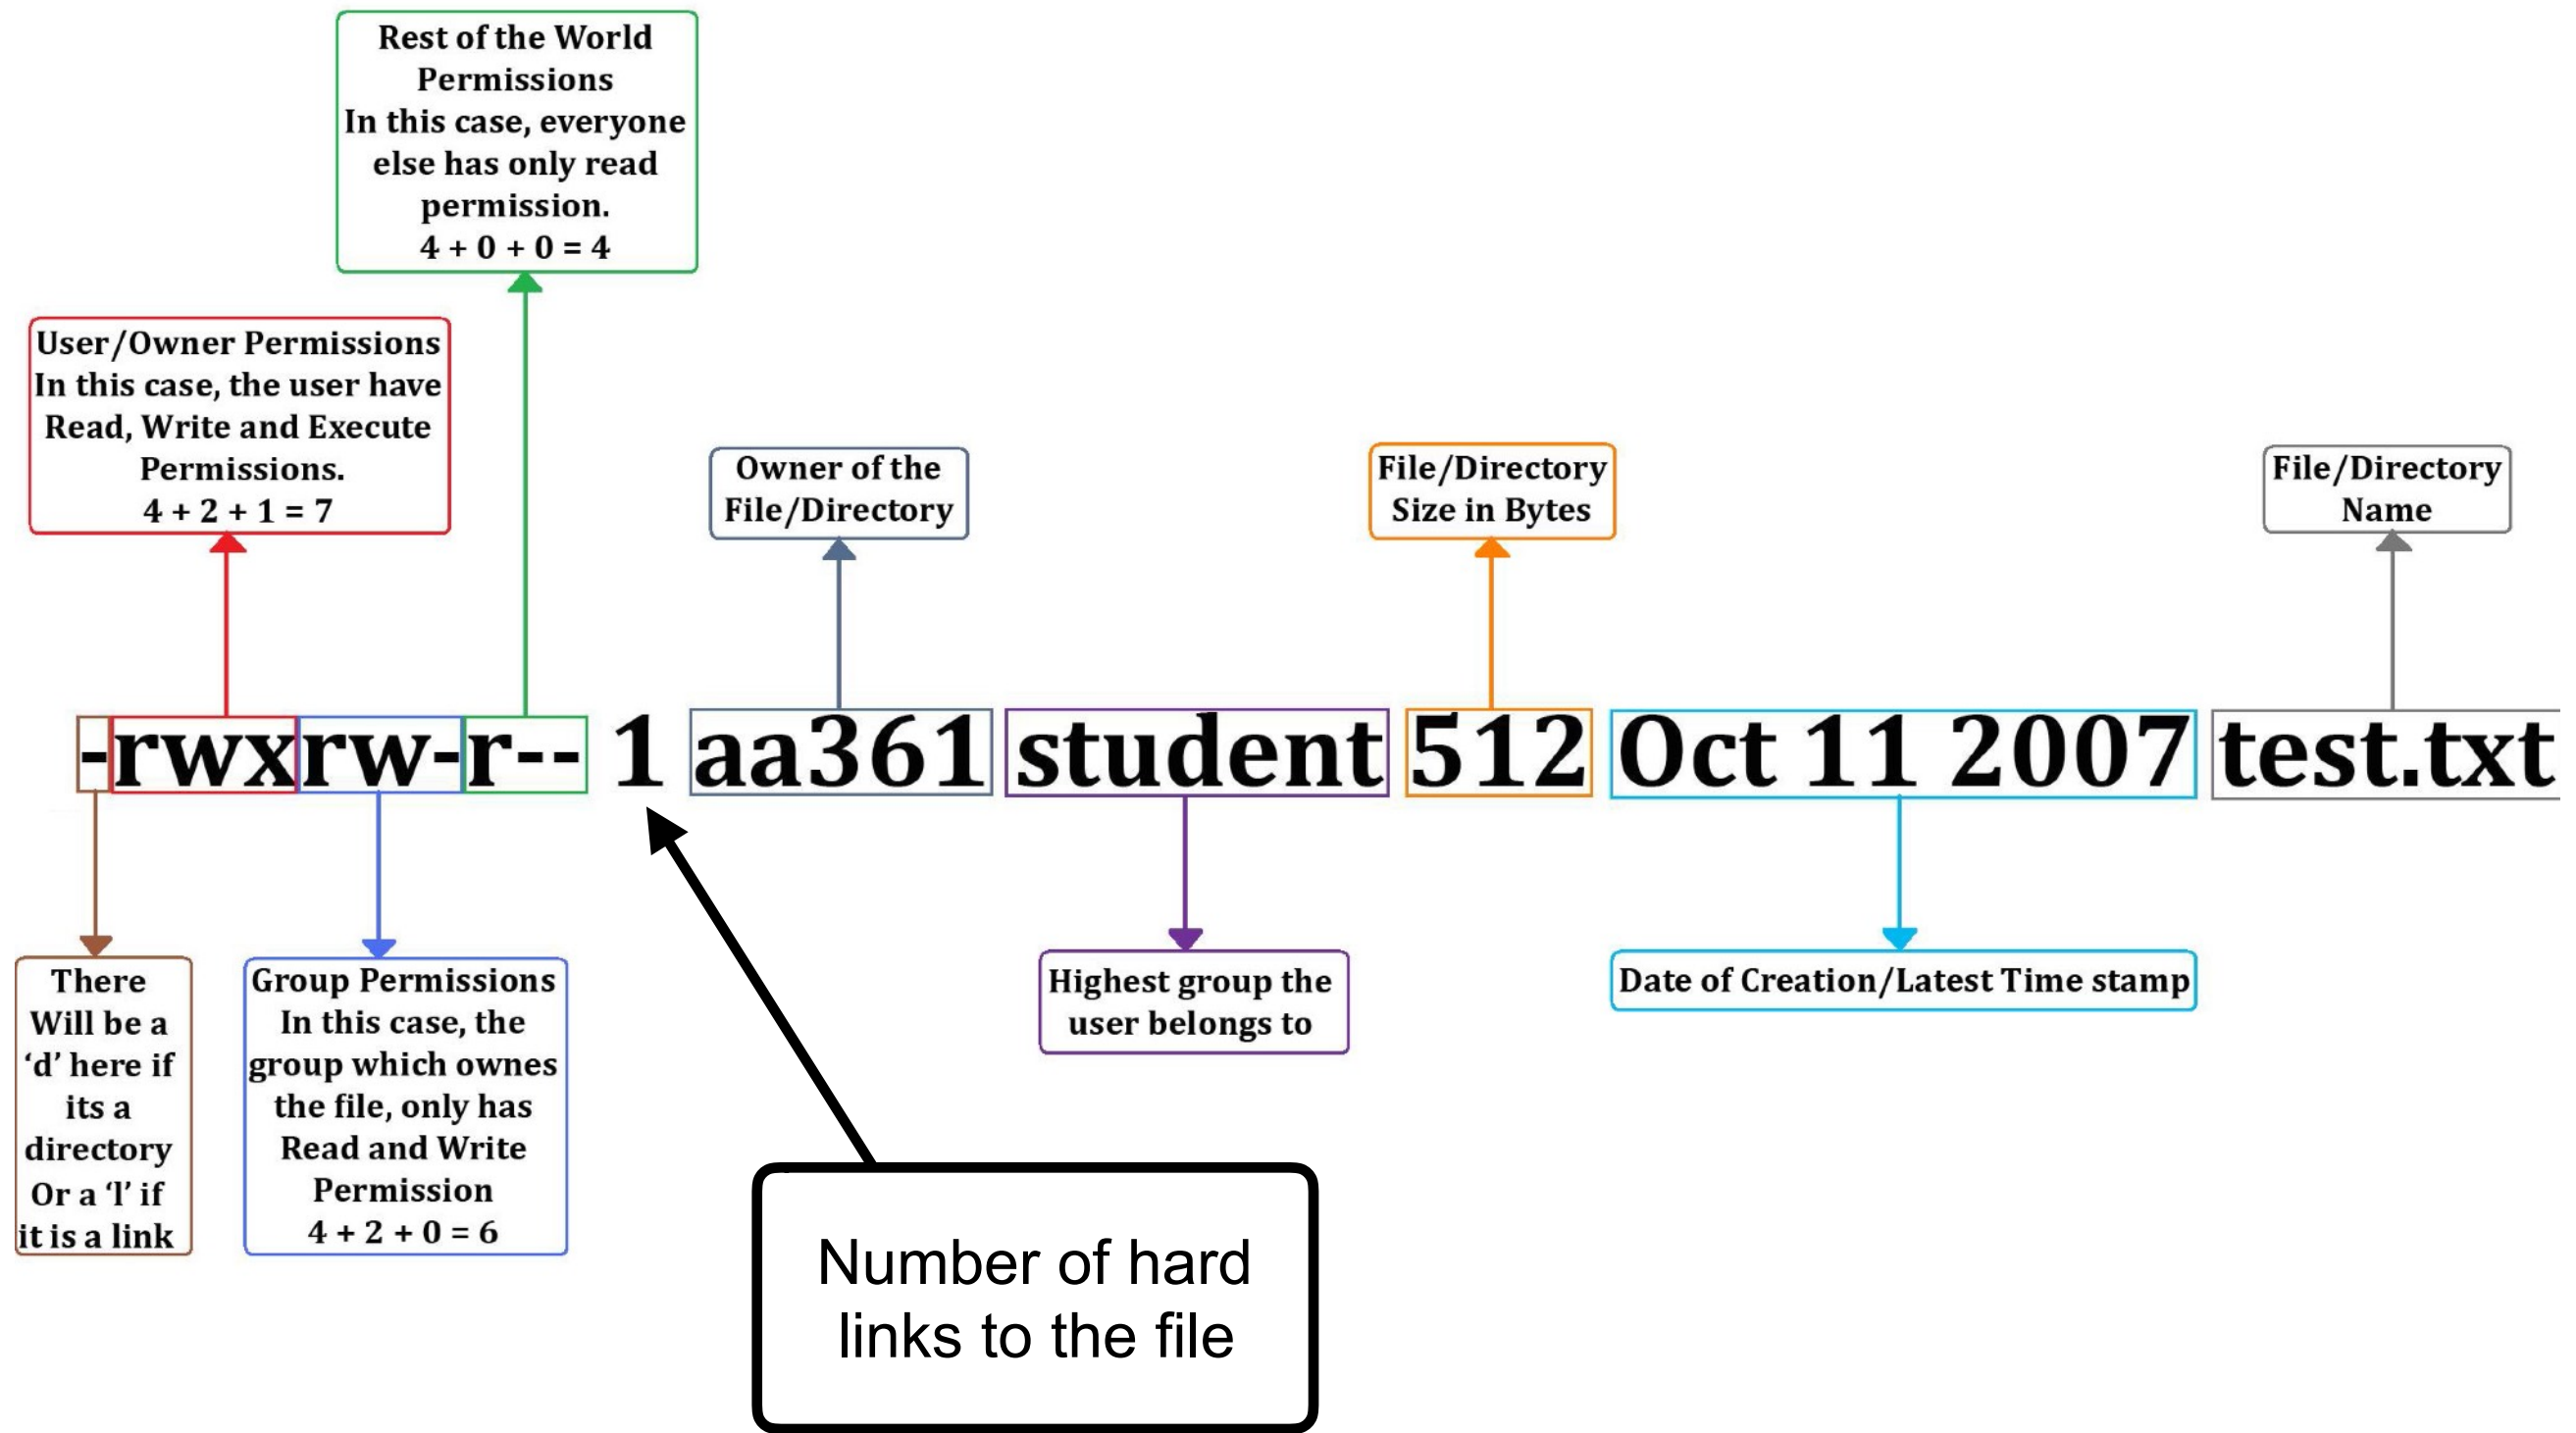
# Unix file permissions

```
$> ls -l
```

```
-rw----- 1 hans it readme.txt  
-rwx-r--r-- 1 karl it script.sh
```

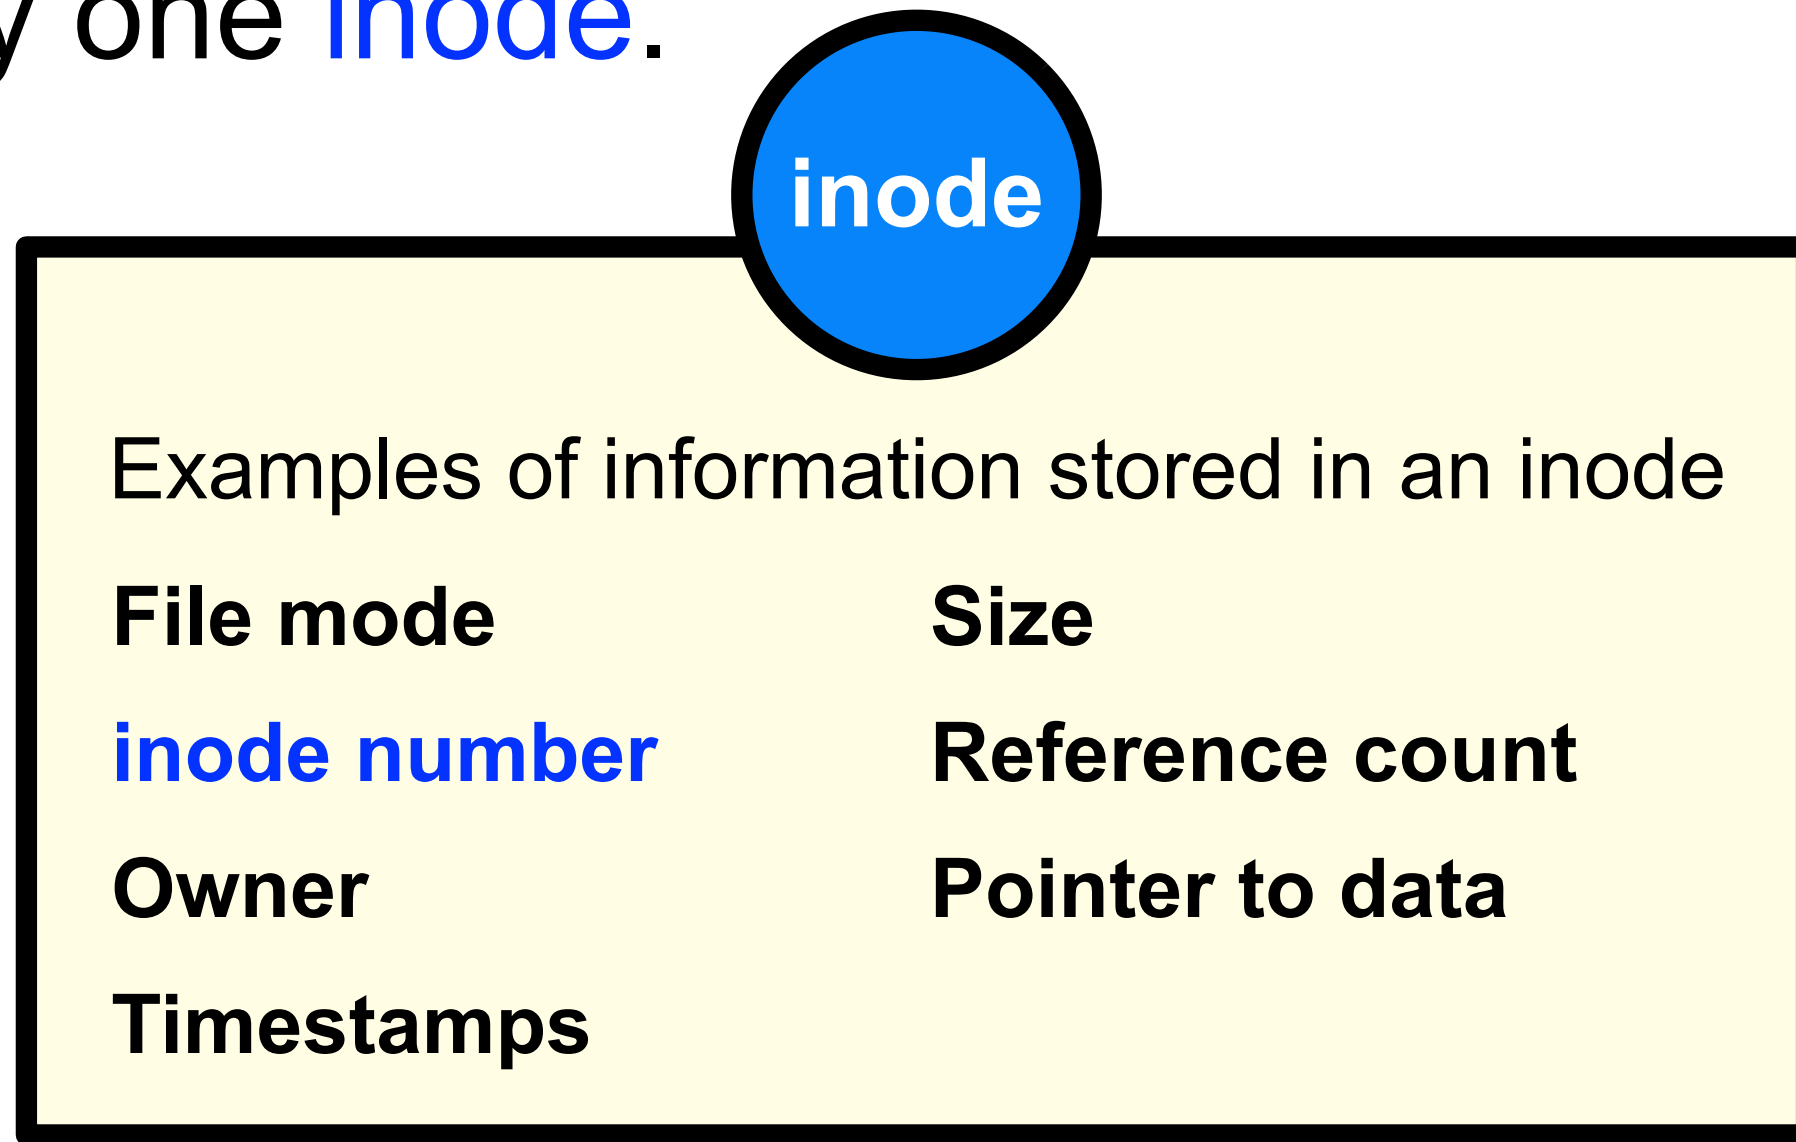
**r** = read      **w** = write      **x** = execute

**- r w x r - - r - -**  
          └──┬──┘   └──┬──┘   └──┬──┘  
          owner   group       all  
                                others



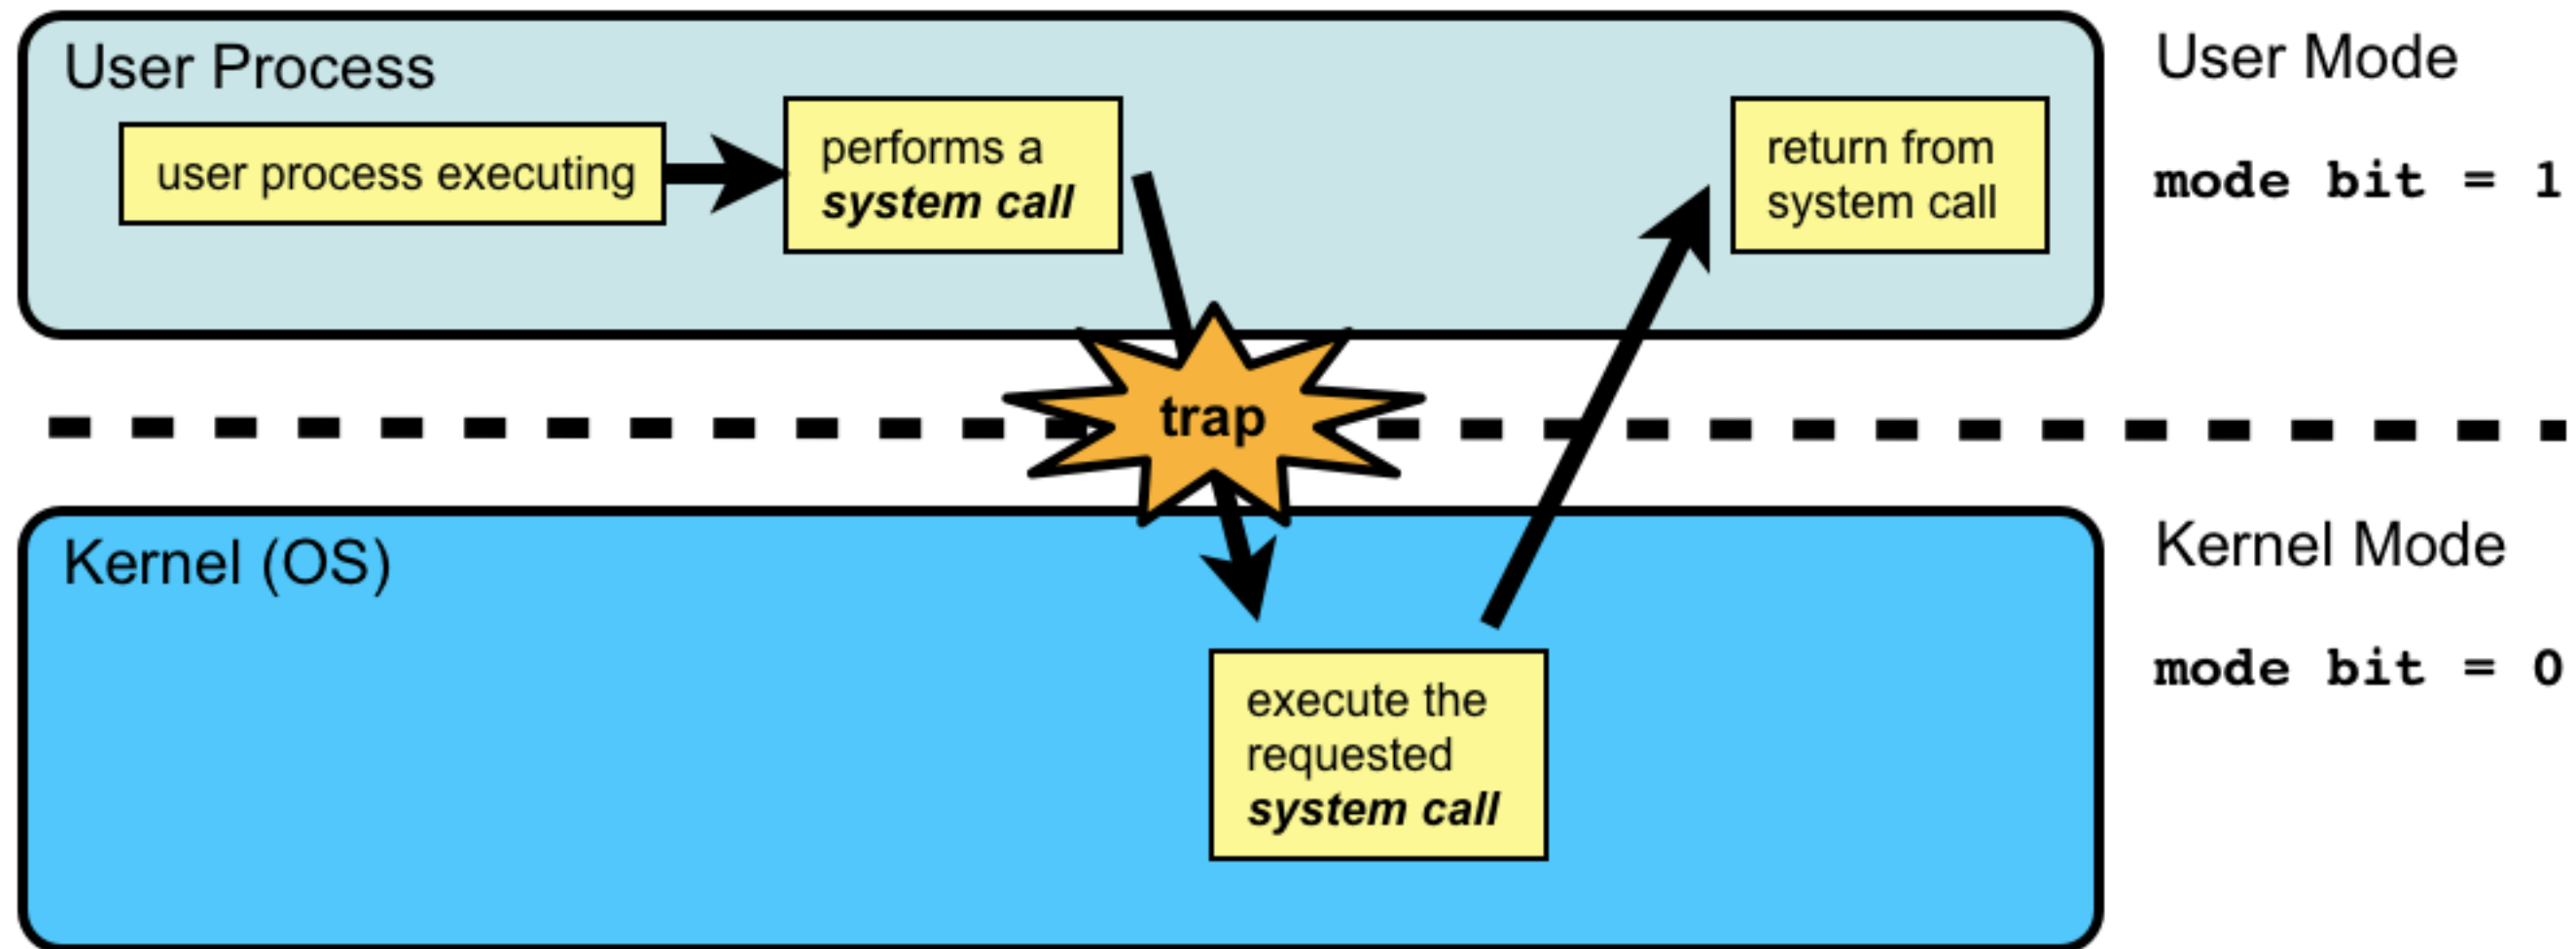


In Unix-like system a file is represented by exactly one **inode**.



- ▶ A file does not have a name.
- ▶ The file is uniquely identified by its inode number.
- ▶ The "file name" is a property of the directory.

To obtain information about files and directories there exist a number **system calls**.



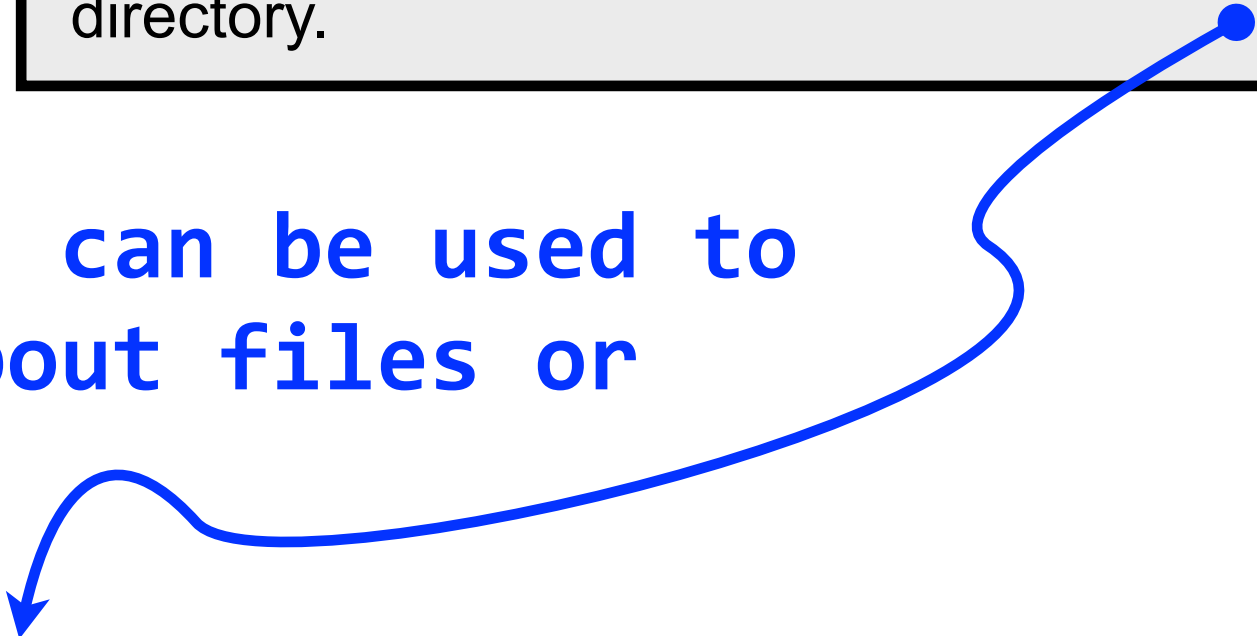
# **C programming**

## **with**

## **stat() & lstat()**

```
#include <sys/types.h>
#include <sys/stat.h>
```


Relative or absolute path to file or directory.



```
// Two system calls that can be used to
// obtain information about files or
// directories.
```

```
int  stat(const char *path, struct stat *buf);
int  lstat(const char *path, struct stat *buf);
```

```
// Both stat() & lstat() returns 0 on success
// and -1 on failure.
```



Must pass a **pointer** to a **stat struct** – will hold result of the system call.

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
```

```
struct stat fstat;
int status;
```

```
status = stat("./file.txt", &fstat);
```

```
printf("MODE = %d \n", (int) fstat.st_mode);
```

```
// Convert the numeric mode to the standard
// string representation such as "-rw-r--r--".
```

```
printf("MODE = %s \n", strmode(fstat.st_mode));
```

Don't forget to use the address-of operator & to get the pointer to **fstat**.

Now the **fstat** struct will be populated with data.

Need to implement a **strmode()** function that converts the numeric mode to string.

See the man page for the stat struct for more information

```
$ man -s2 stat
```

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* inode number */
    mode_t     st_mode;         /* protection */
    nlink_t    st_nlink;        /* number of hard links */
    uid_t      st_uid;          /* user ID of owner */
    gid_t      st_gid;          /* group ID of owner */
    dev_t      st_rdev;         /* device ID (if special file) */
    off_t      st_size;         /* total size, in bytes */
    blksize_t  st_blksize;      /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;       /* number of 512B blocks allocated */
    .
    .
    .
};
```

# User name

How can we get the user name (string) of the file owner?

```
$> ls -l
```


```
-rw----- 1 hans it readme.txt
```

```
-rwx-r--r-- 1 karl it script.sh
```

See the man page for the stat struct for more information

```
$ man -s2 stat
```

```
struct stat {
    dev_t      st_dev;          /* ID of device containing file */
    ino_t      st_ino;          /* inode number */
    mode_t     st_mode;         /* protection */
    nlink_t    st_nlink;        /* number of hard links */
    uid_t      st_uid;          /* user ID of owner */
    gid_t      st_gid;          /* group ID of owner */
    dev_t      st_rdev;         /* device ID (if special file) */
    off_t      st_size;         /* total size, in bytes */
    blksize_t  st_blksize;      /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;       /* number of 512B blocks allocated */
    .
    .
    .
};
```





# getpwnuid()

Convert a numerical user id to username string.

The **getpwnuid()** function shall return a pointer to a **struct passwd** with the structure as defined in **<pwd.h>** with a matching entry if found.

```
struct passwd *getpwnuid(uid_t uid);
```

# pwd.h

The `<pwd.h>` header shall provide a definition for **struct passwd**, which shall include at least the following members:

<code>char</code>	<code>*pw_name</code>	<code>// User's login name.</code>
<code>uid_t</code>	<code>pw_uid</code>	<code>// Numerical user ID.</code>
<code>gid_t</code>	<code>pw_gid</code>	<code>// Numerical group ID.</code>
<code>char</code>	<code>*pw_dir</code>	<code>// Initial working directory.</code>
<code>char</code>	<code>*pw_shell</code>	<code>// Program to use as shell.</code>

# Get user name of the file owner

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <pwd.h>
```

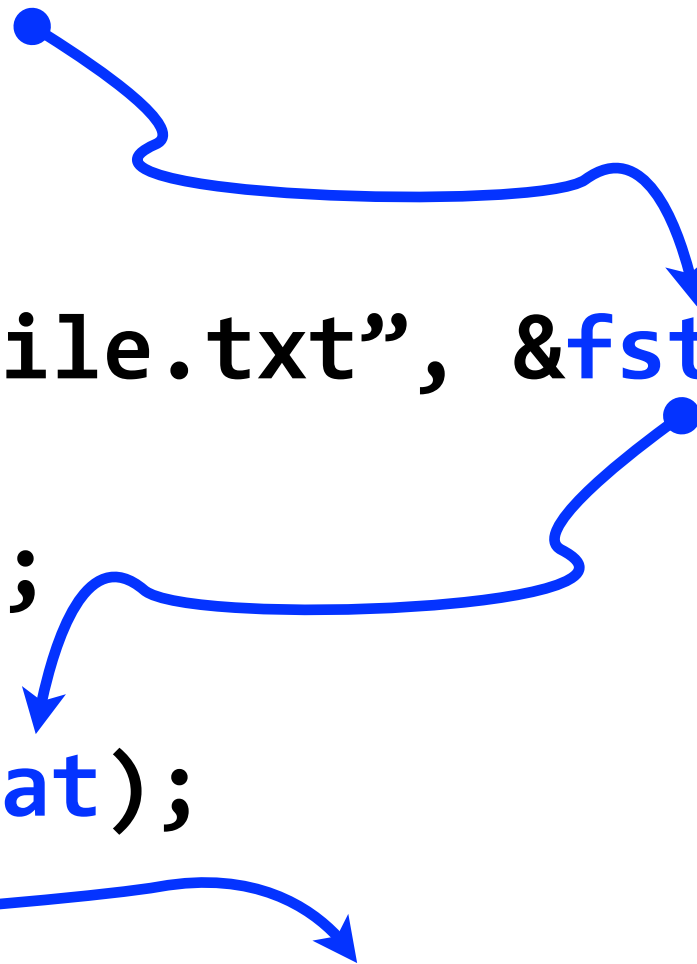
```
struct stat fstat;
int status;
```

```
status = stat("./file.txt", &fstat);
```

```
struct passwd *pwd;
```

```
pwd = getpwuid(fstat);
```

```
printf("USER = %s \n", pwd->pw_name);
```



# Obtaining information about links

```
$> ls -l
```

```
lrw----- 1 karl it link -> file.txt  
-rw-r--r-- 1 karl it file.txt
```

- ▶ Directory entry **link** is a symbolic link to **file.txt**.
- ▶ How to find out if something is a symbolic link?
- ▶ For a symbolic link, to what does the link refer?

# The difference between `stat()` & `lstat()`

```
int  stat(const char *path, struct stat *buf);
```

```
int  lstat(const char *path, struct stat *buf);
```

- ▶ The `lstat()` function shall be equivalent to `stat()`, except when path refers to a symbolic link.
- ▶ In that case `lstat()` shall return information about the link, while `stat()` shall return information about the file the link references.

# Obtaining information about links

```
#define MAX_LINK_LEN 20
```

```
struct stat fstat;  
lstat("link", &fstat);
```

Use **lstat()** so we can get info about symbolic links.

**S\_ISLNK(m)** macro that tests whether file with mode **m** is a symbolic link or not.

```
// Using lstat() we can check if a file  
// is a symbolic link or not:
```

```
if (S_ISLNK(fstat.st_mode)) {  
    char link_buffer[MAX_LINK_LEN];  
    int len;  
    len = readlink("link", link_buffer, MAX_LINK_LEN);  
    link_buffer[len] = 0; // Null terminate  
    printf("link -> %s\n", link_buffer);  
}
```

**readlink()** system call used to get contents of symbolic link.